
Linux Dev-tools Documentation

Release 6.8.0

The kernel development community

Jan 16, 2026

TABLE OF CONTENTS

1	Kernel Testing Guide	3
1.1	Writing and Running Tests	3
1.2	Code Coverage Tools	4
1.3	Dynamic Analysis Tools	4
1.4	Static Analysis Tools	5
2	Checkpatch	7
2.1	Options	7
2.2	Message Levels	10
2.3	Type Descriptions	10
3	Coccinelle	29
3.1	Getting Coccinelle	29
3.2	Supplemental documentation	30
3.3	Using Coccinelle on the Linux kernel	30
3.4	Coccinelle parallelization	31
3.5	Using Coccinelle with a single semantic patch	31
3.6	Controlling Which Files are Processed by Coccinelle	32
3.7	Debugging Coccinelle SmPL patches	32
3.8	.cocciconfig support	33
3.9	Additional flags	34
3.10	SmPL patch specific options	34
3.11	SmPL patch Coccinelle requirements	35
3.12	Proposing new semantic patches	35
3.13	Detailed description of the report mode	35
3.14	Detailed description of the patch mode	36
3.15	Detailed description of the context mode	37
3.16	Detailed description of the org mode	37
4	Sparse	39
4.1	Using sparse for typechecking	39
4.2	Using sparse for lock checking	40
4.3	Getting sparse	40
4.4	Using sparse	40
5	KCOV: code coverage for fuzzing	41
5.1	Prerequisites	41
5.2	Coverage collection	42
5.3	Comparison operands collection	44

5.4	Remote coverage collection	45
6	Using gcov with the Linux kernel	49
6.1	Preparation	49
6.2	Customization	50
6.3	Files	50
6.4	Modules	50
6.5	Separated build and test machines	51
6.6	Note on compilers	52
6.7	Troubleshooting	52
6.8	Appendix A: gather_on_build.sh	52
6.9	Appendix B: gather_on_test.sh	53
7	Kernel Address Sanitizer (KASAN)	55
7.1	Overview	55
7.2	Support	56
7.3	Usage	56
7.4	Implementation details	60
7.5	Shadow memory	61
7.6	For developers	62
8	Kernel Memory Sanitizer (KMSAN)	65
8.1	Usage	65
8.2	Support	67
8.3	How KMSAN works	67
8.4	References	73
9	Undefined Behavior Sanitizer - UBSAN	75
9.1	Report example	75
9.2	Usage	76
9.3	References	76
10	Kernel Memory Leak Detector	77
10.1	Usage	77
10.2	Basic Algorithm	78
10.3	Testing specific sections with kmemleak	79
10.4	Freeing kmemleak internal objects	79
10.5	Kmemleak API	79
10.6	Dealing with false positives/negatives	80
10.7	Limitations and Drawbacks	80
10.8	Testing with kmemleak-test	81
11	Kernel Concurrency Sanitizer (KCSAN)	83
11.1	Usage	83
11.2	Data Races	86
11.3	Race Detection Beyond Data Races	87
11.4	Implementation Details	90
11.5	Alternatives Considered	92
12	Kernel Electric-Fence (KFENCE)	93
12.1	Usage	93
12.2	Implementation Details	97
12.3	Interface	98

12.4	Related Tools	101
13	Debugging kernel and modules via gdb	103
13.1	Requirements	103
13.2	Setup	103
13.3	Examples of using the Linux-provided gdb helpers	104
13.4	List of commands and functions	106
14	Using kgdb, kdb and the kernel debugger internals	107
14.1	Introduction	107
14.2	Compiling a kernel	108
14.3	Kernel Debugger Boot Arguments	109
14.4	Using kdb	113
14.5	Using kgdb / gdb	115
14.6	kgdb and kdb interoperability	116
14.7	kgdb Test Suite	118
14.8	Kernel Debugger Internals	118
14.9	Credits	127
15	Linux Kernel Selftests	129
15.1	Running the selftests (hotplug tests are run in limited mode)	129
15.2	Running a subset of selftests	130
15.3	Running the full range hotplug selftests	131
15.4	Install selftests	131
15.5	Running installed selftests	131
15.6	Timeout for selftests	132
15.7	Packaging selftests	132
15.8	Contributing new tests	133
15.9	Contributing new tests (details)	133
15.10	Test Module	134
15.11	Test Harness	135
16	KUnit - Linux Kernel Unit Testing	147
16.1	Getting Started	147
16.2	KUnit Architecture	152
16.3	Running tests with kunit_tool	156
16.4	Run Tests without kunit_tool	161
16.5	Writing Tests	163
16.6	Common Patterns	167
16.7	API Reference	178
16.8	Test Style and Nomenclature	212
16.9	Frequently Asked Questions	215
16.10	Tips For Running KUnit Tests	217
16.11	Introduction	225
16.12	Unit Testing	225
16.13	How do I use it?	226
17	The Kernel Test Anything Protocol (KTAP), version 1	227
17.1	Version lines	227
17.2	Plan lines	228
17.3	Test case result lines	228
17.4	Diagnostic lines	229

17.5	Unknown lines	230
17.6	Nested tests	230
17.7	Major differences between TAP and KTAP	231
17.8	Example KTAP output	231
17.9	See also:	232
18	UAPI Checker	233
18.1	Options	233
18.2	Examples	234
18.3	Caveats	240
Index		243

This document is a collection of documents about development tools that can be used to work on the kernel. For now, the documents have been pulled together without any significant effort to integrate them into a coherent whole; patches welcome!

A brief overview of testing-specific tools can be found in [*Kernel Testing Guide*](#)

KERNEL TESTING GUIDE

There are a number of different tools for testing the Linux kernel, so knowing when to use each of them can be a challenge. This document provides a rough overview of their differences, and how they fit together.

1.1 Writing and Running Tests

The bulk of kernel tests are written using either the kselftest or KUnit frameworks. These both provide infrastructure to help make running tests and groups of tests easier, as well as providing helpers to aid in writing new tests.

If you're looking to verify the behaviour of the Kernel — particularly specific parts of the kernel — then you'll want to use KUnit or kselftest.

1.1.1 The Difference Between KUnit and kselftest

KUnit ([KUnit - Linux Kernel Unit Testing](#)) is an entirely in-kernel system for "white box" testing: because test code is part of the kernel, it can access internal structures and functions which aren't exposed to userspace.

KUnit tests therefore are best written against small, self-contained parts of the kernel, which can be tested in isolation. This aligns well with the concept of 'unit' testing.

For example, a KUnit test might test an individual kernel function (or even a single codepath through a function, such as an error handling case), rather than a feature as a whole.

This also makes KUnit tests very fast to build and run, allowing them to be run frequently as part of the development process.

There is a KUnit test style guide which may give further pointers in [Test Style and Nomenclature](#). kselftest ([Linux Kernel Selftests](#)), on the other hand, is largely implemented in userspace, and tests are normal userspace scripts or programs.

This makes it easier to write more complicated tests, or tests which need to manipulate the overall system state more (e.g., spawning processes, etc.). However, it's not possible to call kernel functions directly from kselftest. This means that only kernel functionality which is exposed to userspace somehow (e.g. by a syscall, device, filesystem, etc.) can be tested with kselftest. To work around this, some tests include a companion kernel module which exposes more information or functionality. If a test runs mostly or entirely within the kernel, however, KUnit may be the more appropriate tool.

kseltest is therefore suited well to tests of whole features, as these will expose an interface to userspace, which can be tested, but not implementation details. This aligns well with 'system' or 'end-to-end' testing.

For example, all new system calls should be accompanied by kselftest tests.

1.2 Code Coverage Tools

The Linux Kernel supports two different code coverage measurement tools. These can be used to verify that a test is executing particular functions or lines of code. This is useful for determining how much of the kernel is being tested, and for finding corner-cases which are not covered by the appropriate test.

Using gcov with the Linux kernel is GCC's coverage testing tool, which can be used with the kernel to get global or per-module coverage. Unlike KCOV, it does not record per-task coverage. Coverage data can be read from debugfs, and interpreted using the usual gcov tooling.

KCOV: code coverage for fuzzing is a feature which can be built in to the kernel to allow capturing coverage on a per-task level. It's therefore useful for fuzzing and other situations where information about code executed during, for example, a single syscall is useful.

1.3 Dynamic Analysis Tools

The kernel also supports a number of dynamic analysis tools, which attempt to detect classes of issues when they occur in a running kernel. These typically each look for a different class of bugs, such as invalid memory accesses, concurrency issues such as data races, or other undefined behaviour like integer overflows.

Some of these tools are listed below:

- kmemleak detects possible memory leaks. See *Kernel Memory Leak Detector*
- KASAN detects invalid memory accesses such as out-of-bounds and use-after-free errors. See *Kernel Address Sanitizer (KASAN)*
- UBSAN detects behaviour that is undefined by the C standard, like integer overflows. See *Undefined Behavior Sanitizer - UBSAN*
- KCSAN detects data races. See *Kernel Concurrency Sanitizer (KCSAN)*
- KFENCE is a low-overhead detector of memory issues, which is much faster than KASAN and can be used in production. See *Kernel Electric-Fence (KFENCE)*
- lockdep is a locking correctness validator. See Documentation/locking/lockdep-design.rst
- There are several other pieces of debug instrumentation in the kernel, many of which can be found in lib/Kconfig.debug

These tools tend to test the kernel as a whole, and do not "pass" like kselftest or KUnit tests. They can be combined with KUnit or kselftest by running tests on a kernel with these tools enabled: you can then be sure that none of these errors are occurring during the test.

Some of these tools integrate with KUnit or kselftest and will automatically fail tests if an issue is detected.

1.4 Static Analysis Tools

In addition to testing a running kernel, one can also analyze kernel source code directly (**at compile time**) using **static analysis** tools. The tools commonly used in the kernel allow one to inspect the whole source tree or just specific files within it. They make it easier to detect and fix problems during the development process.

Sparse can help test the kernel by performing type-checking, lock checking, value range checking, in addition to reporting various errors and warnings while examining the code. See the [Sparse](#) documentation page for details on how to use it.

Smatch extends Sparse and provides additional checks for programming logic mistakes such as missing breaks in switch statements, unused return values on error checking, forgetting to set an error code in the return of an error path, etc. Smatch also has tests against more serious issues such as integer overflows, null pointer dereferences, and memory leaks. See the project page at <http://smatch.sourceforge.net/>.

Coccinelle is another static analyzer at our disposal. Coccinelle is often used to aid refactoring and collateral evolution of source code, but it can also help to avoid certain bugs that occur in common code patterns. The types of tests available include API tests, tests for correct usage of kernel iterators, checks for the soundness of free operations, analysis of locking behavior, and further tests known to help keep consistent kernel usage. See the [Coccinelle](#) documentation page for details.

Beware, though, that static analysis tools suffer from **false positives**. Errors and warns need to be evaluated carefully before attempting to fix them.

1.4.1 When to use Sparse and Smatch

Sparse does type checking, such as verifying that annotated variables do not cause endianness bugs, detecting places that use `_user` pointers improperly, and analyzing the compatibility of symbol initializers.

Smatch does flow analysis and, if allowed to build the function database, it also does cross function analysis. Smatch tries to answer questions like where is this buffer allocated? How big is it? Can this index be controlled by the user? Is this variable larger than that variable?

It's generally easier to write checks in Smatch than it is to write checks in Sparse. Nevertheless, there are some overlaps between Sparse and Smatch checks.

1.4.2 Strong points of Smatch and Coccinelle

Coccinelle is probably the easiest for writing checks. It works before the pre-processor so it's easier to check for bugs in macros using Coccinelle. Coccinelle also creates patches for you, which no other tool does.

For example, with Coccinelle you can do a mass conversion from `kmalloc(x * size, GFP_KERNEL)` to `kmalloc_array(x, size, GFP_KERNEL)`, and that's really useful. If you just created a Smatch warning and try to push the work of converting on to the maintainers they would be annoyed. You'd have to argue about each warning if can really overflow or not.

Coccinelle does no analysis of variable values, which is the strong point of Smatch. On the other hand, Coccinelle allows you to do simple things in a simple way.

CHECKPATCH

Checkpatch (scripts/checkpatch.pl) is a perl script which checks for trivial style violations in patches and optionally corrects them. Checkpatch can also be run on file contexts and without the kernel tree.

Checkpatch is not always right. Your judgement takes precedence over checkpatch messages. If your code looks better with the violations, then its probably best left alone.

2.1 Options

This section will describe the options checkpatch can be run with.

Usage:

```
./scripts/checkpatch.pl [OPTION]... [FILE]...
```

Available options:

- **-q, --quiet**
Enable quiet mode.
- **-v, --verbose** Enable verbose mode. Additional verbose test descriptions are output so as to provide information on why that particular message is shown.
- **--no-tree**
Run checkpatch without the kernel tree.
- **--no-signoff**

Disable the 'Signed-off-by' line check. The sign-off is a simple line at the end of the explanation for the patch, which certifies that you wrote it or otherwise have the right to pass it on as an open-source patch.

Example:

```
Signed-off-by: Random J Developer <random@developer.example.org>
```

Setting this flag effectively stops a message for a missing signed-off-by line in a patch context.

- **--patch**

Treat FILE as a patch. This is the default option and need not be explicitly specified.

- `--emacs`
Set output to emacs compile window format. This allows emacs users to jump from the error in the compile window directly to the offending line in the patch.
- `--terse`
Output only one line per report.
- `--showfile`
Show the diffed file position instead of the input file position.
- `-g, --git`
Treat FILE as a single commit or a git revision range.

Single commit with:

- <rev>
- <rev>^
- <rev>~n

Multiple commits with:

- <rev1>..<rev2>
- <rev1>...<rev2>
- <rev>-<count>

- `-f, --file`
Treat FILE as a regular source file. This option must be used when running checkpatch on source files in the kernel.

- `--subjective, --strict`

Enable stricter tests in checkpatch. By default the tests emitted as CHECK do not activate by default. Use this flag to activate the CHECK tests.

- `--list-types`

Every message emitted by checkpatch has an associated TYPE. Add this flag to display all the types in checkpatch.

Note that when this flag is active, checkpatch does not read the input FILE, and no message is emitted. Only a list of types in checkpatch is output.

- `--types TYPE(TYPE2...)`

Only display messages with the given types.

Example:

```
./scripts/checkpatch.pl mypatch.patch --types EMAIL SUBJECT,BRACES
```

- `--ignore TYPE(TYPE2...)`

Checkpatch will not emit messages for the specified types.

Example:

```
./scripts/checkpatch.pl mypatch.patch --ignore EMAIL SUBJECT,BRACES
```

- `--show-types`

By default checkpatch doesn't display the type associated with the messages. Set this flag to show the message type in the output.

- `--max-line-length=n`

Set the max line length (default 100). If a line exceeds the specified length, a LONG_LINE message is emitted.

The message level is different for patch and file contexts. For patches, a WARNING is emitted. While a milder CHECK is emitted for files. So for file contexts, the `--strict` flag must also be enabled.

- `--min-conf-desc-length=n`

Set the Kconfig entry minimum description length, if shorter, warn.

- `--tab-size=n`

Set the number of spaces for tab (default 8).

- `--root=PATH`

PATH to the kernel tree root.

This option must be specified when invoking checkpatch from outside the kernel root.

- `--no-summary`

Suppress the per file summary.

- `--mailback`

Only produce a report in case of Warnings or Errors. Milder Checks are excluded from this.

- `--summary-file`

Include the filename in summary.

- `--debug KEY=[0|1]`

Turn on/off debugging of KEY, where KEY is one of 'values', 'possible', 'type', and 'attr' (default is all off).

- `--fix`

This is an EXPERIMENTAL feature. If correctable errors exists, a file <inputfile>.EXPERIMENTAL-checkpatch-fixes is created which has the automatically fixable errors corrected.

- `--fix-inplace`

EXPERIMENTAL - Similar to `--fix` but input file is overwritten with fixes.

DO NOT USE this flag unless you are absolutely sure and you have a backup in place.

- `--ignore-perl-version`

Override checking of perl version. Runtime errors maybe encountered after enabling this flag if the perl version does not meet the minimum specified.

- `--codespell`
Use the codespell dictionary for checking spelling errors.
- `--codespellfile`
Use the specified codespell file. Default is '/usr/share/codespell/dictionary.txt'.
- `--typedefsfile`
Read additional types from this file.
- `--color[=WHEN]`
Use colors 'always', 'never', or only when output is a terminal ('auto'). Default is 'auto'.
- `--kconfig-prefix=WORD`
Use WORD as a prefix for Kconfig symbols (default is *CONFIG_*).
- `-h, --help, --version`
Display the help text.

2.2 Message Levels

Messages in checkpatch are divided into three levels. The levels of messages in checkpatch denote the severity of the error. They are:

- **ERROR**

This is the most strict level. Messages of type ERROR must be taken seriously as they denote things that are very likely to be wrong.

- **WARNING**

This is the next stricter level. Messages of type WARNING requires a more careful review. But it is milder than an ERROR.

- **CHECK**

This is the mildest level. These are things which may require some thought.

2.3 Type Descriptions

This section contains a description of all the message types in checkpatch.

2.3.1 Allocation style

ALLOC_ARRAY_ARGS

The first argument for kcalloc or kmalloc_array should be the number of elements. sizeof() as the first argument is generally wrong.

See: <https://www.kernel.org/doc/html/latest/core-api/memory-allocation.html>

ALLOC_SIZEOF_STRUCT

The allocation style is bad. In general for family of allocation functions using sizeof() to get memory size, constructs like:

```
p = alloc(sizeof(struct foo), ...)
```

should be:

```
p = alloc(sizeof(*p), ...)
```

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#allocating-memory>

ALLOC_WITH_MULTIPLY

Prefer kmalloc_array/kcalloc over kmalloc/kzalloc with a sizeof multiply.

See: <https://www.kernel.org/doc/html/latest/core-api/memory-allocation.html>

2.3.2 API usage

ARCH_DEFINES

Architecture specific defines should be avoided wherever possible.

ARCH_INCLUDE_LINUX

Whenever asm/file.h is included and linux/file.h exists, a conversion can be made when linux/file.h includes asm/file.h. However this is not always the case (See signal.h). This message type is emitted only for includes from arch/.

AVOID_BUG

BUG() or BUG_ON() should be avoided totally. Use WARN() and WARN_ON() instead, and handle the "impossible" error condition as gracefully as possible.

See: <https://www.kernel.org/doc/html/latest/process/deprecated.html#bug-and-bug-on>

CONSIDER_KSTRTO

The simple_strtol(), simple strtoll(), simple strtoul(), and simple strtoull() functions explicitly ignore overflows, which may lead to unexpected results in callers. The respective kstrtol(), kstrtoll(), kstrtoul(), and kstrtoull() functions tend to be the correct replacements.

See: <https://www.kernel.org/doc/html/latest/process/deprecated.html#simple-strtol-simple-strtoll-simple-strtoul-simple-strtoull>

CONSTANT_CONVERSION

Use of __constant_<foo> form is discouraged for the following functions:

```
__constant_cpu_to_be[x]
__constant_cpu_to_le[x]
__constant_be[x]_to_cpu
__constant_le[x]_to_cpu
__constant_htons
__constant_ntohs
```

Using any of these outside of include/uapi/ is not preferred as using the function without `_constant_` is identical when the argument is a constant.

In big endian systems, the macros like `_constant_cpu_to_be32(x)` and `cpu_to_be32(x)` expand to the same expression:

```
#define __constant_cpu_to_be32(x) (((__force __be32)(__u32)(x))
#define __cpu_to_be32(x)           (((__force __be32)(__u32)(x))
```

In little endian systems, the macros `_constant_cpu_to_be32(x)` and `cpu_to_be32(x)` expand to `_constant_swab32` and `_swab32`. `_swab32` has a `_builtin_constant_p` check:

```
#define __swab32(x) \
    (__builtin_constant_p((__u32)(x)) ? \
     __constant_swab32(x) : \
     __fswab32(x)) \
```

So ultimately they have a special case for constants. Similar is the case with all of the macros in the list. Thus using the `_constant_...` forms are unnecessarily verbose and not preferred outside of include/uapi.

See: <https://lore.kernel.org/lkml/1400106425.12666.6.camel@joe-AO725/>

DEPRECATED API

Usage of a deprecated RCU API is detected. It is recommended to replace old flavourful RCU APIs by their new vanilla-RCU counterparts.

The full list of available RCU APIs can be viewed from the kernel docs.

See: <https://www.kernel.org/doc/html/latest/RCU/whatisRCU.html#full-list-of-rcu-apis>

DEPRECATED VARIABLE

`EXTRA_{A,C,CPP,LD}FLAGS` are deprecated and should be replaced by the new flags added via [commit f77bf01425b1](#) ("kbuild: introduce ccflags-y, asflags-y and ldflags-y").

The following conversion scheme maybe used:

<code>EXTRA_AFLAGS</code>	<code>-> asflags-y</code>
<code>EXTRA_CFLAGS</code>	<code>-> ccflags-y</code>
<code>EXTRA_CPPFLAGS</code>	<code>-> cppflags-y</code>
<code>EXTRA_LDFLAGS</code>	<code>-> ldflags-y</code>

See:

1. <https://lore.kernel.org/lkml/20070930191054.GA15876@uranus.ravnborg.org/>

2. <https://lore.kernel.org/lkml/1313384834-24433-12-git-send-email-lacombar@gmail.com/>
3. <https://www.kernel.org/doc/html/latest/kbuild/makefiles.html#compilation-flags>

DEVICE_ATTR_FUNCTIONS

The function names used in DEVICE_ATTR is unusual. Typically, the store and show functions are used with <attr>_store and <attr>_show, where <attr> is a named attribute variable of the device.

Consider the following examples:

```
static DEVICE_ATTR(type, 0444, type_show, NULL);
static DEVICE_ATTR(power, 0644, power_show, power_store);
```

The function names should preferably follow the above pattern.

See: <https://www.kernel.org/doc/html/latest/driver-api/driver-model/device.html#attributes>

DEVICE_ATTR_RO

The DEVICE_ATTR_RO(name) helper macro can be used instead of DEVICE_ATTR(name, 0444, name_show, NULL);

Note that the macro automatically appends _show to the named attribute variable of the device for the show method.

See: <https://www.kernel.org/doc/html/latest/driver-api/driver-model/device.html#attributes>

DEVICE_ATTR_RW

The DEVICE_ATTR_RW(name) helper macro can be used instead of DEVICE_ATTR(name, 0644, name_show, name_store);

Note that the macro automatically appends _show and _store to the named attribute variable of the device for the show and store methods.

See: <https://www.kernel.org/doc/html/latest/driver-api/driver-model/device.html#attributes>

DEVICE_ATTR_WO

The DEVICE_ATTR_WO(name) helper macro can be used instead of DEVICE_ATTR(name, 0200, NULL, name_store);

Note that the macro automatically appends _store to the named attribute variable of the device for the store method.

See: <https://www.kernel.org/doc/html/latest/driver-api/driver-model/device.html#attributes>

DUPLICATED_SYSCTL_CONST

Commit d91bff3011cf ("proc/sysctl: add shared variables for range check") added some shared const variables to be used instead of a local copy in each source file.

Consider replacing the sysctl range checking value with the shared one in include/linux/sysctl.h. The following conversion scheme may be used:

```
&zero      -> SYSCTL_ZERO
&one       -> SYSCTL_ONE
&int_max   -> SYSCTL_INT_MAX
```

See:

1. <https://lore.kernel.org/lkml/20190430180111.10688-1-mcroce@redhat.com/>
2. <https://lore.kernel.org/lkml/20190531131422.14970-1-mcroce@redhat.com/>

ENOSYS

ENOSYS means that a nonexistent system call was called. Earlier, it was wrongly used for things like invalid operations on otherwise valid syscalls. This should be avoided in new code.

See: <https://lore.kernel.org/lkml/5eb299021dec23c1a48fa7d9f2c8b794e967766d.1408730669.git.luto@amacapital.net/>

ENOTSUPP

ENOTSUPP is not a standard error code and should be avoided in new patches. EOPNOTSUPP should be used instead.

See: <https://lore.kernel.org/netdev/20200510182252.GA411829@lunn.ch/>

EXPORT_SYMBOL

EXPORT_SYMBOL should immediately follow the symbol to be exported.

IN_ATOMIC

in_atomic() is not for driver use so any such use is reported as an ERROR. Also in_atomic() is often used to determine if sleeping is permitted, but it is not reliable in this use model. Therefore its use is strongly discouraged.

However, in_atomic() is ok for core kernel use.

See: <https://lore.kernel.org/lkml/20080320201723.b87b3732.akpm@linux-foundation.org/>

LOCKDEP

The lockdep_no_validate class was added as a temporary measure to prevent warnings on conversion of device->sem to device->mutex. It should not be used for any other purpose.

See: <https://lore.kernel.org/lkml/1268959062.9440.467.camel@laptop/>

MALFORMED_INCLUDE

The #include statement has a malformed path. This has happened because the author has included a double slash “//” in the pathname accidentally.

USE_LOCKDEP

lockdep_assert_held() annotations should be preferred over assertions based on spin_is_locked()

See: <https://www.kernel.org/doc/html/latest/locking/lockdep-design.html#annotations>

UAPI_INCLUDE

No #include statements in include/uapi should use a uapi/ path.

USLEEP_RANGE

`usleep_range()` should be preferred over `udelay()`. The proper way of using `usleep_range()` is mentioned in the kernel docs.

See: <https://www.kernel.org/doc/html/latest/timers/timers-howto.html#delays-information-on-the-various-kernel-delay-sleep-mechanisms>

2.3.3 Comments**BLOCK_COMMENT_STYLE**

The comment style is incorrect. The preferred style for multi-line comments is:

```
/*
 * This is the preferred style
 * for multi line comments.
*/
```

The networking comment style is a bit different, with the first line not empty like the former:

```
/* This is the preferred comment style
 * for files in net/ and drivers/net/
 */
```

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#commenting>

C99_COMMENTS

C99 style single line comments (`//`) should not be used. Prefer the block comment style instead.

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#commenting>

DATA_RACE

Applications of `data_race()` should have a comment so as to document the reasoning behind why it was deemed safe.

See: <https://lore.kernel.org/lkml/20200401101714.44781-1-elver@google.com/>

FSF_MAILING_ADDRESS

Kernel maintainers reject new instances of the GPL boilerplate paragraph directing people to write to the FSF for a copy of the GPL, since the FSF has moved in the past and may do so again. So do not write paragraphs about writing to the Free Software Foundation's mailing address.

See: <https://lore.kernel.org/lkml/20131006222342.GT19510@leaf/>

2.3.4 Commit message

BAD_SIGN_OFF

The signed-off-by line does not fall in line with the standards specified by the community.

See: <https://www.kernel.org/doc/html/latest/process/submitting-patches.html#developer-s-certificate-of-origin-1-1>

BAD_STABLE_ADDRESS_STYLE

The email format for stable is incorrect. Some valid options for stable address are:

1. stable@vger.kernel.org
2. stable@kernel.org

For adding version info, the following comment style should be used:

```
stable@vger.kernel.org # version info
```

COMMIT_COMMENT_SYMBOL

Commit log lines starting with a '#' are ignored by git as comments. To solve this problem addition of a single space in front of the log line is enough.

COMMIT_MESSAGE

The patch is missing a commit description. A brief description of the changes made by the patch should be added.

See: <https://www.kernel.org/doc/html/latest/process/submitting-patches.html#describe-your-changes>

EMAIL SUBJECT

Naming the tool that found the issue is not very useful in the subject line. A good subject line summarizes the change that the patch brings.

See: <https://www.kernel.org/doc/html/latest/process/submitting-patches.html#describe-your-changes>

FROM SIGN OFF MISMATCH

The author's email does not match with that in the Signed-off-by: line(s). This can be sometimes caused due to an improperly configured email client.

This message is emitted due to any of the following reasons:

- The email names do not match.
- The email addresses do not match.
- The email subaddresses do not match.
- The email comments do not match.

MISSING_SIGN_OFF

The patch is missing a Signed-off-by line. A signed-off-by line should be added according to Developer's certificate of Origin.

See: <https://www.kernel.org/doc/html/latest/process/submitting-patches.html#sign-your-work-the-developer-s-certificate-of-origin>

NO_AUTHOR_SIGN_OFF

The author of the patch has not signed off the patch. It is required that a simple sign off line should be present at the end of explanation of the patch to denote that the author has written it or otherwise has the rights to pass it on as an open source patch.

See: <https://www.kernel.org/doc/html/latest/process/submitting-patches.html#sign-your-work-the-developer-s-certificate-of-origin>

DIFF_IN_COMMIT_MSG

Avoid having diff content in commit message. This causes problems when one tries to apply a file containing both the changelog and the diff because patch(1) tries to apply the diff which it found in the changelog.

See: <https://lore.kernel.org/lkml/20150611134006.9df79a893e3636019ad2759e@linux-foundation.org/>

GERRIT_CHANGE_ID

To be picked up by gerrit, the footer of the commit message might have a Change-Id like:

```
Change-Id: Ic8aaa0728a43936cd4c6e1ed590e01ba8f0fbf5b
Signed-off-by: A. U. Thor <author@example.com>
```

The Change-Id line must be removed before submitting.

GIT_COMMIT_ID

The proper way to reference a commit id is: commit <12+ chars of sha1> ("<title line>")

An example may be:

```
Commit e21d2170f36602ae2708 ("video: remove unnecessary
platform_set_drvdata()") removed the unnecessary
platform_set_drvdata(), but left the variable "dev" unused,
delete it.
```

See: <https://www.kernel.org/doc/html/latest/process/submitting-patches.html#describe-your-changes>

BAD_FIXES_TAG

The Fixes: tag is malformed or does not follow the community conventions. This can occur if the tag have been split into multiple lines (e.g., when pasted in an email program with word wrapping enabled).

See: <https://www.kernel.org/doc/html/latest/process/submitting-patches.html#describe-your-changes>

2.3.5 Comparison style

ASSIGN_IN_IF

Do not use assignments in if condition. Example:

```
if ((foo = bar(...)) < BAZ) {
```

should be written as:

```
foo = bar(...);
if (foo < BAZ) {
```

BOOL_COMPARISON

Comparisons of A to true and false are better written as A and !A.

See: <https://lore.kernel.org/lkml/1365563834.27174.12.camel@joe-AO722/>

COMPARISON_TO_NULL

Comparisons to NULL in the form (foo == NULL) or (foo != NULL) are better written as (!foo) and (foo).

CONSTANT_COMPARISON

Comparisons with a constant or upper case identifier on the left side of the test should be avoided.

2.3.6 Indentation and Line Breaks

CODE_INDENT

Code indent should use tabs instead of spaces. Outside of comments, documentation and Kconfig, spaces are never used for indentation.

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#indentation>

DEEP_INDENTATION

Indentation with 6 or more tabs usually indicate overly indented code.

It is suggested to refactor excessive indentation of if/else/for/do/while/switch statements.

See: <https://lore.kernel.org/lkml/1328311239.21255.24.camel@joe2Laptop/>

SWITCH_CASE_INDENT_LEVEL

switch should be at the same indent as case. Example:

```
switch (suffix) {
case 'G':
case 'g':
    mem <= 30;
    break;
case 'M':
case 'm':
    mem <= 20;
    break;
case 'K':
```

```

case 'k':
    mem <= 10;
    fallthrough;
default:
    break;
}

```

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#indentation>

LONG_LINE

The line has exceeded the specified maximum length. To use a different maximum line length, the --max-line-length=n option may be added while invoking checkpatch.

Earlier, the default line length was 80 columns. Commit [bdc48fa11e46](#) ("checkpatch/coding-style: deprecate 80-column warning") increased the limit to 100 columns. This is not a hard limit either and it's preferable to stay within 80 columns whenever possible.

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#breaking-long-lines-and-strings>

LONG_LINE_STRING

A string starts before but extends beyond the maximum line length. To use a different maximum line length, the --max-line-length=n option may be added while invoking checkpatch.

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#breaking-long-lines-and-strings>

LONG_LINE_COMMENT

A comment starts before but extends beyond the maximum line length. To use a different maximum line length, the --max-line-length=n option may be added while invoking checkpatch.

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#breaking-long-lines-and-strings>

SPLIT_STRING

Quoted strings that appear as messages in userspace and can be grepped, should not be split across multiple lines.

See: <https://lore.kernel.org/lkml/20120203052727.GA15035@leaf/>

MULTILINE_DEREFERENCE

A single dereferencing identifier spanned on multiple lines like:

```

struct_identifier->member[index].
member = <foo>;

```

is generally hard to follow. It can easily lead to typos and so makes the code vulnerable to bugs.

If fixing the multiple line dereferencing leads to an 80 column violation, then either rewrite the code in a more simple way or if the starting part of the derefer-

encing identifier is the same and used at multiple places then store it in a temporary variable, and use that temporary variable only at all the places. For example, if there are two dereferencing identifiers:

```
member1->member2->member3.foo1;  
member1->member2->member3.foo2;
```

then store the member1->member2->member3 part in a temporary variable. It not only helps to avoid the 80 column violation but also reduces the program size by removing the unnecessary dereferences.

But if none of the above methods work then ignore the 80 column violation because it is much easier to read a dereferencing identifier on a single line.

TRAILING_STATEMENTS

Trailing statements (for example after any conditional) should be on the next line. Statements, such as:

```
if (x == y) break;
```

should be:

```
if (x == y)  
    break;
```

2.3.7 Macros, Attributes and Symbols

ARRAY_SIZE

The ARRAY_SIZE(foo) macro should be preferred over sizeof(foo)/sizeof(foo[0]) for finding number of elements in an array.

The macro is defined in include/linux/kernel.h:

```
#define ARRAY_SIZE(x) (sizeof(x) / sizeof((x)[0]))
```

AVOID_EXTERNS

Function prototypes don't need to be declared extern in .h files. It's assumed by the compiler and is unnecessary.

AVOID_L_PREFIX

Local symbol names that are prefixed with *.L* should be avoided, as this has special meaning for the assembler; a symbol entry will not be emitted into the symbol table. This can prevent *objtool* from generating correct unwind info.

Symbols with STB_LOCAL binding may still be used, and *.L* prefixed local symbol names are still generally usable within a function, but *.L* prefixed local symbol names should not be used to denote the beginning or end of code regions via *SYM_CODE_START_LOCAL/SYM_CODE_END*

BIT_MACRO

Defines like: 1 << <digit> could be BIT(digit). The BIT() macro is defined via include/linux/bits.h:

```
#define BIT(nr)          (1UL << (nr))
```

CONST_READ_MOSTLY

When a variable is tagged with the `_read_mostly` annotation, it is a signal to the compiler that accesses to the variable will be mostly reads and rarely (but NOT never) a write.

`const _read_mostly` does not make any sense as `const` data is already read-only. The `_read_mostly` annotation thus should be removed.

DATE_TIME

It is generally desirable that building the same source code with the same set of tools is reproducible, i.e. the output is always exactly the same.

The kernel does *not* use the `_DATE_` and `_TIME_` macros, and enables warnings if they are used as they can lead to non-deterministic builds.

See: <https://www.kernel.org/doc/html/latest/kbuild/reproducible-builds.html#timestamps>

DEFINE_ARCH_HAS

The `ARCH_HAS_xyz` and `ARCH_HAVE_xyz` patterns are wrong.

For big conceptual features use Kconfig symbols instead. And for smaller things where we have compatibility fallback functions but want architectures able to override them with optimized ones, we should either use weak functions (appropriate for some cases), or the symbol that protects them should be the same symbol we use.

See: https://lore.kernel.org/lkml/CA+55aFycQ9XJvEOsiM3txHL5bjUc8CeKWJNR_H+MiicaddB42Q@mail.gmail.com/

DO_WHILE_MACRO_WITH_TRAILING_SEMICOLON

`do {} while(0)` macros should not have a trailing semicolon.

INIT_ATTRIBUTE

Const init definitions should use `_initconst` instead of `_initdata`.

Similarly init definitions without const require a separate use of `const`.

INLINE_LOCATION

The `inline` keyword should sit between storage class and type.

For example, the following segment:

```
inline static int example_function(void)
{
    ...
}
```

should be:

```
static inline int example_function(void)
{
    ...
}
```

MISPLACED_INIT

It is possible to use section markers on variables in a way which gcc doesn't understand (or at least not the way the developer intended):

```
static struct __initdata samsung_pll_clock exynos4_pll[nr_pll] =  
→{
```

does not put exynos4_pll in the .initdata section. The __initdata marker can be virtually anywhere on the line, except right after "struct". The preferred location is before the "=" sign if there is one, or before the trailing ";" otherwise.

See: <https://lore.kernel.org/lkml/1377655732.3619.19.camel@joe-AO722/>

MULTISTATEMENT_MACRO_USE_DO WHILE

Macros with multiple statements should be enclosed in a do - while block. Same should also be the case for macros starting with *if* to avoid logic defects:

```
#define macrofun(a, b, c)           \
    do {                           \
        if (a == 5)                \
            do_this(b, c);         \
    } while (0)
```

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#macros-enums-and-rtl>

PREFER_FALLTHROUGH

Use the *fallthrough*; pseudo keyword instead of /* *fallthrough* */ like comments.

TRAILING_SEMICOLON

Macro definition should not end with a semicolon. The macro invocation style should be consistent with function calls. This can prevent any unexpected code paths:

```
#define MAC do_something;
```

If this macro is used within a if else statement, like:

```
if (some_condition)  
    MAC;  
  
else  
    do_something;
```

Then there would be a compilation error, because when the macro is expanded there are two trailing semicolons, so the else branch gets orphaned.

See: <https://lore.kernel.org/lkml/1399671106.2912.21.camel@joe-AO725/>

SINGLE_STATEMENT_DO WHILE_MACRO

For the multi-statement macros, it is necessary to use the do-while loop to avoid unpredictable code paths. The do-while loop helps to group the multiple statements into a single one so that a function-like macro can be used as a function only.

But for the single statement macros, it is unnecessary to use the do-while loop. Although the code is syntactically correct but using the do-while loop is redundant. So remove the do-while loop for single statement macros.

WEAK_DECLARATION

Using weak declarations like `_attribute_((weak))` or `_weak` can have unintended link defects. Avoid using them.

2.3.8 Functions and Variables

CAMELCASE

Avoid CamelCase Identifiers.

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#naming>

CONST_CONST

Using `const <type> const *` is generally meant to be written `const <type> * const`.

CONST_STRUCT

Using `const` is generally a good idea. Checkpatch reads a list of frequently used structs that are always or almost always constant.

The existing structs list can be viewed from `scripts/const_structs.checkpatch`.

See: <https://lore.kernel.org/lkml/alpine.DEB.2.10.1608281509480.3321@hadrien/>

EMBEDDED_FUNCTION_NAME

Embedded function names are less appropriate to use as refactoring can cause function renaming. Prefer the use of "%s", `_func_` to embedded function names.

Note that this does not work with `-f` (--file) checkpatch option as it depends on patch context providing the function name.

FUNCTION_ARGUMENTS

This warning is emitted due to any of the following reasons:

1. Arguments for the function declaration do not follow the identifier name. Example:

```
void foo
(int bar, int baz)
```

This should be corrected to:

```
void foo(int bar, int baz)
```

2. Some arguments for the function definition do not have an identifier name. Example:

```
void foo(int)
```

All arguments should have identifier names.

FUNCTION_WITHOUT_ARGS

Function declarations without arguments like:

```
int foo()
```

should be:

```
int foo(void)
```

GLOBAL_INITIALISERS

Global variables should not be initialized explicitly to 0 (or NULL, false, etc.). Your compiler (or rather your loader, which is responsible for zeroing out the relevant sections) automatically does it for you.

INITIALISED_STATIC

Static variables should not be initialized explicitly to zero. Your compiler (or rather your loader) automatically does it for you.

MULTIPLE_ASSIGNMENTS

Multiple assignments on a single line makes the code unnecessarily complicated. So on a single line assign value to a single variable only, this makes the code more readable and helps avoid typos.

RETURN_PARENTHESSES

return is not a function and as such doesn't need parentheses:

```
return (bar);
```

can simply be:

```
return bar;
```

2.3.9 Permissions

DEVICE_ATTR_PERMS

The permissions used in DEVICE_ATTR are unusual. Typically only three permissions are used - 0644 (RW), 0444 (RO) and 0200 (WO).

See: <https://www.kernel.org/doc/html/latest/filesystems/sysfs.html#attributes>

EXECUTE_PERMISSIONS

There is no reason for source files to be executable. The executable bit can be removed safely.

EXPORTED_WORLD_WRITABLE

Exporting world writable sysfs/debugfs files is usually a bad thing. When done arbitrarily they can introduce serious security bugs. In the past, some of the debugfs vulnerabilities would seemingly allow any local user to write arbitrary values into device registers - a situation from which little good can be expected to emerge.

See: <https://lore.kernel.org/linux-arm-kernel/cover.1296818921.git.segoon@openwall.com/>

NON_OCTAL_PERMISSIONS

Permission bits should use 4 digit octal permissions (like 0700 or 0444). Avoid using any other base like decimal.

SYMBOLIC_PERMS

Permission bits in the octal form are more readable and easier to understand than their symbolic counterparts because many command-line tools use this notation. Experienced kernel developers have been using these traditional Unix permission bits for decades and so they find it easier to understand the octal notation than the symbolic macros. For example, it is harder to read S_IWUSR|S_IRUGO than 0644, which obscures the developer's intent rather than clarifying it.

See: https://lore.kernel.org/lkml/CA+55aFw5v23T-zvDZp-MmD_EYxF8WbafwwB59934FV7g21uMGQ@mail.gmail.com/

2.3.10 Spacing and Brackets

ASSIGNMENT_CONTINUATIONS

Assignment operators should not be written at the start of a line but should follow the operand at the previous line.

BRACES

The placement of braces is stylistically incorrect. The preferred way is to put the opening brace last on the line, and put the closing brace first:

```
if (x is true) {
    we do y
}
```

This applies for all non-functional blocks. However, there is one special case, namely functions: they have the opening brace at the beginning of the next line, thus:

```
int function(int x)
{
    body of function
}
```

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#placing-braces-and-spaces>

BRACKET_SPACE

Whitespace before opening bracket '[' is prohibited. There are some exceptions:

1. With a type on the left:

```
int [] a;
```

2. At the beginning of a line for slice initialisers:

```
[0...10] = 5,
```

3. Inside a curly brace:

```
= { [0...10] = 5 }
```

CONCATENATED_STRING

Concatenated elements should have a space in between. Example:

```
printf(KERN_INFO"bar");
```

should be:

```
printf(KERN_INFO "bar");
```

ELSE_AFTER_BRACE

else { should follow the closing block *}* on the same line.

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#placing-braces-and-spaces>

LINE_SPACING

Vertical space is wasted given the limited number of lines an editor window can display when multiple blank lines are used.

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#spaces>

OPEN_BRACE

The opening brace should be following the function definitions on the next line. For any non-functional block it should be on the same line as the last construct.

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#placing-braces-and-spaces>

POINTER_LOCATION

When using pointer data or a function that returns a pointer type, the preferred use of * is adjacent to the data name or function name and not adjacent to the type name. Examples:

```
char *linux_banner;
unsigned long long memparse(char *ptr, char **retptr);
char *match_strdup(substring_t *s);
```

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#spaces>

SPACING

Whitespace style used in the kernel sources is described in kernel docs.

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#spaces>

TRAILING_WHITESPACE

Trailing whitespace should always be removed. Some editors highlight the trailing whitespace and cause visual distractions when editing files.

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#spaces>

UNNECESSARY_PARENTHESSES

Parentheses are not required in the following cases:

1. Function pointer uses:

```
(foo->bar)();
```

could be:

```
foo->bar();
```

2. Comparisons in if:

```
if ((foo->bar) && (foo->baz))
if ((foo == bar))
```

could be:

```
if (foo->bar && foo->baz)
if (foo == bar)
```

3. addressof/dereference single Lvalues:

```
&(foo->bar)
*(foo->bar)
```

could be:

```
&foo->bar
*foo->bar
```

WHILE_AFTER_BRACE

while should follow the closing bracket on the same line:

```
do {
    ...
} while(something);
```

See: <https://www.kernel.org/doc/html/latest/process/coding-style.html#placing-braces-and-spaces>

2.3.11 Others

CONFIG_DESCRIPTION

Kconfig symbols should have a help text which fully describes it.

CORRUPTED_PATCH

The patch seems to be corrupted or lines are wrapped. Please regenerate the patch file before sending it to the maintainer.

CVS_KEYWORD

Since linux moved to git, the CVS markers are no longer used. So, CVS style keywords (\$Id\$, \$Revision\$, \$Log\$) should not be added.

DEFAULT_NO_BREAK

switch default case is sometimes written as "default:;". This can cause new cases added below default to be defective.

A "break;" should be added after empty default statement to avoid unwanted fallthrough.

DOS_LINE_ENDINGS

For DOS-formatted patches, there are extra ^M symbols at the end of the line. These should be removed.

DT_SCHEMA_BINDING_PATCH

DT bindings moved to a json-schema based format instead of freeform text.

See: <https://www.kernel.org/doc/html/latest/devicetree/bindings/writing-schema.html>

DT_SPLIT_BINDING_PATCH

Devicetree bindings should be their own patch. This is because bindings are logically independent from a driver implementation, they have a different maintainer (even though they often are applied via the same tree), and it makes for a cleaner history in the DT only tree created with git-filter-branch.

See: <https://www.kernel.org/doc/html/latest/devicetree/bindings/submitting-patches.html#i-for-patch-submitters>

EMBEDDED_FILENAME

Embedding the complete filename path inside the file isn't particularly useful as often the path is moved around and becomes incorrect.

FILE_PATH_CHANGES

Whenever files are added, moved, or deleted, the MAINTAINERS file patterns can be out of sync or outdated.

So MAINTAINERS might need updating in these cases.

MEMSET

The memset use appears to be incorrect. This may be caused due to badly ordered parameters. Please recheck the usage.

NOT_UNIFIED_DIFF

The patch file does not appear to be in unified-diff format. Please regenerate the patch file before sending it to the maintainer.

PRINTF_OXDECIMAL

Prefixing 0x with decimal output is defective and should be corrected.

SPDX_LICENSE_TAG

The source file is missing or has an improper SPDX identifier tag. The Linux kernel requires the precise SPDX identifier in all source files, and it is thoroughly documented in the kernel docs.

See: <https://www.kernel.org/doc/html/latest/process/license-rules.html>

TYPO SPELLING

Some words may have been misspelled. Consider reviewing them.

COCCINELLE

Coccinelle is a tool for pattern matching and text transformation that has many uses in kernel development, including the application of complex, tree-wide patches and detection of problematic programming patterns.

3.1 Getting Coccinelle

The semantic patches included in the kernel use features and options which are provided by Coccinelle version 1.0.0-rc11 and above. Using earlier versions will fail as the option names used by the Coccinelle files and coccicheck have been updated.

Coccinelle is available through the package manager of many distributions, e.g. :

- Debian
- Fedora
- Ubuntu
- OpenSUSE
- Arch Linux
- NetBSD
- FreeBSD

Some distribution packages are obsolete and it is recommended to use the latest version released from the Coccinelle homepage at <http://coccinelle.lip6.fr/>

Or from Github at:

<https://github.com/coccinelle/coccinelle>

Once you have it, run the following commands:

```
./autogen  
./configure  
make
```

as a regular user, and install it with:

```
sudo make install
```

More detailed installation instructions to build from source can be found at:

<https://github.com/coccinelle/coccinelle/blob/master/install.txt>

3.2 Supplemental documentation

For supplemental documentation refer to the wiki:

<https://bottest.wiki.kernel.org/coccicheck>

The wiki documentation always refers to the linux-next version of the script.

For Semantic Patch Language(SmPL) grammar documentation refer to:

https://coccinelle.gitlabpages.inria.fr/website/docs/main_grammar.html

3.3 Using Coccinelle on the Linux kernel

A Coccinelle-specific target is defined in the top level Makefile. This target is named `coccicheck` and calls the `coccicheck` front-end in the `scripts` directory.

Four basic modes are defined: `patch`, `report`, `context`, and `org`. The mode to use is specified by setting the `MODE` variable with `MODE=<mode>`.

- `patch` proposes a fix, when possible.
- `report` generates a list in the following format: `file:line:column-column: message`
- `context` highlights lines of interest and their context in a diff-like style. Lines of interest are indicated with `-`.
- `org` generates a report in the Org mode format of Emacs.

Note that not all semantic patches implement all modes. For easy use of Coccinelle, the default mode is "report".

Two other modes provide some common combinations of these modes.

- `chain` tries the previous modes in the order above until one succeeds.
- `rep+ctxt` runs successively the `report` mode and the `context` mode. It should be used with the `C` option (described later) which checks the code on a file basis.

3.3.1 Examples

To make a report for every semantic patch, run the following command:

```
make coccicheck MODE=report
```

To produce patches, run:

```
make coccicheck MODE=patch
```

The coccicheck target applies every semantic patch available in the sub-directories of `scripts/coccinelle` to the entire Linux kernel.

For each semantic patch, a commit message is proposed. It gives a description of the problem being checked by the semantic patch, and includes a reference to Coccinelle.

As with any static code analyzer, Coccinelle produces false positives. Thus, reports must be carefully checked, and patches reviewed.

To enable verbose messages set the `V=` variable, for example:

```
make coccicheck MODE=report V=1
```

3.4 Coccinelle parallelization

By default, coccicheck tries to run as parallel as possible. To change the parallelism, set the `J=` variable. For example, to run across 4 CPUs:

```
make coccicheck MODE=report J=4
```

As of Coccinelle 1.0.2 Coccinelle uses Ocaml parmap for parallelization; if support for this is detected you will benefit from parmap parallelization.

When parmap is enabled coccicheck will enable dynamic load balancing by using `--chunksize 1` argument. This ensures we keep feeding threads with work one by one, so that we avoid the situation where most work gets done by only a few threads. With dynamic load balancing, if a thread finishes early we keep feeding it more work.

When parmap is enabled, if an error occurs in Coccinelle, this error value is propagated back, and the return value of the `make coccicheck` command captures this return value.

3.5 Using Coccinelle with a single semantic patch

The optional make variable `COCCI` can be used to check a single semantic patch. In that case, the variable must be initialized with the name of the semantic patch to apply.

For instance:

```
make coccicheck COCCI=<my_SP.cocci> MODE=patch
```

or:

```
make coccicheck COCCI=<my_SP.cocci> MODE=report
```

3.6 Controlling Which Files are Processed by Coccinelle

By default the entire kernel source tree is checked.

To apply Coccinelle to a specific directory, `M=` can be used. For example, to check `drivers/net/wireless/` one may write:

```
make coccicheck M=drivers/net/wireless/
```

To apply Coccinelle on a file basis, instead of a directory basis, the `C` variable is used by the makefile to select which files to work with. This variable can be used to run scripts for the entire kernel, a specific directory, or for a single file.

For example, to check `drivers/bluetooth/bfusb.c`, the value `1` is passed to the `C` variable to check files that make considers need to be compiled.:

```
make C=1 CHECK=scripts/coccicheck drivers/bluetooth/bfusb.o
```

The value `2` is passed to the `C` variable to check files regardless of whether they need to be compiled or not.:

```
make C=2 CHECK=scripts/coccicheck drivers/bluetooth/bfusb.o
```

In these modes, which work on a file basis, there is no information about semantic patches displayed, and no commit message proposed.

This runs every semantic patch in `scripts/coccinelle` by default. The `COCCI` variable may additionally be used to only apply a single semantic patch as shown in the previous section.

The "report" mode is the default. You can select another one with the `MODE` variable explained above.

3.7 Debugging Coccinelle SmPL patches

Using `coccicheck` is best as it provides in the `spatch` command line include options matching the options used when we compile the kernel. You can learn what these options are by using `V=1`; you could then manually run Coccinelle with debug options added.

Alternatively you can debug running Coccinelle against SmPL patches by asking for `stderr` to be redirected to `stderr`. By default `stderr` is redirected to `/dev/null`; if you'd like to capture `stderr` you can specify the `DEBUG_FILE="file.txt"` option to `coccicheck`. For instance:

```
rm -f cocci.err
make coccicheck COCCI=scripts/coccinelle/free/kfree.cocci MODE=report DEBUG_
→FILE=c cocci.err
cat cocci.err
```

You can use `SPFLAGS` to add debugging flags; for instance you may want to add both `--profile` `--show-trying` to `SPFLAGS` when debugging. For example you may want to use:

```
rm -f err.log
export COCCI=scripts/coccinelle/misc/irqf_oneshot.cocci
make coccicheck DEBUG_FILE="err.log" MODE=report SPFLAGS="--profile --show-
→trying" M=../drivers/mfd
```

err.log will now have the profiling information, while stdout will provide some progress information as Coccinelle moves forward with work.

NOTE:

DEBUG_FILE support is only supported when using coccinelle >= 1.0.2.

Currently, DEBUG_FILE support is only available to check folders, and not single files. This is because checking a single file requires spatch to be called twice leading to DEBUG_FILE being set both times to the same value, giving rise to an error.

3.8 .cocciconfig support

Coccinelle supports reading .cocciconfig for default Coccinelle options that should be used every time spatch is spawned. The order of precedence for variables for .cocciconfig is as follows:

- Your current user's home directory is processed first
- Your directory from which spatch is called is processed next
- The directory provided with the --dir option is processed last, if used

Since coccicheck runs through make, it naturally runs from the kernel proper dir; as such the second rule above would be implied for picking up a .cocciconfig when using make coccicheck.

make coccicheck also supports using M= targets. If you do not supply any M= target, it is assumed you want to target the entire kernel. The kernel coccicheck script has:

```
if [ "$KBUILD_EXTMOD" = "" ] ; then
    OPTIONS="--dir $srctree $COCCIINCLUDE"
else
    OPTIONS="--dir $KBUILD_EXTMOD $COCCIINCLUDE"
fi
```

KBUILD_EXTMOD is set when an explicit target with M= is used. For both cases the spatch --dir argument is used, as such third rule applies when whether M= is used or not, and when M= is used the target directory can have its own .cocciconfig file. When M= is not passed as an argument to coccicheck the target directory is the same as the directory from where spatch was called.

If not using the kernel's coccicheck target, keep the above precedence order logic of .cocciconfig reading. If using the kernel's coccicheck target, override any of the kernel's .coccicheck's settings using SPFLAGS.

We help Coccinelle when used against Linux with a set of sensible default options for Linux with our own Linux .cocciconfig. This hints to coccinelle that git can be used for git grep queries over coccigrep. A timeout of 200 seconds should suffice for now.

The options picked up by coccinelle when reading a .cocciconfig do not appear as arguments to spatch processes running on your system. To confirm what options will be used by Coccinelle run:

```
spatch --print-options-only
```

You can override with your own preferred index option by using SPFLAGS. Take note that when there are conflicting options Coccinelle takes precedence for the last options passed. Using .cocciconfig is possible to use idutils, however given the order of precedence followed by Coccinelle, since the kernel now carries its own .cocciconfig, you will need to use SPFLAGS to use idutils if desired. See below section "Additional flags" for more details on how to use idutils.

3.9 Additional flags

Additional flags can be passed to spatch through the SPFLAGS variable. This works as Coccinelle respects the last flags given to it when options are in conflict.

```
make SPFLAGS=--use-glimpse coccicheck
```

Coccinelle supports idutils as well but requires coccinelle $\geq 1.0.6$. When no ID file is specified coccinelle assumes your ID database file is in the file .id-utils.index on the top level of the kernel. Coccinelle carries a script scripts/idutils_index.sh which creates the database with:

```
mkid -i C --output .id-utils.index
```

If you have another database filename you can also just symlink with this name.

```
make SPFLAGS=--use-idutils coccicheck
```

Alternatively you can specify the database filename explicitly, for instance:

```
make SPFLAGS="--use-idutils /full-path/to/ID" coccicheck
```

See spatch --help to learn more about spatch options.

Note that the --use-glimpse and --use-idutils options require external tools for indexing the code. None of them is thus active by default. However, by indexing the code with one of these tools, and according to the cocci file used, spatch could proceed the entire code base more quickly.

3.10 SmPL patch specific options

SmPL patches can have their own requirements for options passed to Coccinelle. SmPL patch-specific options can be provided by providing them at the top of the SmPL patch, for instance:

```
// Options: --no-includes --include-headers
```

3.11 SmPL patch Coccinelle requirements

As Coccinelle features get added some more advanced SmPL patches may require newer versions of Coccinelle. If an SmPL patch requires a minimum version of Coccinelle, this can be specified as follows, as an example if requiring at least Coccinelle >= 1.0.5:

```
// Requires: 1.0.5
```

3.12 Proposing new semantic patches

New semantic patches can be proposed and submitted by kernel developers. For sake of clarity, they should be organized in the sub-directories of `scripts/coccinelle/`.

3.13 Detailed description of the report mode

`report` generates a list in the following format:

```
file:line:column-column: message
```

3.13.1 Example

Running:

```
make coccicheck MODE=report COCCI=scripts/coccinelle/api/err_cast.cocci
```

will execute the following part of the SmPL script:

```
<smpl>
@r depends on !context && !patch && (org || report)@
expression x;
position p;
@@

ERR_PTR@p(PTR_ERR(x))

@script:python depends on report@
p << r.p;
x << r.x;
@@

msg="ERR_CAST can be used with %s" % (x)
coccilib.report.print_report(p[0], msg)
</smpl>
```

This SmPL excerpt generates entries on the standard output, as illustrated below:

```
/home/user/linux/crypto/ctr.c:188:9-16: ERR_CAST can be used with alg  
/home/user/linux/crypto/authenc.c:619:9-16: ERR_CAST can be used with auth  
/home/user/linux/crypto/xts.c:227:9-16: ERR_CAST can be used with alg
```

3.14 Detailed description of the patch mode

When the patch mode is available, it proposes a fix for each problem identified.

3.14.1 Example

Running:

```
make coccicheck MODE=patch COCCI=scripts/coccinelle/api/err_cast.cocci
```

will execute the following part of the SmPL script:

```
<smpl>  
@ depends on !context && patch && !org && !report @  
expression x;  
@@  
  
- ERR_PTR(PTR_ERR(x))  
+ ERR_CAST(x)  
</smpl>
```

This SmPL excerpt generates patch hunks on the standard output, as illustrated below:

```
diff -u -p a/crypto/ctr.c b/crypto/ctr.c  
--- a/crypto/ctr.c 2010-05-26 10:49:38.000000000 +0200  
+++ b/crypto/ctr.c 2010-06-03 23:44:49.000000000 +0200  
@@ -185,7 +185,7 @@ static struct crypto_instance *crypto_ct  
     alg = crypto_attr_alg(tb[1], CRYPTO_ALG_TYPE_CIPHER,  
                           CRYPTO_ALG_TYPE_MASK);  
     if (IS_ERR(alg))  
-         return ERR_PTR(PTR_ERR(alg));  
+         return ERR_CAST(alg);  
  
     /* Block size must be >= 4 bytes. */  
     err = -EINVAL;
```

3.15 Detailed description of the context mode

`context` highlights lines of interest and their context in a diff-like style.

NOTE: The diff-like output generated is NOT an applicable patch. The intent of the `context` mode is to highlight the important lines (annotated with minus, `-`) and gives some surrounding context lines around. This output can be used with the `diff` mode of Emacs to review the code.

3.15.1 Example

Running:

```
make coccicheck MODE=context COCCI=scripts/coccinelle/api/err_cast.cocci
```

will execute the following part of the SmPL script:

```
<smpl>
@ depends on context && !patch && !org && !report@
expression x;
@@

* ERR_PTR(PTR_ERR(x))
</smpl>
```

This SmPL excerpt generates diff hunks on the standard output, as illustrated below:

```
diff -u -p /home/user/linux/crypto/ctr.c /tmp/nothing
--- /home/user/linux/crypto/ctr.c    2010-05-26 10:49:38.000000000 +0200
+++ /tmp/nothing
@@ -185,7 +185,6 @@ static struct crypto_instance *crypto_ct
     alg = crypto_attr_alg(tb[1], CRYPTO_ALG_TYPE_CIPHER,
                           CRYPTO_ALG_TYPE_MASK);
     if (IS_ERR(alg))
-        return ERR_PTR(PTR_ERR(alg));
 
     /* Block size must be >= 4 bytes. */
     err = -EINVAL;
```

3.16 Detailed description of the org mode

`org` generates a report in the Org mode format of Emacs.

3.16.1 Example

Running:

```
make coccicheck MODE=org COCCI=scripts/coccinelle/api/err_cast.cocci
```

will execute the following part of the SmPL script:

```
<smpl>
@r depends on !context && !patch && (org || report)@
expression x;
position p;
@@

ERR_PTR@p(PTR_ERR(x))

@script:python depends on org@
p << r.p;
x << r.x;
@@

msg="ERR_CAST can be used with %s" % (x)
msg_safe=msg.replace("[","@(").replace("]",")")
coccilib.org.print_todo(p[0], msg_safe)
</smpl>
```

This SmPL excerpt generates Org entries on the standard output, as illustrated below:

```
* TODO [[view:/home/user/linux/crypto/ctr.c::face=ovl-
  ↳ face1::linb=188::colb=9::cole=16][ERR_CAST can be used with alg]]
* TODO [[view:/home/user/linux/crypto/authenc.c::face=ovl-
  ↳ face1::linb=619::colb=9::cole=16][ERR_CAST can be used with auth]]
* TODO [[view:/home/user/linux/crypto/xts.c::face=ovl-
  ↳ face1::linb=227::colb=9::cole=16][ERR_CAST can be used with alg]]
```

SPARSE

Sparse is a semantic checker for C programs; it can be used to find a number of potential problems with kernel code. See <https://lwn.net/Articles/689907/> for an overview of sparse; this document contains some kernel-specific sparse information. More information on sparse, mainly about its internals, can be found in its official pages at <https://sparse.docs.kernel.org>.

4.1 Using sparse for typechecking

"__bitwise" is a type attribute, so you have to do something like this:

```
typedef int __bitwise pm_request_t;

enum pm_request {
    PM_SUSPEND = (__force pm_request_t) 1,
    PM_RESUME = (__force pm_request_t) 2
};
```

which makes PM_SUSPEND and PM_RESUME "bitwise" integers (the "__force" is there because sparse will complain about casting to/from a bitwise type, but in this case we really _do_ want to force the conversion). And because the enum values are all the same type, now "enum pm_request" will be that type too.

And with gcc, all the "__bitwise"/"__force stuff" goes away, and it all ends up looking just like integers to gcc.

Quite frankly, you don't need the enum there. The above all really just boils down to one special "int __bitwise" type.

So the simpler way is to just do:

```
typedef int __bitwise pm_request_t;

#define PM_SUSPEND ((__force pm_request_t) 1)
#define PM_RESUME ((__force pm_request_t) 2)
```

and you now have all the infrastructure needed for strict typechecking.

One small note: the constant integer "0" is special. You can use a constant zero as a bitwise integer type without sparse ever complaining. This is because "bitwise" (as the name implies) was designed for making sure that bitwise types don't get mixed up (little-endian vs big-endian vs cpu-endian vs whatever), and there the constant "0" really _is_ special.

4.2 Using sparse for lock checking

The following macros are undefined for gcc and defined during a sparse run to use the "context" tracking feature of sparse, applied to locking. These annotations tell sparse when a lock is held, with regard to the annotated function's entry and exit.

`_must_hold` - The specified lock is held on function entry and exit.

`_acquires` - The specified lock is held on function exit, but not entry.

`_releases` - The specified lock is held on function entry, but not exit.

If the function enters and exits without the lock held, acquiring and releasing the lock inside the function in a balanced way, no annotation is needed. The three annotations above are for cases where sparse would otherwise report a context imbalance.

4.3 Getting sparse

You can get tarballs of the latest released versions from: <https://www.kernel.org/pub/software-devel/sparse/dist/>

Alternatively, you can get snapshots of the latest development version of sparse using git to clone:

```
git://git.kernel.org/pub/scm-devel/sparse/sparse.git
```

Once you have it, just do:

```
make  
make install
```

as a regular user, and it will install sparse in your `~/bin` directory.

4.4 Using sparse

Do a kernel make with "make C=1" to run sparse on all the C files that get recompiled, or use "make C=2" to run sparse on the files whether they need to be recompiled or not. The latter is a fast way to check the whole tree if you have already built it.

The optional make variable CF can be used to pass arguments to sparse. The build system passes `-Wbitwise` to sparse automatically.

Note that sparse defines the `_CHECKER_` preprocessor symbol.

KCOV: CODE COVERAGE FOR FUZZING

KCOV collects and exposes kernel code coverage information in a form suitable for coverage-guided fuzzing. Coverage data of a running kernel is exported via the `kcov` debugfs file. Coverage collection is enabled on a task basis, and thus KCOV can capture precise coverage of a single system call.

Note that KCOV does not aim to collect as much coverage as possible. It aims to collect more or less stable coverage that is a function of syscall inputs. To achieve this goal, it does not collect coverage in soft/hard interrupts (unless remote coverage collection is enabled, see below) and from some inherently non-deterministic parts of the kernel (e.g. scheduler, locking).

Besides collecting code coverage, KCOV can also collect comparison operands. See the "Comparison operands collection" section for details.

Besides collecting coverage data from syscall handlers, KCOV can also collect coverage for annotated parts of the kernel executing in background kernel tasks or soft interrupts. See the "Remote coverage collection" section for details.

5.1 Prerequisites

KCOV relies on compiler instrumentation and requires GCC 6.1.0 or later or any Clang version supported by the kernel.

Collecting comparison operands is supported with GCC 8+ or with Clang.

To enable KCOV, configure the kernel with:

```
CONFIG_KCOV=y
```

To enable comparison operands collection, set:

```
CONFIG_KCOV_ENABLE_COMPARISONS=y
```

Coverage data only becomes accessible once debugfs has been mounted:

```
mount -t debugfs none /sys/kernel/debug
```

5.2 Coverage collection

The following program demonstrates how to use KCOV to collect coverage for a single syscall from within a test program:

```
#include <stdio.h>
#include <stddef.h>
#include <stdint.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/ioctl.h>
#include <sys/mman.h>
#include <unistd.h>
#include <fcntl.h>
#include <linux/types.h>

#define KCOV_INIT_TRACE           _IOR('c', 1, unsigned long)
#define KCOV_ENABLE                _IO('c', 100)
#define KCOV_DISABLE               _IO('c', 101)
#define COVER_SIZE          (64<<10)

#define KCOV_TRACE_PC   0
#define KCOV_TRACE_CMP  1

int main(int argc, char **argv)
{
    int fd;
    unsigned long *cover, n, i;

    /* A single fd descriptor allows coverage collection on a single
     * thread.
     */
    fd = open("/sys/kernel/debug/kcov", O_RDWR);
    if (fd == -1)
        perror("open"), exit(1);
    /* Setup trace mode and trace size. */
    if (ioctl(fd, KCOV_INIT_TRACE, COVER_SIZE))
        perror("ioctl"), exit(1);
    /* Mmap buffer shared between kernel- and user-space. */
    cover = (unsigned long*)mmap(NULL, COVER_SIZE * sizeof(unsigned long),
                                PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    if ((void*)cover == MAP_FAILED)
        perror("mmap"), exit(1);
    /* Enable coverage collection on the current thread. */
    if (ioctl(fd, KCOV_ENABLE, KCOV_TRACE_PC))
        perror("ioctl"), exit(1);
    /* Reset coverage from the tail of the ioctl() call. */
    __atomic_store_n(&cover[0], 0, __ATOMIC_RELAXED);
    /* Call the target syscall call. */
    read(-1, NULL, 0);
```

```

/* Read number of PCs collected. */
n = __atomic_load_n(&cover[0], __ATOMIC_RELAXED);
for (i = 0; i < n; i++)
    printf("0x%lx\n", cover[i + 1]);
/* Disable coverage collection for the current thread. After this call
 * coverage can be enabled for a different thread.
 */
if (ioctl(fd, KCOV_DISABLE, 0))
    perror("ioctl"), exit(1);
/* Free resources. */
if (munmap(cover, COVER_SIZE * sizeof(unsigned long)))
    perror("munmap"), exit(1);
if (close(fd))
    perror("close"), exit(1);
return 0;
}

```

After piping through `addr2line` the output of the program looks as follows:

```

SyS_read
fs/read_write.c:562
__fdget_pos
fs/file.c:774
__fget_light
fs/file.c:746
__fget_light
fs/file.c:750
__fget_light
fs/file.c:760
__fdget_pos
fs/file.c:784
SyS_read
fs/read_write.c:562

```

If a program needs to collect coverage from several threads (independently), it needs to open `/sys/kernel/debug/kcov` in each thread separately.

The interface is fine-grained to allow efficient forking of test processes. That is, a parent process opens `/sys/kernel/debug/kcov`, enables trace mode, mmap's coverage buffer, and then forks child processes in a loop. The child processes only need to enable coverage (it gets disabled automatically when a thread exits).

5.3 Comparison operands collection

Comparison operands collection is similar to coverage collection:

```
/* Same includes and defines as above. */

/* Number of 64-bit words per record. */
#define KCOV_WORDS_PER_CMP 4

/*
 * The format for the types of collected comparisons.
 *
 * Bit 0 shows whether one of the arguments is a compile-time constant.
 * Bits 1 & 2 contain log2 of the argument size, up to 8 bytes.
 */

#define KCOV_CMP_CONST          (1 << 0)
#define KCOV_CMP_SIZE(n)         (((n) << 1)
#define KCOV_CMP_MASK           KCOV_CMP_SIZE(3)

int main(int argc, char **argv)
{
    int fd;
    uint64_t *cover, type, arg1, arg2, is_const, size;
    unsigned long n, i;

    fd = open("/sys/kernel/debug/kcov", O_RDWR);
    if (fd == -1)
        perror("open"), exit(1);
    if (ioctl(fd, KCOV_INIT_TRACE, COVER_SIZE))
        perror("ioctl"), exit(1);
    /*
     * Note that the buffer pointer is of type uint64_t*, because all
     * the comparison operands are promoted to uint64_t.
     */
    cover = (uint64_t *)mmap(NULL, COVER_SIZE * sizeof(unsigned long),
                           PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    if ((void*)cover == MAP_FAILED)
        perror("mmap"), exit(1);
    /* Note KCOV_TRACE_CMP instead of KCOV_TRACE_PC. */
    if (ioctl(fd, KCOV_ENABLE, KCOV_TRACE_CMP))
        perror("ioctl"), exit(1);
    __atomic_store_n(&cover[0], 0, __ATOMIC_RELAXED);
    read(-1, NULL, 0);
    /* Read number of comparisons collected. */
    n = __atomic_load_n(&cover[0], __ATOMIC_RELAXED);
    for (i = 0; i < n; i++) {
        uint64_t ip;

        type = cover[i * KCOV_WORDS_PER_CMP + 1];
        /* arg1 and arg2 - operands of the comparison. */
    }
}
```

```

        arg1 = cover[i * KCOV_WORDS_PER_CMP + 2];
        arg2 = cover[i * KCOV_WORDS_PER_CMP + 3];
        /* ip - caller address. */
        ip = cover[i * KCOV_WORDS_PER_CMP + 4];
        /* size of the operands. */
        size = 1 << ((type & KCOV_CMP_MASK) >> 1);
        /* is_const - true if either operand is a compile-time constant.*/
        is_const = type & KCOV_CMP_CONST;
        printf("ip: 0x%lx type: 0x%lx, arg1: 0x%lx, arg2: 0x%lx, "
               "size: %lu, %s\n",
               ip, type, arg1, arg2, size,
               is_const ? "const" : "non-const");
    }
    if (ioctl(fd, KCOV_DISABLE, 0))
        perror("ioctl"), exit(1);
    /* Free resources. */
    if (munmap(cover, COVER_SIZE * sizeof(unsigned long)))
        perror("munmap"), exit(1);
    if (close(fd))
        perror("close"), exit(1);
    return 0;
}

```

Note that the KCOV modes (collection of code coverage or comparison operands) are mutually exclusive.

5.4 Remote coverage collection

Besides collecting coverage data from handlers of syscalls issued from a userspace process, KCOV can also collect coverage for parts of the kernel executing in other contexts - so-called "remote" coverage.

Using KCOV to collect remote coverage requires:

1. Modifying kernel code to annotate the code section from where coverage should be collected with `kcov_remote_start` and `kcov_remote_stop`.
2. Using `KCOV_REMOTE_ENABLE` instead of `KCOV_ENABLE` in the userspace process that collects coverage.

Both `kcov_remote_start` and `kcov_remote_stop` annotations and the `KCOV_REMOTE_ENABLE` ioctl accept handles that identify particular coverage collection sections. The way a handle is used depends on the context where the matching code section executes.

KCOV supports collecting remote coverage from the following contexts:

1. Global kernel background tasks. These are the tasks that are spawned during kernel boot in a limited number of instances (e.g. one `USB_hub_event` worker is spawned per one `USB_HCD`).
2. Local kernel background tasks. These are spawned when a userspace process interacts with some kernel interface and are usually killed when the process exits (e.g. `vhost` workers).

3. Soft interrupts.

For #1 and #3, a unique global handle must be chosen and passed to the corresponding `kcov_remote_start` call. Then a userspace process must pass this handle to `KCOV_REMOTE_ENABLE` in the `handles` array field of the `kcov_remote_arg` struct. This will attach the used KCOV device to the code section referenced by this handle. Multiple global handles identifying different code sections can be passed at once.

For #2, the userspace process instead must pass a non-zero handle through the `common_handle` field of the `kcov_remote_arg` struct. This common handle gets saved to the `kcov_handle` field in the current `task_struct` and needs to be passed to the newly spawned local tasks via custom kernel code modifications. Those tasks should in turn use the passed handle in their `kcov_remote_start` and `kcov_remote_stop` annotations.

KCOV follows a predefined format for both global and common handles. Each handle is a `u64` integer. Currently, only the one top and the lower 4 bytes are used. Bytes 4-7 are reserved and must be zero.

For global handles, the top byte of the handle denotes the id of a subsystem this handle belongs to. For example, KCOV uses 1 as the USB subsystem id. The lower 4 bytes of a global handle denote the id of a task instance within that subsystem. For example, each `hub_event` worker uses the USB bus number as the task instance id.

For common handles, a reserved value 0 is used as a subsystem id, as such handles don't belong to a particular subsystem. The lower 4 bytes of a common handle identify a collective instance of all local tasks spawned by the userspace process that passed a common handle to `KCOV_REMOTE_ENABLE`.

In practice, any value can be used for common handle instance id if coverage is only collected from a single userspace process on the system. However, if common handles are used by multiple processes, unique instance ids must be used for each process. One option is to use the process id as the common handle instance id.

The following program demonstrates using KCOV to collect coverage from both local tasks spawned by the process and the global task that handles USB bus #1:

```
/* Same includes and defines as above. */

struct kcov_remote_arg {
    __u32          trace_mode;
    __u32          area_size;
    __u32          num_handles;
    __aligned_u64  common_handle;
    __aligned_u64  handles[0];
};

#define KCOV_INIT_TRACE           _IOR('c', 1, unsigned long)
#define KCOV_DISABLE              _IO('c', 101)
#define KCOV_REMOTE_ENABLE         _IOW('c', 102, struct kcov_remote_arg)

#define COVER_SIZE   (64 << 10)
#define KCOV_TRACE_PC      0
#define KCOV_SUBSYSTEM_COMMON (0x00ull << 56)
```

```

#define KCOV_SUBSYSTEM_USB  (0x01ull << 56)
#define KCOV_SUBSYSTEM_MASK (0xfffull << 56)
#define KCOV_INSTANCE_MASK  (0xfffffffffull)

static inline __u64 kcov_remote_handle(__u64 subsys, __u64 inst)
{
    if (subsys & ~KCOV_SUBSYSTEM_MASK || inst & ~KCOV_INSTANCE_MASK)
        return 0;
    return subsys | inst;
}

#define KCOV_COMMON_ID      0x42
#define KCOV_USB_BUS_NUM    1

int main(int argc, char **argv)
{
    int fd;
    unsigned long *cover, n, i;
    struct kcov_remote_arg *arg;

    fd = open("/sys/kernel/debug/kcov", O_RDWR);
    if (fd == -1)
        perror("open"), exit(1);
    if (ioctl(fd, KCOV_INIT_TRACE, COVER_SIZE))
        perror("ioctl"), exit(1);
    cover = (unsigned long*)mmap(NULL, COVER_SIZE * sizeof(unsigned long),
                                 PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    if ((void*)cover == MAP_FAILED)
        perror("mmap"), exit(1);

    /* Enable coverage collection via common handle and from USB bus #1. */
    arg = calloc(1, sizeof(*arg) + sizeof(uint64_t));
    if (!arg)
        perror("calloc"), exit(1);
    arg->trace_mode = KCOV_TRACE_PC;
    arg->area_size = COVER_SIZE;
    arg->num_handles = 1;
    arg->common_handle = kcov_remote_handle(KCOV_SUBSYSTEM_COMMON,
                                              KCOV_COMMON_ID);
    arg->handles[0] = kcov_remote_handle(KCOV_SUBSYSTEM_USB,
                                          KCOV_USB_BUS_NUM);
    if (ioctl(fd, KCOV_REMOTE_ENABLE, arg))
        perror("ioctl"), free(arg), exit(1);
    free(arg);

    /*
     * Here the user needs to trigger execution of a kernel code section
     * that is either annotated with the common handle, or to trigger some
     * activity on USB bus #1.
     */
}

```

```
 */
sleep(2);

n = __atomic_load_n(&cover[0], __ATOMIC_RELAXED);
for (i = 0; i < n; i++)
    printf("0x%lx\n", cover[i + 1]);
if (ioctl(fd, KCOV_DISABLE, 0))
    perror("ioctl"), exit(1);
if (munmap(cover, COVER_SIZE * sizeof(unsigned long)))
    perror("munmap"), exit(1);
if (close(fd))
    perror("close"), exit(1);
return 0;
}
```

USING GCOV WITH THE LINUX KERNEL

gcov profiling kernel support enables the use of GCC's coverage testing tool `gcov` with the Linux kernel. Coverage data of a running kernel is exported in gcov-compatible format via the "gcov" debugfs directory. To get coverage data for a specific file, change to the kernel build directory and use gcov with the -o option as follows (requires root):

```
# cd /tmp/linux-out
# gcov -o /sys/kernel/debug/gcov/tmp/linux-out/kernel spinlock.c
```

This will create source code files annotated with execution counts in the current directory. In addition, graphical gcov front-ends such as `lcov` can be used to automate the process of collecting data for the entire kernel and provide coverage overviews in HTML format.

Possible uses:

- debugging (has this line been reached at all?)
- test improvement (how do I change my test to cover these lines?)
- minimizing kernel configurations (do I need this option if the associated code is never run?)

6.1 Preparation

Configure the kernel with:

```
CONFIG_DEBUG_FS=y
CONFIG_GCOV_KERNEL=y
```

and to get coverage data for the entire kernel:

```
CONFIG_GCOV_PROFILE_ALL=y
```

Note that kernels compiled with profiling flags will be significantly larger and run slower. Also `CONFIG_GCOV_PROFILE_ALL` may not be supported on all architectures.

Profiling data will only become accessible once debugfs has been mounted:

```
mount -t debugfs none /sys/kernel/debug
```

6.2 Customization

To enable profiling for specific files or directories, add a line similar to the following to the respective kernel Makefile:

- For a single file (e.g. main.o):

```
GCov_PROFILE_main.o := y
```

- For all files in one directory:

```
GCov_PROFILE := y
```

To exclude files from being profiled even when CONFIG_GCOV_PROFILE_ALL is specified, use:

```
GCov_PROFILE_main.o := n
```

and:

```
GCov_PROFILE := n
```

Only files which are linked to the main kernel image or are compiled as kernel modules are supported by this mechanism.

6.3 Files

The gcov kernel support creates the following files in debugfs:

/sys/kernel/debug/gcov

Parent directory for all gcov-related files.

/sys/kernel/debug/gcov/reset

Global reset file: resets all coverage data to zero when written to.

/sys/kernel/debug/gcov/path/to/compile/dir/file.gcda

The actual gcov data file as understood by the gcov tool. Resets file coverage data to zero when written to.

/sys/kernel/debug/gcov/path/to/compile/dir/file.gcno

Symbolic link to a static data file required by the gcov tool. This file is generated by gcc when compiling with option `-ftest-coverage`.

6.4 Modules

Kernel modules may contain cleanup code which is only run during module unload time. The gcov mechanism provides a means to collect coverage data for such code by keeping a copy of the data associated with the unloaded module. This data remains available through debugfs. Once the module is loaded again, the associated coverage counters are initialized with the data from its previous instantiation.

This behavior can be deactivated by specifying the `gcov_persist` kernel parameter:

```
gcov_persist=0
```

At run-time, a user can also choose to discard data for an unloaded module by writing to its data file or the global reset file.

6.5 Separated build and test machines

The gcov kernel profiling infrastructure is designed to work out-of-the box for setups where kernels are built and run on the same machine. In cases where the kernel runs on a separate machine, special preparations must be made, depending on where the gcov tool is used:

- gcov is run on the TEST machine

The gcov tool version on the test machine must be compatible with the gcc version used for kernel build. Also the following files need to be copied from build to test machine:

from the source tree:

- all C source files + headers

from the build tree:

- all C source files + headers
- all .gcda and .gcno files
- all links to directories

It is important to note that these files need to be placed into the exact same file system location on the test machine as on the build machine. If any of the path components is symbolic link, the actual directory needs to be used instead (due to make's CURDIR handling).

- gcov is run on the BUILD machine

The following files need to be copied after each test case from test to build machine:

from the gcov directory in sysfs:

- all .gcda files
- all links to .gcno files

These files can be copied to any location on the build machine. gcov must then be called with the -o option pointing to that directory.

Example directory setup on the build machine:

```
/tmp/linux:    kernel source tree
/tmp/out:      kernel build directory as specified by make O=
/tmp/coverage: location of the files copied from the test machine

[user@build] cd /tmp/out
[user@build] gcov -o /tmp/coverage/tmp/out/init main.c
```

6.6 Note on compilers

GCC and LLVM gcov tools are not necessarily compatible. Use `gcov` to work with GCC-generated `.gcno` and `.gcda` files, and use `llvm-cov` for Clang.

Build differences between GCC and Clang gcov are handled by Kconfig. It automatically selects the appropriate gcov format depending on the detected toolchain.

6.7 Troubleshooting

Problem

Compilation aborts during linker step.

Cause

Profiling flags are specified for source files which are not linked to the main kernel or which are linked by a custom linker procedure.

Solution

Exclude affected source files from profiling by specifying `GCNV_PROFILE := n` or `GCNV_PROFILE_basename.o := n` in the corresponding Makefile.

Problem

Files copied from sysfs appear empty or incomplete.

Cause

Due to the way `seq_file` works, some tools such as `cp` or `tar` may not correctly copy files from sysfs.

Solution

Use `cat` to read `.gcda` files and `cp -d` to copy links. Alternatively use the mechanism shown in Appendix B.

6.8 Appendix A: gather_on_build.sh

Sample script to gather coverage meta files on the build machine (see *Separated build and test machines a.*):

```
#!/bin/bash

KSRC=$1
KOBJ=$2
DEST=$3

if [ -z "$KSRC" ] || [ -z "$KOBJ" ] || [ -z "$DEST" ]; then
    echo "Usage: $0 <ksrc directory> <kobj directory> <output.tar.gz>" >&2
    exit 1
fi

KSRC=$(cd $KSRC; printf "all:\n\t@echo \$\{CURDIR\}\n" | make -f -)
KOBJ=$(cd $KOBJ; printf "all:\n\t@echo \$\{CURDIR\}\n" | make -f -)
```

```

find $KSRC $KOBJ \(
    -name '*.[ch]' -o -name '*.[ch]' -o -type l \) -a \
    -perm /u+r,g+r | tar cfz $DEST -P -T -
if [ $? -eq 0 ] ; then
    echo "$DEST successfully created, copy to test system and unpack with:"
    echo "  tar xfz $DEST -P"
else
    echo "Could not create file $DEST"
fi

```

6.9 Appendix B: gather_on_test.sh

Sample script to gather coverage data files on the test machine (see *Separated build and test machines b.*):

```

#!/bin/bash -e

DEST=$1
GCDA=/sys/kernel/debug/gcov

if [ -z "$DEST" ] ; then
    echo "Usage: $0 <output.tar.gz>" >&2
    exit 1
fi

TEMPDIR=$(mktemp -d)
echo Collecting data...
find $GCDA -type d -exec mkdir -p $TEMPDIR/\{\} \;
find $GCDA -name '*.[ch]' -exec sh -c 'cat < $0 > '$TEMPDIR'/$0' {} \;
find $GCDA -name '*.[ch]' -exec sh -c 'cp -d $0 '$TEMPDIR'/$0' {} \;
tar czf $DEST -C $TEMPDIR sys
rm -rf $TEMPDIR

echo "$DEST successfully created, copy to build system and unpack with:"
echo "  tar xfz $DEST"

```


KERNEL ADDRESS SANITIZER (KASAN)

7.1 Overview

Kernel Address Sanitizer (KASAN) is a dynamic memory safety error detector designed to find out-of-bounds and use-after-free bugs.

KASAN has three modes:

1. Generic KASAN
2. Software Tag-Based KASAN
3. Hardware Tag-Based KASAN

Generic KASAN, enabled with `CONFIG_KASAN_GENERIC`, is the mode intended for debugging, similar to userspace ASan. This mode is supported on many CPU architectures, but it has significant performance and memory overheads.

Software Tag-Based KASAN or SW_TAGS KASAN, enabled with `CONFIG_KASAN_SW_TAGS`, can be used for both debugging and dogfood testing, similar to userspace HWASan. This mode is only supported for arm64, but its moderate memory overhead allows using it for testing on memory-restricted devices with real workloads.

Hardware Tag-Based KASAN or HW_TAGS KASAN, enabled with `CONFIG_KASAN_HW_TAGS`, is the mode intended to be used as an in-field memory bug detector or as a security mitigation. This mode only works on arm64 CPUs that support MTE (Memory Tagging Extension), but it has low memory and performance overheads and thus can be used in production.

For details about the memory and performance impact of each KASAN mode, see the descriptions of the corresponding Kconfig options.

The Generic and the Software Tag-Based modes are commonly referred to as the software modes. The Software Tag-Based and the Hardware Tag-Based modes are referred to as the tag-based modes.

7.2 Support

7.2.1 Architectures

Generic KASAN is supported on x86_64, arm, arm64, powerpc, riscv, s390, xtensa, and loongarch, and the tag-based KASAN modes are supported only on arm64.

7.2.2 Compilers

Software KASAN modes use compile-time instrumentation to insert validity checks before every memory access and thus require a compiler version that provides support for that. The Hardware Tag-Based mode relies on hardware to perform these checks but still requires a compiler version that supports the memory tagging instructions.

Generic KASAN requires GCC version 8.3.0 or later or any Clang version supported by the kernel.

Software Tag-Based KASAN requires GCC 11+ or any Clang version supported by the kernel.

Hardware Tag-Based KASAN requires GCC 10+ or Clang 12+.

7.2.3 Memory types

Generic KASAN supports finding bugs in all of slab, page_alloc, vmap, vmalloc, stack, and global memory.

Software Tag-Based KASAN supports slab, page_alloc, vmalloc, and stack memory.

Hardware Tag-Based KASAN supports slab, page_alloc, and non-executable vmalloc memory.

For slab, both software KASAN modes support SLUB and SLAB allocators, while Hardware Tag-Based KASAN only supports SLUB.

7.3 Usage

To enable KASAN, configure the kernel with:

```
CONFIG_KASAN=y
```

and choose between CONFIG_KASAN_GENERIC (to enable Generic KASAN), CONFIG_KASAN_SW_TAGS (to enable Software Tag-Based KASAN), and CONFIG_KASAN_HW_TAGS (to enable Hardware Tag-Based KASAN).

For the software modes, also choose between CONFIG_KASAN_OUTLINE and CONFIG_KASAN_INLINE. Outline and inline are compiler instrumentation types. The former produces a smaller binary while the latter is up to 2 times faster.

To include alloc and free stack traces of affected slab objects into reports, enable CONFIG_STACKTRACE. To include alloc and free stack traces of affected physical pages, enable CONFIG_PAGE_OWNER and boot with page_owner=on.

7.3.1 Boot parameters

KASAN is affected by the generic `panic_on_warn` command line parameter. When it is enabled, KASAN panics the kernel after printing a bug report.

By default, KASAN prints a bug report only for the first invalid memory access. With `kasan_multi_shot`, KASAN prints a report on every invalid access. This effectively disables `panic_on_warn` for KASAN reports.

Alternatively, independent of `panic_on_warn`, the `kasan.fault=` boot parameter can be used to control panic and reporting behaviour:

- `kasan.fault=report`, `=panic`, or `=panic_on_write` controls whether to only print a KASAN report, panic the kernel, or panic the kernel on invalid writes only (default: `report`). The panic happens even if `kasan_multi_shot` is enabled. Note that when using asynchronous mode of Hardware Tag-Based KASAN, `kasan.fault=panic_on_write` always panics on asynchronously checked accesses (including reads).

Software and Hardware Tag-Based KASAN modes (see the section about various modes below) support altering stack trace collection behavior:

- `kasan.stacktrace=off` or `=on` disables or enables alloc and free stack traces collection (default: `on`).
- `kasan.stack_ring_size=<number of entries>` specifies the number of entries in the stack ring (default: 32768).

Hardware Tag-Based KASAN mode is intended for use in production as a security mitigation. Therefore, it supports additional boot parameters that allow disabling KASAN altogether or controlling its features:

- `kasan=off` or `=on` controls whether KASAN is enabled (default: `on`).
- `kasan.mode=sync`, `=async` or `=asymm` controls whether KASAN is configured in synchronous, asynchronous or asymmetric mode of execution (default: `sync`). Synchronous mode: a bad access is detected immediately when a tag check fault occurs. Asynchronous mode: a bad access detection is delayed. When a tag check fault occurs, the information is stored in hardware (in the TFSR_EL1 register for arm64). The kernel periodically checks the hardware and only reports tag faults during these checks. Asymmetric mode: a bad access is detected synchronously on reads and asynchronously on writes.
- `kasan.vmalloc=off` or `=on` disables or enables tagging of vmalloc allocations (default: `on`).
- `kasan.page_alloc.sample=<sampling interval>` makes KASAN tag only every Nth page_alloc allocation with the order equal or greater than `kasan.page_alloc.sample.order`, where N is the value of the `sample` parameter (default: 1, or tag every such allocation). This parameter is intended to mitigate the performance overhead introduced by KASAN. Note that enabling this parameter makes Hardware Tag-Based KASAN skip checks of allocations chosen by sampling and thus miss bad accesses to these allocations. Use the default value for accurate bug detection.
- `kasan.page_alloc.sample.order=<minimum page order>` specifies the minimum order of allocations that are affected by sampling (default: 3). Only applies when `kasan.page_alloc.sample` is set to a value greater than 1. This parameter is intended to allow sampling only large page_alloc allocations, which is the biggest source of the performance overhead.

7.3.2 Error reports

A typical KASAN report looks like this:

```
=====
BUG: KASAN: slab-out-of-bounds in kmalloc_oob_right+0xa8/0xbc [test_kasan]
Write of size 1 at addr ffff8801f44ec37b by task insmod/2760

CPU: 1 PID: 2760 Comm: insmod Not tainted 4.19.0-rc3+ #698
Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS 1.10.2-1 04/01/2014
Call Trace:
dump_stack+0x94/0xd8
print_address_description+0x73/0x280
kasan_report+0x144/0x187
__asan_report_store1_noabort+0x17/0x20
kmalloc_oob_right+0xa8/0xbc [test_kasan]
kmalloc_tests_init+0x16/0x700 [test_kasan]
do_one_initcall+0xa5/0x3ae
do_init_module+0x1b6/0x547
load_module+0x75df/0x8070
__do_sys_init_module+0x1c6/0x200
__x64_sys_init_module+0x6e/0xb0
do_syscall_64+0x9f/0x2c0
entry_SYSCALL_64_after_hwframe+0x44/0xa9
RIP: 0033:0x7f96443109da
RSP: 002b:00007ffcf0b51b08 EFLAGS: 00000202 ORIG_RAX: 00000000000000af
RAX: ffffffff7fffffd da RBX: 000055dc3ee521a0 RCX: 00007f96443109da
RDX: 00007f96445cff88 RSI: 0000000000057a50 RDI: 00007f9644992000
RBP: 000055dc3ee510b0 R08: 0000000000000003 R09: 0000000000000000
R10: 00007f964430cd0a R11: 0000000000000202 R12: 00007f96445cff88
R13: 000055dc3ee51090 R14: 0000000000000000 R15: 0000000000000000

Allocated by task 2760:
save_stack+0x43/0xd0
kasan_kmalloc+0xa7/0xd0
kmem_cache_alloc_trace+0xe1/0x1b0
kmalloc_oob_right+0x56/0xbc [test_kasan]
kmalloc_tests_init+0x16/0x700 [test_kasan]
do_one_initcall+0xa5/0x3ae
do_init_module+0x1b6/0x547
load_module+0x75df/0x8070
__do_sys_init_module+0x1c6/0x200
__x64_sys_init_module+0x6e/0xb0
do_syscall_64+0x9f/0x2c0
entry_SYSCALL_64_after_hwframe+0x44/0xa9

Freed by task 815:
save_stack+0x43/0xd0
__kasan_slab_free+0x135/0x190
kasan_slab_free+0xe/0x10
kfree+0x93/0x1a0
```

```
umh_complete+0x6a/0xa0
call_usermodehelper_exec_async+0x4c3/0x640
ret_from_fork+0x35/0x40
```

The buggy address belongs to the object at ffff8801f44ec300
which belongs to the cache kmalloc-128 of size 128
The buggy address is located 123 bytes inside of
128-byte region [ffff8801f44ec300, ffff8801f44ec380)
The buggy address belongs to the page:
page:fffffea0007d13b00 count:1 mapcount:0 mapping:ffff8801f7001640 index:0x0
flags: 0x200000000000100(slab)
raw: 0200000000000100 ffffffea0007d11dc0 0000001a0000001a ffff8801f7001640
raw: 0000000000000000 000000080150015 00000001ffffffff 0000000000000000
page dumped because: kasan: bad access detected

Memory state around the buggy address:

```
ffff8801f44ec200: fc fc fc fc fc fc fc fb fb fb fb fb fb fb fb fb
ffff8801f44ec280: fb fb fb fb fb fb fc fc fc fc fc fc fc fc fc
>ffff8801f44ec300: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 03
                                                               ^
ffff8801f44ec380: fc fc fc fc fc fc fc fb fb fb fb fb fb fb fb
ffff8801f44ec400: fb fb fb fb fb fc fc fc fc fc fc fc fc fc fc
=====
```

The report header summarizes what kind of bug happened and what kind of access caused it. It is followed by a stack trace of the bad access, a stack trace of where the accessed memory was allocated (in case a slab object was accessed), and a stack trace of where the object was freed (in case of a use-after-free bug report). Next comes a description of the accessed slab object and the information about the accessed memory page.

In the end, the report shows the memory state around the accessed address. Internally, KASAN tracks memory state separately for each memory granule, which is either 8 or 16 aligned bytes depending on KASAN mode. Each number in the memory state section of the report shows the state of one of the memory granules that surround the accessed address.

For Generic KASAN, the size of each memory granule is 8. The state of each granule is encoded in one shadow byte. Those 8 bytes can be accessible, partially accessible, freed, or be a part of a redzone. KASAN uses the following encoding for each shadow byte: 00 means that all 8 bytes of the corresponding memory region are accessible; number N (1 <= N <= 7) means that the first N bytes are accessible, and other (8 - N) bytes are not; any negative value indicates that the entire 8-byte word is inaccessible. KASAN uses different negative values to distinguish between different kinds of inaccessible memory like redzones or freed memory (see mm/kasan/kasan.h).

In the report above, the arrow points to the shadow byte 03, which means that the accessed address is partially accessible.

For tag-based KASAN modes, this last report section shows the memory tags around the accessed address (see the *Implementation details* section).

Note that KASAN bug titles (like `slab-out-of-bounds` or `use-after-free`) are best-effort: KASAN prints the most probable bug type based on the limited information it has. The actual type of the bug might be different.

Generic KASAN also reports up to two auxiliary call stack traces. These stack traces point to

places in code that interacted with the object but that are not directly present in the bad access stack trace. Currently, this includes `call_rcu()` and workqueue queuing.

7.4 Implementation details

7.4.1 Generic KASAN

Software KASAN modes use shadow memory to record whether each byte of memory is safe to access and use compile-time instrumentation to insert shadow memory checks before each memory access.

Generic KASAN dedicates 1/8th of kernel memory to its shadow memory (16TB to cover 128TB on x86_64) and uses direct mapping with a scale and offset to translate a memory address to its corresponding shadow address.

Here is the function which translates an address to its corresponding shadow address:

```
static inline void *kasan_mem_to_shadow(const void *addr)
{
    return (void *)((unsigned long)addr >> KASAN_SHADOW_SCALE_SHIFT)
           + KASAN_SHADOW_OFFSET;
}
```

where `KASAN_SHADOW_SCALE_SHIFT = 3`.

Compile-time instrumentation is used to insert memory access checks. Compiler inserts function calls (`__asan_load*(addr)`, `__asan_store*(addr)`) before each memory access of size 1, 2, 4, 8, or 16. These functions check whether memory accesses are valid or not by checking corresponding shadow memory.

With inline instrumentation, instead of making function calls, the compiler directly inserts the code to check shadow memory. This option significantly enlarges the kernel, but it gives an x1.1-x2 performance boost over the outline-instrumented kernel.

Generic KASAN is the only mode that delays the reuse of freed objects via quarantine (see `mm/kasan/quarantine.c` for implementation).

7.4.2 Software Tag-Based KASAN

Software Tag-Based KASAN uses a software memory tagging approach to checking access validity. It is currently only implemented for the arm64 architecture.

Software Tag-Based KASAN uses the Top Byte Ignore (TBI) feature of arm64 CPUs to store a pointer tag in the top byte of kernel pointers. It uses shadow memory to store memory tags associated with each 16-byte memory cell (therefore, it dedicates 1/16th of the kernel memory for shadow memory).

On each memory allocation, Software Tag-Based KASAN generates a random tag, tags the allocated memory with this tag, and embeds the same tag into the returned pointer.

Software Tag-Based KASAN uses compile-time instrumentation to insert checks before each memory access. These checks make sure that the tag of the memory that is being accessed is

equal to the tag of the pointer that is used to access this memory. In case of a tag mismatch, Software Tag-Based KASAN prints a bug report.

Software Tag-Based KASAN also has two instrumentation modes (outline, which emits callbacks to check memory accesses; and inline, which performs the shadow memory checks inline). With outline instrumentation mode, a bug report is printed from the function that performs the access check. With inline instrumentation, a `brk` instruction is emitted by the compiler, and a dedicated `brk` handler is used to print bug reports.

Software Tag-Based KASAN uses 0xFF as a match-all pointer tag (accesses through pointers with the 0xFF pointer tag are not checked). The value 0xFE is currently reserved to tag freed memory regions.

7.4.3 Hardware Tag-Based KASAN

Hardware Tag-Based KASAN is similar to the software mode in concept but uses hardware memory tagging support instead of compiler instrumentation and shadow memory.

Hardware Tag-Based KASAN is currently only implemented for arm64 architecture and based on both arm64 Memory Tagging Extension (MTE) introduced in ARMv8.5 Instruction Set Architecture and Top Byte Ignore (TBI).

Special arm64 instructions are used to assign memory tags for each allocation. Same tags are assigned to pointers to those allocations. On every memory access, hardware makes sure that the tag of the memory that is being accessed is equal to the tag of the pointer that is used to access this memory. In case of a tag mismatch, a fault is generated, and a report is printed.

Hardware Tag-Based KASAN uses 0xFF as a match-all pointer tag (accesses through pointers with the 0xFF pointer tag are not checked). The value 0xFE is currently reserved to tag freed memory regions.

If the hardware does not support MTE (pre ARMv8.5), Hardware Tag-Based KASAN will not be enabled. In this case, all KASAN boot parameters are ignored.

Note that enabling `CONFIG_KASAN_HW_TAGS` always results in in-kernel TBI being enabled. Even when `kasan.mode=off` is provided or when the hardware does not support MTE (but supports TBI).

Hardware Tag-Based KASAN only reports the first found bug. After that, MTE tag checking gets disabled.

7.5 Shadow memory

The contents of this section are only applicable to software KASAN modes.

The kernel maps memory in several different parts of the address space. The range of kernel virtual addresses is large: there is not enough real memory to support a real shadow region for every address that could be accessed by the kernel. Therefore, KASAN only maps real shadow for certain parts of the address space.

7.5.1 Default behaviour

By default, architectures only map real memory over the shadow region for the linear mapping (and potentially other small areas). For all other areas - such as `vmalloc` and `vmemmap` space - a single read-only page is mapped over the shadow area. This read-only shadow page declares all memory accesses as permitted.

This presents a problem for modules: they do not live in the linear mapping but in a dedicated module space. By hooking into the module allocator, KASAN temporarily maps real shadow memory to cover them. This allows detection of invalid accesses to module globals, for example.

This also creates an incompatibility with `VMAP_STACK`: if the stack lives in `vmalloc` space, it will be shadowed by the read-only page, and the kernel will fault when trying to set up the shadow data for stack variables.

7.5.2 CONFIG_KASAN_VMALLOC

With `CONFIG_KASAN_VMALLOC`, KASAN can cover `vmalloc` space at the cost of greater memory usage. Currently, this is supported on x86, arm64, riscv, s390, and powerpc.

This works by hooking into `vmalloc` and `vmap` and dynamically allocating real shadow memory to back the mappings.

Most mappings in `vmalloc` space are small, requiring less than a full page of shadow space. Allocating a full shadow page per mapping would therefore be wasteful. Furthermore, to ensure that different mappings use different shadow pages, mappings would have to be aligned to `KASAN_GRANULE_SIZE * PAGE_SIZE`.

Instead, KASAN shares backing space across multiple mappings. It allocates a backing page when a mapping in `vmalloc` space uses a particular page of the shadow region. This page can be shared by other `vmalloc` mappings later on.

KASAN hooks into the `vmap` infrastructure to lazily clean up unused shadow memory.

To avoid the difficulties around swapping mappings around, KASAN expects that the part of the shadow region that covers the `vmalloc` space will not be covered by the early shadow page but will be left unmapped. This will require changes in arch-specific code.

This allows `VMAP_STACK` support on x86 and can simplify support of architectures that do not have a fixed module region.

7.6 For developers

7.6.1 Ignoring accesses

Software KASAN modes use compiler instrumentation to insert validity checks. Such instrumentation might be incompatible with some parts of the kernel, and therefore needs to be disabled.

Other parts of the kernel might access metadata for allocated objects. Normally, KASAN detects and reports such accesses, but in some cases (e.g., in memory allocators), these accesses are valid.

For software KASAN modes, to disable instrumentation for a specific file or directory, add a `KASAN_SANITIZE` annotation to the respective kernel Makefile:

- For a single file (e.g., `main.o`):

```
KASAN_SANITIZE_main.o := n
```

- For all files in one directory:

```
KASAN_SANITIZE := n
```

For software KASAN modes, to disable instrumentation on a per-function basis, use the KASAN-specific `__no_sanitize_address` function attribute or the generic `noinstr` one.

Note that disabling compiler instrumentation (either on a per-file or a per-function basis) makes KASAN ignore the accesses that happen directly in that code for software KASAN modes. It does not help when the accesses happen indirectly (through calls to instrumented functions) or with Hardware Tag-Based KASAN, which does not use compiler instrumentation.

For software KASAN modes, to disable KASAN reports in a part of the kernel code for the current task, annotate this part of the code with a `kasan_disable_current()/kasan_enable_current()` section. This also disables the reports for indirect accesses that happen through function calls.

For tag-based KASAN modes, to disable access checking, use `kasan_reset_tag()` or `page_kasan_tag_reset()`. Note that temporarily disabling access checking via `page_kasan_tag_reset()` requires saving and restoring the per-page KASAN tag via `page_kasan_tag/page_kasan_tag_set`.

7.6.2 Tests

There are KASAN tests that allow verifying that KASAN works and can detect certain types of memory corruptions. The tests consist of two parts:

1. Tests that are integrated with the KUnit Test Framework. Enabled with `CONFIG_KASAN_KUNIT_TEST`. These tests can be run and partially verified automatically in a few different ways; see the instructions below.
2. Tests that are currently incompatible with KUnit. Enabled with `CONFIG_KASAN_MODULE_TEST` and can only be run as a module. These tests can only be verified manually by loading the kernel module and inspecting the kernel log for KASAN reports.

Each KUnit-compatible KASAN test prints one of multiple KASAN reports if an error is detected. Then the test prints its number and status.

When a test passes:

```
ok 28 - kmalloc_double_kzfree
```

When a test fails due to a failed `kmalloc`:

```
# kmalloc_large_oob_right: ASSERTION FAILED at lib/test_kasan.c:163
Expected ptr is not null, but is
not ok 4 - kmalloc_large_oob_right
```

When a test fails due to a missing KASAN report:

```
# kmalloc_double_kzfree: EXPECTATION FAILED at lib/test_kasan.c:974
KASAN failure expected in "kfree_sensitive(ptr)", but none occurred
not ok 44 - kmalloc_double_kzfree
```

At the end the cumulative status of all KASAN tests is printed. On success:

```
ok 1 - kasan
```

Or, if one of the tests failed:

```
not ok 1 - kasan
```

There are a few ways to run KUnit-compatible KASAN tests.

1. Loadable module

With `CONFIG_KUNIT` enabled, KASAN-KUnit tests can be built as a loadable module and run by loading `test_kasan.ko` with `insmod` or `modprobe`.

2. Built-In

With `CONFIG_KUNIT` built-in, KASAN-KUnit tests can be built-in as well. In this case, the tests will run at boot as a late-init call.

3. Using `kunit_tool`

With `CONFIG_KUNIT` and `CONFIG_KASAN_KUNIT_TEST` built-in, it is also possible to use `kunit_tool` to see the results of KUnit tests in a more readable way. This will not print the KASAN reports of the tests that passed. See [KUnit documentation](#) for more up-to-date information on `kunit_tool`.

KERNEL MEMORY SANITIZER (KMSAN)

KMSAN is a dynamic error detector aimed at finding uses of uninitialized values. It is based on compiler instrumentation, and is quite similar to the userspace [MemorySanitizer](#) tool.

An important note is that KMSAN is not intended for production use, because it drastically increases kernel memory footprint and slows the whole system down.

8.1 Usage

8.1.1 Building the kernel

In order to build a kernel with KMSAN you will need a fresh Clang (14.0.6+). Please refer to [LLVM documentation](#) for the instructions on how to build Clang.

Now configure and build the kernel with CONFIG_KMSAN enabled.

8.1.2 Example report

Here is an example of a KMSAN report:

```
=====
BUG: KMSAN: uninit-value in test_uninit_kmsan_check_memory+0x1be/0x380 [kmsan_
↳ test]
test_uninit_kmsan_check_memory+0x1be/0x380 mm/kmsan/kmsan_test.c:273
kunit_run_case_internal lib/kunit/test.c:333
kunit_try_run_case+0x206/0x420 lib/kunit/test.c:374
kunit_generic_run_threadfn_adapter+0x6d/0xc0 lib/kunit/try-catch.c:28
kthread+0x721/0x850 kernel/kthread.c:327
ret_from_fork+0x1f/0x30 ???:?
```



```
Uninit was stored to memory at:
do_uninit_local_array+0xfa/0x110 mm/kmsan/kmsan_test.c:260
test_uninit_kmsan_check_memory+0x1a2/0x380 mm/kmsan/kmsan_test.c:271
kunit_run_case_internal lib/kunit/test.c:333
kunit_try_run_case+0x206/0x420 lib/kunit/test.c:374
kunit_generic_run_threadfn_adapter+0x6d/0xc0 lib/kunit/try-catch.c:28
kthread+0x721/0x850 kernel/kthread.c:327
ret_from_fork+0x1f/0x30 ???:?
```

```

Local variable uninit created at:
do_uninit_local_array+0x4a/0x110 mm/kmsan/kmsan_test.c:256
test_uninit_kmsan_check_memory+0x1a2/0x380 mm/kmsan/kmsan_test.c:271

Bytes 4-7 of 8 are uninitialized
Memory access of size 8 starts at ffff888083fe3da0

CPU: 0 PID: 6731 Comm: kunit_try_catch Tainted: G     B      E      5.16.0-rc3+
        ↳#104
Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS 1.14.0-2 04/01/2014
=====

```

The report says that the local variable `uninit` was created uninitialized in `do_uninit_local_array()`. The third stack trace corresponds to the place where this variable was created.

The first stack trace shows where the `uninit` value was used (in `test_uninit_kmsan_check_memory()`). The tool shows the bytes which were left uninitialized in the local variable, as well as the stack where the value was copied to another memory location before use.

A use of uninitialized value `v` is reported by KMSAN in the following cases:

- in a condition, e.g. `if (v) { ... };`
- in an indexing or pointer dereferencing, e.g. `array[v]` or `*v`;
- when it is copied to userspace or hardware, e.g. `copy_to_user(..., &v, ...)`;
- when it is passed as an argument to a function, and `CONFIG_KMSAN_CHECK_PARAM_RETVAL` is enabled (see below).

The mentioned cases (apart from copying data to userspace or hardware, which is a security issue) are considered undefined behavior from the C11 Standard point of view.

8.1.3 Disabling the instrumentation

A function can be marked with `__no_kmsan_checks`. Doing so makes KMSAN ignore uninitialized values in that function and mark its output as initialized. As a result, the user will not get KMSAN reports related to that function.

Another function attribute supported by KMSAN is `__no_sanitize_memory`. Applying this attribute to a function will result in KMSAN not instrumenting it, which can be helpful if we do not want the compiler to interfere with some low-level code (e.g. that marked with `noinstr` which implicitly adds `__no_sanitize_memory`).

This however comes at a cost: stack allocations from such functions will have incorrect shadow/origin values, likely leading to false positives. Functions called from non-instrumented code may also receive incorrect metadata for their parameters.

As a rule of thumb, avoid using `__no_sanitize_memory` explicitly.

It is also possible to disable KMSAN for a single file (e.g. `main.o`):

```
KMSAN_SANITIZE_main.o := n
```

or for the whole directory:

```
KMSAN_SANITIZE := n
```

in the Makefile. Think of this as applying `__no_sanitize_memory` to every function in the file or directory. Most users won't need `KMSAN_SANITIZE`, unless their code gets broken by KMSAN (e.g. runs at early boot time).

8.2 Support

In order for KMSAN to work the kernel must be built with Clang, which so far is the only compiler that has KMSAN support. The kernel instrumentation pass is based on the userspace [MemorySanitizer tool](#).

The runtime library only supports `x86_64` at the moment.

8.3 How KMSAN works

8.3.1 KMSAN shadow memory

KMSAN associates a metadata byte (also called shadow byte) with every byte of kernel memory. A bit in the shadow byte is set iff the corresponding bit of the kernel memory byte is uninitialized. Marking the memory uninitialized (i.e. setting its shadow bytes to `0xff`) is called poisoning, marking it initialized (setting the shadow bytes to `0x00`) is called unpoisoning.

When a new variable is allocated on the stack, it is poisoned by default by instrumentation code inserted by the compiler (unless it is a stack variable that is immediately initialized). Any new heap allocation done without `__GFP_ZERO` is also poisoned.

Compiler instrumentation also tracks the shadow values as they are used along the code. When needed, instrumentation code invokes the runtime library in `mm/kmsan/` to persist shadow values.

The shadow value of a basic or compound type is an array of bytes of the same length. When a constant value is written into memory, that memory is unpoisoned. When a value is read from memory, its shadow memory is also obtained and propagated into all the operations which use that value. For every instruction that takes one or more values the compiler generates code that calculates the shadow of the result depending on those values and their shadows.

Example:

```
int a = 0xff; // i.e. 0x000000ff
int b;
int c = a | b;
```

In this case the shadow of `a` is `0`, shadow of `b` is `0xffffffff`, shadow of `c` is `0xfffffff0`. This means that the upper three bytes of `c` are uninitialized, while the lower byte is initialized.

8.3.2 Origin tracking

Every four bytes of kernel memory also have a so-called origin mapped to them. This origin describes the point in program execution at which the uninitialized value was created. Every origin is associated with either the full allocation stack (for heap-allocated memory), or the function containing the uninitialized variable (for locals).

When an uninitialized variable is allocated on stack or heap, a new origin value is created, and that variable's origin is filled with that value. When a value is read from memory, its origin is also read and kept together with the shadow. For every instruction that takes one or more values, the origin of the result is one of the origins corresponding to any of the uninitialized inputs. If a poisoned value is written into memory, its origin is written to the corresponding storage as well.

Example 1:

```
int a = 42;
int b;
int c = a + b;
```

In this case the origin of `b` is generated upon function entry, and is stored to the origin of `c` right before the addition result is written into memory.

Several variables may share the same origin address, if they are stored in the same four-byte chunk. In this case every write to either variable updates the origin for all of them. We have to sacrifice precision in this case, because storing origins for individual bits (and even bytes) would be too costly.

Example 2:

```
int combine(short a, short b) {
    union ret_t {
        int i;
        short s[2];
    } ret;
    ret.s[0] = a;
    ret.s[1] = b;
    return ret.i;
}
```

If `a` is initialized and `b` is not, the shadow of the result would be `0xffff0000`, and the origin of the result would be the origin of `b`. `ret.s[0]` would have the same origin, but it will never be used, because that variable is initialized.

If both function arguments are uninitialized, only the origin of the second argument is preserved.

Origin chaining

To ease debugging, KMSAN creates a new origin for every store of an uninitialized value to memory. The new origin references both its creation stack and the previous origin the value had. This may cause increased memory consumption, so we limit the length of origin chains in the runtime.

8.3.3 Clang instrumentation API

Clang instrumentation pass inserts calls to functions defined in `mm/kmsan/instrumentation.c` into the kernel code.

Shadow manipulation

For every memory access the compiler emits a call to a function that returns a pair of pointers to the shadow and origin addresses of the given memory:

```
typedef struct {
    void *shadow, *origin;
} shadow_origin_ptr_t

shadow_origin_ptr_t __msan_metadata_ptr_for_load_{1,2,4,8}(void *addr)
shadow_origin_ptr_t __msan_metadata_ptr_for_store_{1,2,4,8}(void *addr)
shadow_origin_ptr_t __msan_metadata_ptr_for_load_n(void *addr, uintptr_t size)
shadow_origin_ptr_t __msan_metadata_ptr_for_store_n(void *addr, uintptr_t size)
```

The function name depends on the memory access size.

The compiler makes sure that for every loaded value its shadow and origin values are read from memory. When a value is stored to memory, its shadow and origin are also stored using the metadata pointers.

Handling locals

A special function is used to create a new origin value for a local variable and set the origin of that variable to that value:

```
void __msan_poison_alloca(void *addr, uintptr_t size, char *descr)
```

Access to per-task data

At the beginning of every instrumented function KMSAN inserts a call to `__msan_get_context_state()`:

```
kmsan_context_state *__msan_get_context_state(void)
```

`kmsan_context_state` is declared in `include/linux/kmsan.h`:

```
struct kmsan_context_state {
    char param_tls[KMSAN_PARAM_SIZE];
    char retval_tls[KMSAN_RETVAL_SIZE];
    char va_arg_tls[KMSAN_PARAM_SIZE];
    char va_arg_origin_tls[KMSAN_PARAM_SIZE];
    u64 va_arg_overflow_size_tls;
    char param_origin_tls[KMSAN_PARAM_SIZE];
    depot_stack_handle_t retval_origin_tls;
};
```

This structure is used by KMSAN to pass parameter shadows and origins between instrumented functions (unless the parameters are checked immediately by `CONFIG_KMSAN_CHECK_PARAM_RETVAL`).

Passing uninitialized values to functions

Clang's MemorySanitizer instrumentation has an option, `-fsanitize-memory-param-retval`, which makes the compiler check function parameters passed by value, as well as function return values.

The option is controlled by `CONFIG_KMSAN_CHECK_PARAM_RETVAL`, which is enabled by default to let KMSAN report uninitialized values earlier. Please refer to the [LKML discussion](#) for more details.

Because of the way the checks are implemented in LLVM (they are only applied to parameters marked as `noundef`), not all parameters are guaranteed to be checked, so we cannot give up the metadata storage in `kmsan_context_state`.

String functions

The compiler replaces calls to `memcpy()`/`memmove()`/`memset()` with the following functions. These functions are also called when data structures are initialized or copied, making sure shadow and origin values are copied alongside with the data:

```
void * __msan_memcpy(void *dst, void *src, uintptr_t n)
void * __msan_memmove(void *dst, void *src, uintptr_t n)
void * __msan_memset(void *dst, int c, uintptr_t n)
```

Error reporting

For each use of a value the compiler emits a shadow check that calls `__msan_warning()` in the case that value is poisoned:

```
void __msan_warning(u32 origin)
```

`__msan_warning()` causes KMSAN runtime to print an error report.

Inline assembly instrumentation

KMSAN instruments every inline assembly output with a call to:

```
void __msan_instrument_asm_store(void *addr, uintptr_t size)
```

, which unpoisons the memory region.

This approach may mask certain errors, but it also helps to avoid a lot of false positives in bitwise operations, atomics etc.

Sometimes the pointers passed into inline assembly do not point to valid memory. In such cases they are ignored at runtime.

8.3.4 Runtime library

The code is located in `mm/kmsan/`.

Per-task KMSAN state

Every `task_struct` has an associated KMSAN task state that holds the KMSAN context (see above) and a per-task flag disallowing KMSAN reports:

```
struct kmsan_context {
    ...
    bool allow_reporting;
    struct kmsan_context_state cstate;
    ...
}

struct task_struct {
    ...
    struct kmsan_context kmsan;
    ...
}
```

KMSAN contexts

When running in a kernel task context, KMSAN uses `current->kmsan.cstate` to hold the metadata for function parameters and return values.

But in the case the kernel is running in the interrupt, softirq or NMI context, where `current` is unavailable, KMSAN switches to per-cpu interrupt state:

```
DEFINE_PER_CPU(struct kmsan_ctx, kmsan_percpu_ctx);
```

Metadata allocation

There are several places in the kernel for which the metadata is stored.

1. Each `struct page` instance contains two pointers to its shadow and origin pages:

```
struct page {
    ...
    struct page *shadow, *origin;
    ...
};
```

At boot-time, the kernel allocates shadow and origin pages for every available kernel page. This is done quite late, when the kernel address space is already fragmented, so normal data pages may arbitrarily interleave with the metadata pages.

This means that in general for two contiguous memory pages their shadow/origin pages may not be contiguous. Consequently, if a memory access crosses the boundary of a memory block, accesses to shadow/origin memory may potentially corrupt other pages or read incorrect values from them.

In practice, contiguous memory pages returned by the same `alloc_pages()` call will have contiguous metadata, whereas if these pages belong to two different allocations their metadata pages can be fragmented.

For the kernel data (`.data`, `.bss` etc.) and percpu memory regions there also are no guarantees on metadata contiguity.

In the case `__msan_metadata_ptr_for_XXX_YYY()` hits the border between two pages with non-contiguous metadata, it returns pointers to fake shadow/origin regions:

```
char dummy_load_page[PAGE_SIZE] __attribute__((aligned(PAGE_SIZE)));
char dummy_store_page[PAGE_SIZE] __attribute__((aligned(PAGE_SIZE)));
```

`dummy_load_page` is zero-initialized, so reads from it always yield zeroes. All stores to `dummy_store_page` are ignored.

2. For `vmalloc` memory and modules, there is a direct mapping between the memory range, its shadow and origin. KMSAN reduces the `vmalloc` area by 3/4, making only the first quarter available to `vmalloc()`. The second quarter of the `vmalloc` area contains shadow memory for the first quarter, the third one holds the origins. A small part of the fourth quarter contains shadow and origins for the kernel modules. Please refer to `arch/x86/include/asm/pgtable_64_types.h` for more details.

When an array of pages is mapped into a contiguous virtual memory space, their shadow and origin pages are similarly mapped into contiguous regions.

8.4 References

E. Stepanov, K. Serebryany. MemorySanitizer: fast detector of uninitialized memory use in C++. In Proceedings of CGO 2015.

UNDEFINED BEHAVIOR SANITIZER - UBSAN

UBSAN is a runtime undefined behaviour checker.

UBSAN uses compile-time instrumentation to catch undefined behavior (UB). Compiler inserts code that perform certain kinds of checks before operations that may cause UB. If check fails (i.e. UB detected) `_ubsan_handle_*` function called to print error message.

GCC has that feature since 4.9.x [1] (see `-fsanitize=undefined` option and its suboptions). GCC 5.x has more checkers implemented [2].

9.1 Report example

```
=====
UBSAN: Undefined behaviour in ../include/linux/bitops.h:110:33
shift exponent 32 is to large for 32-bit type 'unsigned int'
CPU: 0 PID: 0 Comm: swapper Not tainted 4.4.0-rc1+ #26
0000000000000000 ffffffff82403cc8 ffffffff815e6cd6 0000000000000001
ffffffff82403cf8 ffffffff82403ce0 ffffffff8163a5ed 0000000000000020
ffffffff82403d78 ffffffff8163ac2b ffffffff815f0001 0000000000000002
Call Trace:
[<fffffff815e6cd6>] dump_stack+0x45/0x5f
[<fffffff8163a5ed>] ubsan_epilogue+0xd/0x40
[<fffffff8163ac2b>] __ubsan_handle_shift_out_of_bounds+0xeb/0x130
[<fffffff815f0001>] ? radix_tree_gang_lookup_slot+0x51/0x150
[<fffffff8173c586>] _mix_pool_bytes+0x1e6/0x480
[<fffffff83105653>] ? dmi_walk_early+0x48/0x5c
[<fffffff8173c881>] add_device_randomness+0x61/0x130
[<fffffff83105b35>] ? dmi_save_one_device+0xaa/0xaa
[<fffffff83105653>] dmi_walk_early+0x48/0x5c
[<fffffff831066ae>] dmi_scan_machine+0x278/0x4b4
[<fffffff8111d58a>] ? vprintk_default+0x1a/0x20
[<fffffff830ad120>] ? early_idt_handler_array+0x120/0x120
[<fffffff830b2240>] setup_arch+0x405/0xc2c
[<fffffff830ad120>] ? early_idt_handler_array+0x120/0x120
[<fffffff830ae053>] start_kernel+0x83/0x49a
[<fffffff830ad120>] ? early_idt_handler_array+0x120/0x120
[<fffffff830ad386>] x86_64_start_reservations+0x2a/0x2c
[<fffffff830ad4f3>] x86_64_start_kernel+0x16b/0x17a
=====
```

9.2 Usage

To enable UBSAN configure kernel with:

```
CONFIG_UBSAN=y
```

and to check the entire kernel:

```
CONFIG_UBSAN_SANITIZE_ALL=y
```

To enable instrumentation for specific files or directories, add a line similar to the following to the respective kernel Makefile:

- For a single file (e.g. main.o):

```
UBSAN_SANITIZE_main.o := y
```

- For all files in one directory:

```
UBSAN_SANITIZE := y
```

To exclude files from being instrumented even if CONFIG_UBSAN_SANITIZE_ALL=y, use:

```
UBSAN_SANITIZE_main.o := n
```

and:

```
UBSAN_SANITIZE := n
```

Detection of unaligned accesses controlled through the separate option - CONFIG_UBSAN_ALIGNMENT. It's off by default on architectures that support unaligned accesses (CONFIG_HAVE_EFFICIENT_UNALIGNED_ACCESS=y). One could still enable it in config, just note that it will produce a lot of UBSAN reports.

9.3 References

KERNEL MEMORY LEAK DETECTOR

Kmemleak provides a way of detecting possible kernel memory leaks in a way similar to a [tracing garbage collector](#), with the difference that the orphan objects are not freed but only reported via /sys/kernel/debug/kmemleak. A similar method is used by the Valgrind tool (memcheck -- leak-check) to detect the memory leaks in user-space applications.

10.1 Usage

CONFIG_DEBUG_KMEMLEAK in "Kernel hacking" has to be enabled. A kernel thread scans the memory every 10 minutes (by default) and prints the number of new unreferenced objects found. If the debugfs isn't already mounted, mount with:

```
# mount -t debugfs nodev /sys/kernel/debug/
```

To display the details of all the possible scanned memory leaks:

```
# cat /sys/kernel/debug/kmemleak
```

To trigger an intermediate memory scan:

```
# echo scan > /sys/kernel/debug/kmemleak
```

To clear the list of all current possible memory leaks:

```
# echo clear > /sys/kernel/debug/kmemleak
```

New leaks will then come up upon reading /sys/kernel/debug/kmemleak again.

Note that the orphan objects are listed in the order they were allocated and one object at the beginning of the list may cause other subsequent objects to be reported as orphan.

Memory scanning parameters can be modified at run-time by writing to the /sys/kernel/debug/kmemleak file. The following parameters are supported:

- **off**
 disable kmemleak (irreversible)
- **stack=on**
 enable the task stacks scanning (default)
- **stack=off**
 disable the tasks stacks scanning

- **scan=on**
start the automatic memory scanning thread (default)
- **scan=off**
stop the automatic memory scanning thread
- **scan=<secs>**
set the automatic memory scanning period in seconds (default 600, 0 to stop the automatic scanning)
- **scan**
trigger a memory scan
- **clear**
clear list of current memory leak suspects, done by marking all current reported unreferenced objects grey, or free all kmemleak objects if kmemleak has been disabled.
- **dump=<addr>**
dump information about the object found at <addr>

Kmemleak can also be disabled at boot-time by passing `kmemleak=off` on the kernel command line.

Memory may be allocated or freed before kmemleak is initialised and these actions are stored in an early log buffer. The size of this buffer is configured via the `CONFIG_DEBUG_KMEMLEAK_MEM_POOL_SIZE` option.

If `CONFIG_DEBUG_KMEMLEAK_DEFAULT_OFF` are enabled, the kmemleak is disabled by default. Passing `kmemleak=on` on the kernel command line enables the function.

If you are getting errors like "Error while writing to stdout" or "write_loop: Invalid argument", make sure kmemleak is properly enabled.

10.2 Basic Algorithm

The memory allocations via `kmalloc()`, `vmalloc()`, `kmem_cache_alloc()` and friends are traced and the pointers, together with additional information like size and stack trace, are stored in a rbtree. The corresponding freeing function calls are tracked and the pointers removed from the kmemleak data structures.

An allocated block of memory is considered orphan if no pointer to its start address or to any location inside the block can be found by scanning the memory (including saved registers). This means that there might be no way for the kernel to pass the address of the allocated block to a freeing function and therefore the block is considered a memory leak.

The scanning algorithm steps:

1. mark all objects as white (remaining white objects will later be considered orphan)
2. scan the memory starting with the data section and stacks, checking the values against the addresses stored in the rbtree. If a pointer to a white object is found, the object is added to the gray list
3. scan the gray objects for matching addresses (some white objects can become gray and added at the end of the gray list) until the gray set is finished

4. the remaining white objects are considered orphan and reported via /sys/kernel/debug/kmemleak

Some allocated memory blocks have pointers stored in the kernel's internal data structures and they cannot be detected as orphans. To avoid this, kmemleak can also store the number of values pointing to an address inside the block address range that need to be found so that the block is not considered a leak. One example is `_vmalloc()`.

10.3 Testing specific sections with kmemleak

Upon initial bootup your /sys/kernel/debug/kmemleak output page may be quite extensive. This can also be the case if you have very buggy code when doing development. To work around these situations you can use the 'clear' command to clear all reported unreferenced objects from the /sys/kernel/debug/kmemleak output. By issuing a 'scan' after a 'clear' you can find new unreferenced objects; this should help with testing specific sections of code.

To test a critical section on demand with a clean kmemleak do:

```
# echo clear > /sys/kernel/debug/kmemleak
... test your kernel or modules ...
# echo scan > /sys/kernel/debug/kmemleak
```

Then as usual to get your report with:

```
# cat /sys/kernel/debug/kmemleak
```

10.4 Freeing kmemleak internal objects

To allow access to previously found memory leaks after kmemleak has been disabled by the user or due to a fatal error, internal kmemleak objects won't be freed when kmemleak is disabled, and those objects may occupy a large part of physical memory.

In this situation, you may reclaim memory with:

```
# echo clear > /sys/kernel/debug/kmemleak
```

10.5 Kmemleak API

See the include/linux/kmemleak.h header for the functions prototype.

- `kmemleak_init` - initialize kmemleak
- `kmemleak_alloc` - notify of a memory block allocation
- `kmemleak_alloc_percpu` - notify of a percpu memory block allocation
- `kmemleak_vmalloc` - notify of a `vmalloc()` memory allocation
- `kmemleak_free` - notify of a memory block freeing
- `kmemleak_free_part` - notify of a partial memory block freeing

- `kmemleak_free_percpu` - notify of a percpu memory block freeing
- `kmemleak_update_trace` - update object allocation stack trace
- `kmemleak_not_leak` - mark an object as not a leak
- `kmemleak_ignore` - do not scan or report an object as leak
- `kmemleak_scan_area` - add scan areas inside a memory block
- `kmemleak_no_scan` - do not scan a memory block
- `kmemleak_erase` - erase an old value in a pointer variable
- `kmemleak_alloc_recursive` - as `kmemleak_alloc` but checks the recursiveness
- `kmemleak_free_recursive` - as `kmemleak_free` but checks the recursiveness

The following functions take a physical address as the object pointer and only perform the corresponding action if the address has a lowmem mapping:

- `kmemleak_alloc_phys`
- `kmemleak_free_part_phys`
- `kmemleak_ignore_phys`

10.6 Dealing with false positives/negatives

The false negatives are real memory leaks (orphan objects) but not reported by kmemleak because values found during the memory scanning point to such objects. To reduce the number of false negatives, kmemleak provides the `kmemleak_ignore`, `kmemleak_scan_area`, `kmemleak_no_scan` and `kmemleak_erase` functions (see above). The task stacks also increase the amount of false negatives and their scanning is not enabled by default.

The false positives are objects wrongly reported as being memory leaks (orphan). For objects known not to be leaks, kmemleak provides the `kmemleak_not_leak` function. The `kmemleak_ignore` could also be used if the memory block is known not to contain other pointers and it will no longer be scanned.

Some of the reported leaks are only transient, especially on SMP systems, because of pointers temporarily stored in CPU registers or stacks. Kmemleak defines `MSECS_MIN_AGE` (defaulting to 1000) representing the minimum age of an object to be reported as a memory leak.

10.7 Limitations and Drawbacks

The main drawback is the reduced performance of memory allocation and freeing. To avoid other penalties, the memory scanning is only performed when the `/sys/kernel/debug/kmemleak` file is read. Anyway, this tool is intended for debugging purposes where the performance might not be the most important requirement.

To keep the algorithm simple, kmemleak scans for values pointing to any address inside a block's address range. This may lead to an increased number of false negatives. However, it is likely that a real memory leak will eventually become visible.

Another source of false negatives is the data stored in non-pointer values. In a future version, kmemleak could only scan the pointer members in the allocated structures. This feature would solve many of the false negative cases described above.

The tool can report false positives. These are cases where an allocated block doesn't need to be freed (some cases in the init_call functions), the pointer is calculated by other methods than the usual container_of macro or the pointer is stored in a location not scanned by kmemleak.

Page allocations and ioremap are not tracked.

10.8 Testing with kmemleak-test

To check if you have all set up to use kmemleak, you can use the kmemleak-test module, a module that deliberately leaks memory. Set CONFIG_SAMPLE_KMEMLEAK as module (it can't be used as built-in) and boot the kernel with kmemleak enabled. Load the module and perform a scan with:

```
# modprobe kmemleak-test
# echo scan > /sys/kernel/debug/kmemleak
```

Note that the you may not get results instantly or on the first scanning. When kmemleak gets results, it'll log kmemleak: <count of leaks> new suspected memory leaks. Then read the file to see them:

```
# cat /sys/kernel/debug/kmemleak
unreferenced object 0xfffff89862ca702e8 (size 32):
  comm "modprobe", pid 2088, jiffies 4294680594 (age 375.486s)
  hex dump (first 32 bytes):
    6b 6b
  backtrace:
    [<00000000e0a73ec7>] 0xffffffffc01d2036
    [<000000000c5d2a46>] do_one_initcall+0x41/0x1df
    [<0000000046db7e0a>] do_init_module+0x55/0x200
    [<000000000542b9814>] load_module+0x203c/0x2480
    [<000000000c2850256>] __do_sys_finit_module+0xba/0xe0
    [<0000000006564e7ef>] do_syscall_64+0x43/0x110
    [<0000000007c873fa6>] entry_SYSCALL_64_after_hwframe+0x44/0xa9
...
```

Removing the module with rmmod kmemleak_test should also trigger some kmemleak results.

KERNEL CONCURRENCY SANITIZER (KCSAN)

The Kernel Concurrency Sanitizer (KCSAN) is a dynamic race detector, which relies on compile-time instrumentation, and uses a watchpoint-based sampling approach to detect races. KCSAN's primary purpose is to detect *data races*.

11.1 Usage

KCSAN is supported by both GCC and Clang. With GCC we require version 11 or later, and with Clang also require version 11 or later.

To enable KCSAN configure the kernel with:

```
CONFIG_KCSAN = y
```

KCSAN provides several other configuration options to customize behaviour (see the respective help text in `lib/Kconfig.kcsan` for more info).

11.1.1 Error reports

A typical data race report looks like this:

```
=====
BUG: KCSAN: data-race in test_kernel_read / test_kernel_write

write to 0xfffffffffc009a628 of 8 bytes by task 487 on cpu 0:
test_kernel_write+0x1d/0x30
access_thread+0x89/0xd0
kthread+0x23e/0x260
ret_from_fork+0x22/0x30

read to 0xfffffffffc009a628 of 8 bytes by task 488 on cpu 6:
test_kernel_read+0x10/0x20
access_thread+0x89/0xd0
kthread+0x23e/0x260
ret_from_fork+0x22/0x30

value changed: 0x0000000000000009a6 -> 0x0000000000000009b2
```

Reported by Kernel Concurrency Sanitizer on:

```
CPU: 6 PID: 488 Comm: access_thread Not tainted 5.12.0-rc2+ #1
Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS 1.14.0-2 04/01/2014
=====
```

The header of the report provides a short summary of the functions involved in the race. It is followed by the access types and stack traces of the 2 threads involved in the data race. If KCSAN also observed a value change, the observed old value and new value are shown on the "value changed" line respectively.

The other less common type of data race report looks like this:

```
=====
BUG: KCSAN: data-race in test_kernel_rmw_array+0x71/0xd0

race at unknown origin, with read to 0xffffffffc009bdb0 of 8 bytes by task 515_u
→ on cpu 2:
test_kernel_rmw_array+0x71/0xd0
access_thread+0x89/0xd0
kthread+0x23e/0x260
ret_from_fork+0x22/0x30

value changed: 0x0000000000002328 -> 0x0000000000002329

Reported by Kernel Concurrency Sanitizer on:
CPU: 2 PID: 515 Comm: access_thread Not tainted 5.12.0-rc2+ #1
Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS 1.14.0-2 04/01/2014
=====
```

This report is generated where it was not possible to determine the other racing thread, but a race was inferred due to the data value of the watched memory location having changed. These reports always show a "value changed" line. A common reason for reports of this type are missing instrumentation in the racing thread, but could also occur due to e.g. DMA accesses. Such reports are shown only if `CONFIG_KCSAN_REPORT_RACE_UNKNOWN_ORIGIN=y`, which is enabled by default.

11.1.2 Selective analysis

It may be desirable to disable data race detection for specific accesses, functions, compilation units, or entire subsystems. For static blacklisting, the below options are available:

- KCSAN understands the `data_race(expr)` annotation, which tells KCSAN that any data races due to accesses in `expr` should be ignored and resulting behaviour when encountering a data race is deemed safe. Please see "[Marking Shared-Memory Accesses](#)" in the [LKMM](#) for more information.
- Disabling data race detection for entire functions can be accomplished by using the function attribute `_no_kcsan`:

```
_no_kcsan
void foo(void) {
    ...
}
```

To dynamically limit for which functions to generate reports, see the *DebugFS interface* blacklist/whitelist feature.

- To disable data race detection for a particular compilation unit, add to the `Makefile`:

```
KCSAN_SANITIZE_file.o := n
```

- To disable data race detection for all compilation units listed in a `Makefile`, add to the respective `Makefile`:

```
KCSAN_SANITIZE := n
```

Furthermore, it is possible to tell KCSAN to show or hide entire classes of data races, depending on preferences. These can be changed via the following Kconfig options:

- `CONFIG_KCSAN_REPORT_VALUE_CHANGE_ONLY`: If enabled and a conflicting write is observed via a watchpoint, but the data value of the memory location was observed to remain unchanged, do not report the data race.
- `CONFIG_KCSAN_ASSUME_PLAIN_WRITES_ATOMIC`: Assume that plain aligned writes up to word size are atomic by default. Assumes that such writes are not subject to unsafe compiler optimizations resulting in data races. The option causes KCSAN to not report data races due to conflicts where the only plain accesses are aligned writes up to word size.
- `CONFIG_KCSAN_PERMISSIVE`: Enable additional permissive rules to ignore certain classes of common data races. Unlike the above, the rules are more complex involving value-change patterns, access type, and address. This option depends on `CONFIG_KCSAN_REPORT_VALUE_CHANGE_ONLY=y`. For details please see the `kernel/kcsan/permissive.h`. Testers and maintainers that only focus on reports from specific subsystems and not the whole kernel are recommended to disable this option.

To use the strictest possible rules, select `CONFIG_KCSAN_STRICT=y`, which configures KCSAN to follow the Linux-kernel memory consistency model (LKMM) as closely as possible.

11.1.3 DebugFS interface

The file `/sys/kernel/debug/kcsan` provides the following interface:

- Reading `/sys/kernel/debug/kcsan` returns various runtime statistics.
- Writing `on` or `off` to `/sys/kernel/debug/kcsan` allows turning KCSAN on or off, respectively.
- Writing `!some_func_name` to `/sys/kernel/debug/kcsan` adds `some_func_name` to the report filter list, which (by default) blacklists reporting data races where either one of the top stackframes are a function in the list.
- Writing either `blacklist` or `whitelist` to `/sys/kernel/debug/kcsan` changes the report filtering behaviour. For example, the blacklist feature can be used to silence frequently occurring data races; the whitelist feature can help with reproduction and testing of fixes.

11.1.4 Tuning performance

Core parameters that affect KCSAN's overall performance and bug detection ability are exposed as kernel command-line arguments whose defaults can also be changed via the corresponding Kconfig options.

- `kcsan.skip_watch` (`CONFIG_KCSAN_SKIP_WATCH`): Number of per-CPU memory operations to skip, before another watchpoint is set up. Setting up watchpoints more frequently will result in the likelihood of races to be observed to increase. This parameter has the most significant impact on overall system performance and race detection ability.
- `kcsan.udelay_task` (`CONFIG_KCSAN_UDELAY_TASK`): For tasks, the microsecond delay to stall execution after a watchpoint has been set up. Larger values result in the window in which we may observe a race to increase.
- `kcsan.udelay_interrupt` (`CONFIG_KCSAN_UDELAY_INTERRUPT`): For interrupts, the microsecond delay to stall execution after a watchpoint has been set up. Interrupts have tighter latency requirements, and their delay should generally be smaller than the one chosen for tasks.

They may be tweaked at runtime via `/sys/module/kcsan/parameters/`.

11.2 Data Races

In an execution, two memory accesses form a *data race* if they *conflict*, they happen concurrently in different threads, and at least one of them is a *plain access*; they *conflict* if both access the same memory location, and at least one is a write. For a more thorough discussion and definition, see "[Plain Accesses and Data Races](#)" in the LKMM.

11.2.1 Relationship with the Linux-Kernel Memory Consistency Model (LKMM)

The LKMM defines the propagation and ordering rules of various memory operations, which gives developers the ability to reason about concurrent code. Ultimately this allows to determine the possible executions of concurrent code, and if that code is free from data races.

KCSAN is aware of *marked atomic operations* (`READ_ONCE`, `WRITE_ONCE`, `atomic_*`, etc.), and a subset of ordering guarantees implied by memory barriers. With `CONFIG_KCSAN_WEAK_MEMORY=y`, KCSAN models load or store buffering, and can detect missing `smp_mb()`, `smp_wmb()`, `smp_rmb()`, `smp_store_release()`, and all `atomic_*` operations with equivalent implied barriers.

Note, KCSAN will not report all data races due to missing memory ordering, specifically where a memory barrier would be required to prohibit subsequent memory operation from reordering before the barrier. Developers should therefore carefully consider the required memory ordering requirements that remain unchecked.

11.3 Race Detection Beyond Data Races

For code with complex concurrency design, race-condition bugs may not always manifest as data races. Race conditions occur if concurrently executing operations result in unexpected system behaviour. On the other hand, data races are defined at the C-language level. The following macros can be used to check properties of concurrent code where bugs would not manifest as data races.

ASSERT_EXCLUSIVE_WRITER

```
ASSERT_EXCLUSIVE_WRITER (var)
    assert no concurrent writes to var
```

Parameters

var
variable to assert on

Description

Assert that there are no concurrent writes to **var**; other readers are allowed. This assertion can be used to specify properties of concurrent code, where violation cannot be detected as a normal data race.

For example, if we only have a single writer, but multiple concurrent readers, to avoid data races, all these accesses must be marked; even concurrent marked writes racing with the single writer are bugs. Unfortunately, due to being marked, they are no longer data races. For cases like these, we can use the macro as follows:

```
void writer(void) {
    spin_lock(&update_foo_lock);
    ASSERT_EXCLUSIVE_WRITER(shared_foo);
    WRITE_ONCE(shared_foo, ...);
    spin_unlock(&update_foo_lock);
}
void reader(void) {
    // update_foo_lock does not need to be held!
    ... = READ_ONCE(shared_foo);
}
```

Note

`ASSERT_EXCLUSIVE_WRITER_SCOPED()`, if applicable, performs more thorough checking if a clear scope where no concurrent writes are expected exists.

ASSERT_EXCLUSIVE_WRITER_SCOPED

```
ASSERT_EXCLUSIVE_WRITER_SCOPED (var)
    assert no concurrent writes to var in scope
```

Parameters

var
variable to assert on

Description

Scoped variant of `ASSERT_EXCLUSIVE_WRITER()`.

Assert that there are no concurrent writes to `var` for the duration of the scope in which it is introduced. This provides a better way to fully cover the enclosing scope, compared to multiple `ASSERT_EXCLUSIVE_WRITER()`, and increases the likelihood for KCSAN to detect racing accesses.

For example, it allows finding race-condition bugs that only occur due to state changes within the scope itself:

```
void writer(void) {
    spin_lock(&update_foo_lock);
    {
        ASSERT_EXCLUSIVE_WRITER_SCOPED(shared_foo);
        WRITE_ONCE(shared_foo, 42);
        ...
        // shared_foo should still be 42 here!
    }
    spin_unlock(&update_foo_lock);
}
void buggy(void) {
    if (READ_ONCE(shared_foo) == 42)
        WRITE_ONCE(shared_foo, 1); // bug!
}
```

ASSERT_EXCLUSIVE_ACCESS

`ASSERT_EXCLUSIVE_ACCESS (var)`

assert no concurrent accesses to `var`

Parameters

`var`

variable to assert on

Description

Assert that there are no concurrent accesses to `var` (no readers nor writers). This assertion can be used to specify properties of concurrent code, where violation cannot be detected as a normal data race.

For example, where exclusive access is expected after determining no other users of an object are left, but the object is not actually freed. We can check that this property actually holds as follows:

```
if (refcount_dec_and_test(&obj->refcnt)) {
    ASSERT_EXCLUSIVE_ACCESS(*obj);
    do_some_cleanup(obj);
    release_for_reuse(obj);
}
```

1. `ASSERT_EXCLUSIVE_ACCESS_SCOPED()`, if applicable, performs more thorough checking if a clear scope where no concurrent accesses are expected exists.

2. For cases where the object is freed, KASAN is a better fit to detect use-after-free bugs.

Note

`ASSERT_EXCLUSIVE_ACCESS_SCOPED`

`ASSERT_EXCLUSIVE_ACCESS_SCOPED (var)`

assert no concurrent accesses to **var** in scope

Parameters

var

variable to assert on

Description

Scoped variant of [`ASSERT_EXCLUSIVE_ACCESS\(\)`](#).

Assert that there are no concurrent accesses to **var** (no readers nor writers) for the entire duration of the scope in which it is introduced. This provides a better way to fully cover the enclosing scope, compared to multiple [`ASSERT_EXCLUSIVE_ACCESS\(\)`](#), and increases the likelihood for KCSAN to detect racing accesses.

`ASSERT_EXCLUSIVE_BITS`

`ASSERT_EXCLUSIVE_BITS (var, mask)`

assert no concurrent writes to subset of bits in **var**

Parameters

var

variable to assert on

mask

only check for modifications to bits set in **mask**

Description

Bit-granular variant of [`ASSERT_EXCLUSIVE_WRITER\(\)`](#).

Assert that there are no concurrent writes to a subset of bits in **var**; concurrent readers are permitted. This assertion captures more detailed bit-level properties, compared to the other (word granularity) assertions. Only the bits set in **mask** are checked for concurrent modifications, while ignoring the remaining bits, i.e. concurrent writes (or reads) to \sim **mask** bits are ignored.

Use this for variables, where some bits must not be modified concurrently, yet other bits are expected to be modified concurrently.

For example, variables where, after initialization, some bits are read-only, but other bits may still be modified concurrently. A reader may wish to assert that this is true as follows:

```
ASSERT_EXCLUSIVE_BITS(flags, READ_ONLY_MASK);
foo = (READ_ONCE(flags) & READ_ONLY_MASK) >> READ_ONLY_SHIFT;
```

```
ASSERT_EXCLUSIVE_BITS(flags, READ_ONLY_MASK);
foo = (flags & READ_ONLY_MASK) >> READ_ONLY_SHIFT;
```

Another example, where this may be used, is when certain bits of **var** may only be modified when holding the appropriate lock, but other bits may still be modified concurrently. Writers, where other bits may change concurrently, could use the assertion as follows:

```
spin_lock(&foo_lock);
ASSERT_EXCLUSIVE_BITS(flags, FOO_MASK);
old_flags = flags;
new_flags = (old_flags & ~FOO_MASK) | (new_foo << FOO_SHIFT);
if (cmpxchg(&flags, old_flags, new_flags) != old_flags) { ... }
spin_unlock(&foo_lock);
```

Note

The access that immediately follows `ASSERT_EXCLUSIVE_BITS()` is assumed to access the masked bits only, and KCSAN optimistically assumes it is therefore safe, even in the presence of data races, and marking it with `READ_ONCE()` is optional from KCSAN's point-of-view. We caution, however, that it may still be advisable to do so, since we cannot reason about all compiler optimizations when it comes to bit manipulations (on the reader and writer side). If you are sure nothing can go wrong, we can write the above simply as:

11.4 Implementation Details

KCSAN relies on observing that two accesses happen concurrently. Crucially, we want to (a) increase the chances of observing races (especially for races that manifest rarely), and (b) be able to actually observe them. We can accomplish (a) by injecting various delays, and (b) by using address watchpoints (or breakpoints).

If we deliberately stall a memory access, while we have a watchpoint for its address set up, and then observe the watchpoint to fire, two accesses to the same address just raced. Using hardware watchpoints, this is the approach taken in [DataCollider](#). Unlike DataCollider, KCSAN does not use hardware watchpoints, but instead relies on compiler instrumentation and "soft watchpoints".

In KCSAN, watchpoints are implemented using an efficient encoding that stores access type, size, and address in a long; the benefits of using "soft watchpoints" are portability and greater flexibility. KCSAN then relies on the compiler instrumenting plain accesses. For each instrumented plain access:

1. Check if a matching watchpoint exists; if yes, and at least one access is a write, then we encountered a racing access.
2. Periodically, if no matching watchpoint exists, set up a watchpoint and stall for a small randomized delay.
3. Also check the data value before the delay, and re-check the data value after delay; if the values mismatch, we infer a race of unknown origin.

To detect data races between plain and marked accesses, KCSAN also annotates marked accesses, but only to check if a watchpoint exists; i.e. KCSAN never sets up a watchpoint on marked accesses. By never setting up watchpoints for marked operations, if all accesses to a variable that is accessed concurrently are properly marked, KCSAN will never trigger a watchpoint and therefore never report the accesses.

11.4.1 Modeling Weak Memory

KCSAN's approach to detecting data races due to missing memory barriers is based on modeling access reordering (with `CONFIG_KCSAN_WEAK_MEMORY=y`). Each plain memory access for which a watchpoint is set up, is also selected for simulated reordering within the scope of its function (at most 1 in-flight access).

Once an access has been selected for reordering, it is checked along every other access until the end of the function scope. If an appropriate memory barrier is encountered, the access will no longer be considered for simulated reordering.

When the result of a memory operation should be ordered by a barrier, KCSAN can then detect data races where the conflict only occurs as a result of a missing barrier. Consider the example:

```
int x, flag;
void T1(void)
{
    x = 1;                      // data race!
    WRITE_ONCE(flag, 1);        // correct: smp_store_release(&flag, 1)
}
void T2(void)
{
    while (!READ_ONCE(flag));   // correct: smp_load_acquire(&flag)
    ... = x;                   // data race!
}
```

When weak memory modeling is enabled, KCSAN can consider `x` in `T1` for simulated reordering. After the write of `flag`, `x` is again checked for concurrent accesses: because `T2` is able to proceed after the write of `flag`, a data race is detected. With the correct barriers in place, `x` would not be considered for reordering after the proper release of `flag`, and no data race would be detected.

Deliberate trade-offs in complexity but also practical limitations mean only a subset of data races due to missing memory barriers can be detected. With currently available compiler support, the implementation is limited to modeling the effects of "buffering" (delaying accesses), since the runtime cannot "prefetch" accesses. Also recall that watchpoints are only set up for plain accesses, and the only access type for which KCSAN simulates reordering. This means reordering of marked accesses is not modeled.

A consequence of the above is that acquire operations do not require barrier instrumentation (no prefetching). Furthermore, marked accesses introducing address or control dependencies do not require special handling (the marked access cannot be reordered, later dependent accesses cannot be prefetched).

11.4.2 Key Properties

1. **Memory Overhead:** The overall memory overhead is only a few MiB depending on configuration. The current implementation uses a small array of longs to encode watchpoint information, which is negligible.
2. **Performance Overhead:** KCSAN's runtime aims to be minimal, using an efficient watchpoint encoding that does not require acquiring any shared locks in the fast-path. For kernel boot on a system with 8 CPUs:
 - 5.0x slow-down with the default KCSAN config;
 - 2.8x slow-down from runtime fast-path overhead only (set very large `KCSAN_SKIP_WATCH` and unset `KCSAN_SKIP_WATCH_RANDOMIZE`).
3. **Annotation Overheads:** Minimal annotations are required outside the KCSAN runtime. As a result, maintenance overheads are minimal as the kernel evolves.
4. **Detects Racy Writes from Devices:** Due to checking data values upon setting up watchpoints, racy writes from devices can also be detected.
5. **Memory Ordering:** KCSAN is aware of only a subset of LKMM ordering rules; this may result in missed data races (false negatives).
6. **Analysis Accuracy:** For observed executions, due to using a sampling strategy, the analysis is *unsound* (false negatives possible), but aims to be complete (no false positives).

11.5 Alternatives Considered

An alternative data race detection approach for the kernel can be found in the [Kernel Thread Sanitizer \(KTSAN\)](#). KTSAN is a happens-before data race detector, which explicitly establishes the happens-before order between memory operations, which can then be used to determine data races as defined in [Data Races](#).

To build a correct happens-before relation, KTSAN must be aware of all ordering rules of the LKMM and synchronization primitives. Unfortunately, any omission leads to large numbers of false positives, which is especially detrimental in the context of the kernel which includes numerous custom synchronization mechanisms. To track the happens-before relation, KTSAN's implementation requires metadata for each memory location (shadow memory), which for each page corresponds to 4 pages of shadow memory, and can translate into overhead of tens of GiB on a large system.

KERNEL ELECTRIC-FENCE (KFENCE)

Kernel Electric-Fence (KFENCE) is a low-overhead sampling-based memory safety error detector. KFENCE detects heap out-of-bounds access, use-after-free, and invalid-free errors.

KFENCE is designed to be enabled in production kernels, and has near zero performance overhead. Compared to KASAN, KFENCE trades performance for precision. The main motivation behind KFENCE's design, is that with enough total uptime KFENCE will detect bugs in code paths not typically exercised by non-production test workloads. One way to quickly achieve a large enough total uptime is when the tool is deployed across a large fleet of machines.

12.1 Usage

To enable KFENCE, configure the kernel with:

```
CONFIG_KFENCE=y
```

To build a kernel with KFENCE support, but disabled by default (to enable, set kfence.sample_interval to non-zero value), configure the kernel with:

```
CONFIG_KFENCE=y
CONFIG_KFENCE_SAMPLE_INTERVAL=0
```

KFENCE provides several other configuration options to customize behaviour (see the respective help text in lib/Kconfig.kfence for more info).

12.1.1 Tuning performance

The most important parameter is KFENCE's sample interval, which can be set via the kernel boot parameter kfence.sample_interval in milliseconds. The sample interval determines the frequency with which heap allocations will be guarded by KFENCE. The default is configurable via the Kconfig option CONFIG_KFENCE_SAMPLE_INTERVAL. Setting kfence.sample_interval=0 disables KFENCE.

The sample interval controls a timer that sets up KFENCE allocations. By default, to keep the real sample interval predictable, the normal timer also causes CPU wake-ups when the system is completely idle. This may be undesirable on power-constrained systems. The boot parameter kfence.deferrable=1 instead switches to a "deferrable" timer which does not force CPU wake-ups on idle systems, at the risk of unpredictable sample intervals. The default is configurable via the Kconfig option CONFIG_KFENCE_DEFERRABLE.

Warning: The KUnit test suite is very likely to fail when using a deferrable timer since it currently causes very unpredictable sample intervals.

The KFENCE memory pool is of fixed size, and if the pool is exhausted, no further KFENCE allocations occur. With CONFIG_KFENCE_NUM_OBJECTS (default 255), the number of available guarded objects can be controlled. Each object requires 2 pages, one for the object itself and the other one used as a guard page; object pages are interleaved with guard pages, and every object page is therefore surrounded by two guard pages.

The total memory dedicated to the KFENCE memory pool can be computed as:

```
( #objects + 1 ) * 2 * PAGE_SIZE
```

Using the default config, and assuming a page size of 4 KiB, results in dedicating 2 MiB to the KFENCE memory pool.

Note: On architectures that support huge pages, KFENCE will ensure that the pool is using pages of size PAGE_SIZE. This will result in additional page tables being allocated.

12.1.2 Error reports

A typical out-of-bounds access looks like this:

```
=====
BUG: KFENCE: out-of-bounds read in test_out_of_bounds_read+0xa6/0x234

Out-of-bounds read at 0xfffff8c3f2e291fff (1B left of kfence-#72):
  test_out_of_bounds_read+0xa6/0x234
  kunit_try_run_case+0x61/0xa0
  kunit_generic_run_threadfn_adapter+0x16/0x30
  kthread+0x176/0x1b0
  ret_from_fork+0x22/0x30

kfence-#72: 0xfffff8c3f2e292000-0xfffff8c3f2e29201f, size=32, cache=kmalloc-32

allocated by task 484 on cpu 0 at 32.919330s:
  test_alloc+0xfe/0x738
  test_out_of_bounds_read+0x9b/0x234
  kunit_try_run_case+0x61/0xa0
  kunit_generic_run_threadfn_adapter+0x16/0x30
  kthread+0x176/0x1b0
  ret_from_fork+0x22/0x30

CPU: 0 PID: 484 Comm: kunit_try_catch Not tainted 5.13.0-rc3+ #7
Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS 1.14.0-2 04/01/2014
=====
```

The header of the report provides a short summary of the function involved in the access. It is followed by more detailed information about the access and its origin. Note that, real kernel addresses are only shown when using the kernel command line option no_hash_pointers.

Use-after-free accesses are reported as:

```
=====
BUG: KFENCE: use-after-free read in test_use_after_free_read+0xb3/0x143

Use-after-free read at 0xfffff8c3f2e2a0000 (in kfence-#79):
test_use_after_free_read+0xb3/0x143
kunit_try_run_case+0x61/0xa0
kunit_generic_run_threadfn_adapter+0x16/0x30
kthread+0x176/0x1b0
ret_from_fork+0x22/0x30

kfence-#79: 0xfffff8c3f2e2a0000-0xfffff8c3f2e2a001f, size=32, cache=kmalloc-32

allocated by task 488 on cpu 2 at 33.871326s:
test_alloc+0xfe/0x738
test_use_after_free_read+0x76/0x143
kunit_try_run_case+0x61/0xa0
kunit_generic_run_threadfn_adapter+0x16/0x30
kthread+0x176/0x1b0
ret_from_fork+0x22/0x30

freed by task 488 on cpu 2 at 33.871358s:
test_use_after_free_read+0xa8/0x143
kunit_try_run_case+0x61/0xa0
kunit_generic_run_threadfn_adapter+0x16/0x30
kthread+0x176/0x1b0
ret_from_fork+0x22/0x30

CPU: 2 PID: 488 Comm: kunit_try_catch Tainted: G      B      5.13.0-rc3+
      #7
Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS 1.14.0-2 04/01/2014
=====
```

KFENCE also reports on invalid frees, such as double-frees:

```
=====
BUG: KFENCE: invalid free in test_double_free+0xdc/0x171

Invalid free of 0xfffff8c3f2e2a4000 (in kfence-#81):
test_double_free+0xdc/0x171
kunit_try_run_case+0x61/0xa0
kunit_generic_run_threadfn_adapter+0x16/0x30
kthread+0x176/0x1b0
ret_from_fork+0x22/0x30

kfence-#81: 0xfffff8c3f2e2a4000-0xfffff8c3f2e2a401f, size=32, cache=kmalloc-32

allocated by task 490 on cpu 1 at 34.175321s:
test_alloc+0xfe/0x738
test_double_free+0x76/0x171
kunit_try_run_case+0x61/0xa0
=====
```

```

kunit_generic_run_threadfn_adapter+0x16/0x30
kthread+0x176/0x1b0
ret_from_fork+0x22/0x30

freed by task 490 on cpu 1 at 34.175348s:
test_double_free+0xa8/0x171
kunit_try_run_case+0x61/0xa0
kunit_generic_run_threadfn_adapter+0x16/0x30
kthread+0x176/0x1b0
ret_from_fork+0x22/0x30

CPU: 1 PID: 490 Comm: kunit_tryCatch Tainted: G     B      5.13.0-rc3+
#7
Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS 1.14.0-2 04/01/2014
=====
```

KFENCE also uses pattern-based redzones on the other side of an object's guard page, to detect out-of-bounds writes on the unprotected side of the object. These are reported on frees:

```

=====
BUG: KFENCE: memory corruption in test_kmalloc_aligned_oob_write+0xef/0x184

Corrupted memory at 0xfffff8c3f2e33aff9 [ 0xac . . . . . ] (in kfence-#156):
test_kmalloc_aligned_oob_write+0xef/0x184
kunit_try_run_case+0x61/0xa0
kunit_generic_run_threadfn_adapter+0x16/0x30
kthread+0x176/0x1b0
ret_from_fork+0x22/0x30

kfence-#156: 0xfffff8c3f2e33afb0-0xfffff8c3f2e33aff8, size=73, cache=kmalloc-96

allocated by task 502 on cpu 7 at 42.159302s:
test_alloc+0xfe/0x738
test_kmalloc_aligned_oob_write+0x57/0x184
kunit_try_run_case+0x61/0xa0
kunit_generic_run_threadfn_adapter+0x16/0x30
kthread+0x176/0x1b0
ret_from_fork+0x22/0x30

CPU: 7 PID: 502 Comm: kunit_tryCatch Tainted: G     B      5.13.0-rc3+
#7
Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS 1.14.0-2 04/01/2014
=====
```

For such errors, the address where the corruption occurred as well as the invalidly written bytes (offset from the address) are shown; in this representation, '.' denote untouched bytes. In the example above 0xac is the value written to the invalid address at offset 0, and the remaining '.' denote that no following bytes have been touched. Note that, real values are only shown if the kernel was booted with `no_hash_pointers`; to avoid information disclosure otherwise, '!' is used instead to denote invalidly written bytes.

And finally, KFENCE may also report on invalid accesses to any protected page where it was not possible to determine an associated object, e.g. if adjacent object pages had not yet been allocated:

```
=====
BUG: KFENCE: invalid read in test_invalid_access+0x26/0xe0

Invalid read at 0xfffffffffb670b00a:
test_invalid_access+0x26/0xe0
kunit_try_run_case+0x51/0x85
kunit_generic_run_threadfn_adapter+0x16/0x30
kthread+0x137/0x160
ret_from_fork+0x22/0x30

CPU: 4 PID: 124 Comm: kunit_try_catch Tainted: G      W      5.8.0-rc6+ #7
Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS 1.13.0-1 04/01/2014
=====
```

12.1.3 DebugFS interface

Some debugging information is exposed via debugfs:

- The file `/sys/kernel/debug/kfence/stats` provides runtime statistics.
- The file `/sys/kernel/debug/kfence/objects` provides a list of objects allocated via KFENCE, including those already freed but protected.

12.2 Implementation Details

Guarded allocations are set up based on the sample interval. After expiration of the sample interval, the next allocation through the main allocator (SLAB or SLUB) returns a guarded allocation from the KFENCE object pool (allocation sizes up to PAGE_SIZE are supported). At this point, the timer is reset, and the next allocation is set up after the expiration of the interval.

When using `CONFIG_KFENCE_STATIC_KEYS=y`, KFENCE allocations are "gated" through the main allocator's fast-path by relying on static branches via the static keys infrastructure. The static branch is toggled to redirect the allocation to KFENCE. Depending on sample interval, target workloads, and system architecture, this may perform better than the simple dynamic branch. Careful benchmarking is recommended.

KFENCE objects each reside on a dedicated page, at either the left or right page boundaries selected at random. The pages to the left and right of the object page are "guard pages", whose attributes are changed to a protected state, and cause page faults on any attempted access. Such page faults are then intercepted by KFENCE, which handles the fault gracefully by reporting an out-of-bounds access, and marking the page as accessible so that the faulting code can (wrongly) continue executing (set `panic_on_warn` to panic instead).

To detect out-of-bounds writes to memory within the object's page itself, KFENCE also uses pattern-based redzones. For each object page, a redzone is set up for all non-object memory. For typical alignments, the redzone is only required on the unguarded side of an object. Because KFENCE must honor the cache's requested alignment, special alignments may result in unprotected gaps on either side of an object, all of which are redzoned.

The following figure illustrates the page layout:

xxxxxxxxxx	0 :	xxxxxxxxxx	:	0	xxxxxxxxxx
xxxxxxxxxx	B :	xxxxxxxxxx	:	B	xxxxxxxxxx
x GUARD x	J : RED-	x GUARD x	RED-	J	x GUARD x
xxxxxxxxxx	E : ZONE	xxxxxxxxxx	ZONE	E	xxxxxxxxxx
xxxxxxxxxx	C :	xxxxxxxxxx	:	C	xxxxxxxxxx
xxxxxxxxxx	T :	xxxxxxxxxx	:	T	xxxxxxxxxx

Upon deallocation of a KFENCE object, the object's page is again protected and the object is marked as freed. Any further access to the object causes a fault and KFENCE reports a use-after-free access. Freed objects are inserted at the tail of KFENCE's freelist, so that the least recently freed objects are reused first, and the chances of detecting use-after-frees of recently freed objects is increased.

If pool utilization reaches 75% (default) or above, to reduce the risk of the pool eventually being fully occupied by allocated objects yet ensure diverse coverage of allocations, KFENCE limits currently covered allocations of the same source from further filling up the pool. The "source" of an allocation is based on its partial allocation stack trace. A side-effect is that this also limits frequent long-lived allocations (e.g. pagecache) of the same source filling up the pool permanently, which is the most common risk for the pool becoming full and the sampled allocation rate dropping to zero. The threshold at which to start limiting currently covered allocations can be configured via the boot parameter `kfence.skip_covered_thresh` (pool usage%).

12.3 Interface

The following describes the functions which are used by allocators as well as page handling code to set up and deal with KFENCE allocations.

bool `is_kfence_address`(const void *addr)
check if an address belongs to KFENCE pool

Parameters

const void *addr
address to check

Return

true or false depending on whether the address is within the KFENCE object range.

Description

KFENCE objects live in a separate page range and are not to be intermixed with regular heap objects (e.g. KFENCE objects must never be added to the allocator freelist). Failing to do so may and will result in heap corruptions, therefore `is_kfence_address()` must be used to check whether an object requires specific handling.

Note

This function may be used in fast-paths, and is performance critical. Future changes should take this into account; for instance, we want to avoid introducing another load and therefore

need to keep KFENCE_POOL_SIZE a constant (until immediate patching support is added to the kernel).

```
void kfence_shutdown_cache(struct kmem_cache *s)
    handle shutdown_cache() for KFENCE objects
```

Parameters

struct kmem_cache *s
cache being shut down

Description

Before shutting down a cache, one must ensure there are no remaining objects allocated from it. Because KFENCE objects are not referenced from the cache directly, we need to check them here.

Note that shutdown_cache() is internal to SL*B, and kmem_cache_destroy() does not return if allocated objects still exist: it prints an error message and simply aborts destruction of a cache, leaking memory.

If the only such objects are KFENCE objects, we will not leak the entire cache, but instead try to provide more useful debug info by making allocated objects "zombie allocations". Objects may then still be used or freed (which is handled gracefully), but usage will result in showing KFENCE error reports which include stack traces to the user of the object, the original allocation site, and caller to shutdown_cache().

```
void *kfence_alloc(struct kmem_cache *s, size_t size, gfp_t flags)
    allocate a KFENCE object with a low probability
```

Parameters

struct kmem_cache *s
struct kmem_cache with object requirements

size_t size
exact size of the object to allocate (can be less than **s->size** e.g. for kmalloc caches)

gfp_t flags
GFP flags

Return

- NULL - must proceed with allocating as usual,
- non-NUL - pointer to a KFENCE object.

Description

kfence_alloc() should be inserted into the heap allocation fast path, allowing it to transparently return KFENCE-allocated objects with a low probability using a static branch (the probability is controlled by the kfence.sample_interval boot parameter).

```
size_t kfence_ksize(const void *addr)
    get actual amount of memory allocated for a KFENCE object
```

Parameters

const void *addr
pointer to a heap object

Return

- 0 - not a KFENCE object, must call `_ksize()` instead,
- non-0 - this many bytes can be accessed without causing a memory error.

Description

`kfence_ksize()` returns the number of bytes requested for a KFENCE object at allocation time. This number may be less than the object size of the corresponding struct `kmem_cache`.

`void *kfence_object_start(const void *addr)`
 find the beginning of a KFENCE object

Parameters

`const void *addr`
 address within a KFENCE-allocated object

Return

address of the beginning of the object.

Description

SL[AU]B-allocated objects are laid out within a page one by one, so it is easy to calculate the beginning of an object given a pointer inside it and the object size. The same is not true for KFENCE, which places a single object at either end of the page. This helper function is used to find the beginning of a KFENCE-allocated object.

`void __kfence_free(void *addr)`
 release a KFENCE heap object to KFENCE pool

Parameters

`void *addr`
 object to be freed

Description

Requires: `is_kfence_address(addr)`

Release a KFENCE object and mark it as freed.

`bool kfence_free(void *addr)`
 try to release an arbitrary heap object to KFENCE pool

Parameters

`void *addr`
 object to be freed

Return

- false - object doesn't belong to KFENCE pool and was ignored,
- true - object was released to KFENCE pool.

Description

Release a KFENCE object and mark it as freed. May be called on any object, even non-KFENCE objects, to simplify integration of the hooks into the allocator's free codepath. The allocator must check the return value to determine if it was a KFENCE object or not.

```
bool kfence_handle_page_fault(unsigned long addr, bool is_write, struct pt_regs *regs)
    perform page fault handling for KFENCE pages
```

Parameters

unsigned long addr
faulting address

bool is_write
is access a write

struct pt_regs *regs
current struct pt_regs (can be NULL, but shows full stack trace)

Return

- false - address outside KFENCE pool,
- true - page fault handled by KFENCE, no additional handling required.

Description

A page fault inside KFENCE pool indicates a memory error, such as an out-of-bounds access, a use-after-free or an invalid memory access. In these cases KFENCE prints an error message and marks the offending page as present, so that the kernel can proceed.

12.4 Related Tools

In userspace, a similar approach is taken by [GWP-ASan](#). GWP-ASan also relies on guard pages and a sampling strategy to detect memory unsafety bugs at scale. KFENCE's design is directly influenced by GWP-ASan, and can be seen as its kernel sibling. Another similar but non-sampling approach, that also inspired the name "KFENCE", can be found in the userspace [Electric Fence Malloc Debugger](#).

In the kernel, several tools exist to debug memory access errors, and in particular KASAN can detect all bug classes that KFENCE can detect. While KASAN is more precise, relying on compiler instrumentation, this comes at a performance cost.

It is worth highlighting that KASAN and KFENCE are complementary, with different target environments. For instance, KASAN is the better debugging-aid, where test cases or reproducers exists: due to the lower chance to detect the error, it would require more effort using KFENCE to debug. Deployments at scale that cannot afford to enable KASAN, however, would benefit from using KFENCE to discover bugs due to code paths not exercised by test cases or fuzzers.

DEBUGGING KERNEL AND MODULES VIA GDB

The kernel debugger kgdb, hypervisors like QEMU or JTAG-based hardware interfaces allow to debug the Linux kernel and its modules during runtime using gdb. Gdb comes with a powerful scripting interface for python. The kernel provides a collection of helper scripts that can simplify typical kernel debugging steps. This is a short tutorial about how to enable and use them. It focuses on QEMU/KVM virtual machines as target, but the examples can be transferred to the other gdb stubs as well.

13.1 Requirements

- gdb 7.2+ (recommended: 7.4+) with python support enabled (typically true for distributions)

13.2 Setup

- Create a virtual Linux machine for QEMU/KVM (see www.linux-kvm.org and www.qemu.org for more details). For cross-development, <https://landley.net/aboriginal/bin> keeps a pool of machine images and toolchains that can be helpful to start from.
- Build the kernel with CONFIG_GDB_SCRIPTS enabled, but leave CONFIG_DEBUG_INFO_REDUCED off. If your architecture supports CONFIG_FRAME_POINTER, keep it enabled.
- Install that kernel on the guest, turn off KASLR if necessary by adding "nokaslr" to the kernel command line. Alternatively, QEMU allows to boot the kernel directly using -kernel, -append, -initrd command line switches. This is generally only useful if you do not depend on modules. See QEMU documentation for more details on this mode. In this case, you should build the kernel with CONFIG_RANDOMIZE_BASE disabled if the architecture supports KASLR.
- Build the gdb scripts (required on kernels v5.1 and above):

```
make scripts_gdb
```

- Enable the gdb stub of QEMU/KVM, either
 - at VM startup time by appending "-s" to the QEMU command line
- or
- during runtime by issuing "gdbserver" from the QEMU monitor console

- cd /path/to/linux-build
- Start gdb: gdb vmlinux

Note: Some distros may restrict auto-loading of gdb scripts to known safe directories. In case gdb reports to refuse loading vmlinux-gdb.py, add:

```
add-auto-load-safe-path /path/to/linux-build
```

to `~/.gdbinit`. See `gdb help` for more details.

- Attach to the booted guest:

```
(gdb) target remote :1234
```

13.3 Examples of using the Linux-provided gdb helpers

- Load module (and main kernel) symbols:

```
(gdb) lx-symbols
loading vmlinux
scanning for modules in /home/user/linux/build
loading @0xfffffffffa0020000: /home/user/linux/build/net/netfilter/xt_
    ↵tcpudp.ko
loading @0xfffffffffa0016000: /home/user/linux/build/net/netfilter/xt_
    ↵pkttype.ko
loading @0xfffffffffa0002000: /home/user/linux/build/net/netfilter/xt_limit.
    ↵ko
loading @0xfffffffffa00ca000: /home/user/linux/build/net/packet/af_packet.ko
loading @0xfffffffffa003c000: /home/user/linux/build/fs/fuse/fuse.ko
...
loading @0xfffffffffa0000000: /home/user/linux/build/drivers/ata/ata_
    ↵generic.ko
```

- Set a breakpoint on some not yet loaded module function, e.g.:

```
(gdb) b btrfs_init_sysfs
Function "btrfs_init_sysfs" not defined.
Make breakpoint pending on future shared library load? (y or [n]) y
Breakpoint 1 (btrfs_init_sysfs) pending.
```

- Continue the target:

```
(gdb) c
```

- Load the module on the target and watch the symbols being loaded as well as the breakpoint hit:

```
loading @0xfffffffffa0034000: /home/user/linux/build/lib/libcrc32c.ko
loading @0xfffffffffa0050000: /home/user/linux/build/lib/lzo/lzo_compress.ko
loading @0xfffffffffa006e000: /home/user/linux/build/lib/zlib_deflate/zlib_
    ↵deflate.ko
```

```
loading @0xfffffffffa01b1000: /home/user/linux/build/fs/btrfs/btrfs.ko
Breakpoint 1, btrfs_init_sysfs () at /home/user/linux/fs/btrfs/sysfs.c:36
36          btrfs_kset = kset_create_and_add("btrfs", NULL, fs_kobj);
```

- Dump the log buffer of the target kernel:

```
(gdb) lx-dmesg
[    0.000000] Initializing cgroup subsys cpuset
[    0.000000] Initializing cgroup subsys cpu
[    0.000000] Linux version 3.8.0-rc4-dbg+ (...
[    0.000000] Command line: root=/dev/sda2 resume=/dev/sda1 vga=0x314
[    0.000000] e820: BIOS-provided physical RAM map:
[    0.000000]   BIOS-e820: [mem 0x0000000000000000-0x000000000000fbff] u
  ↳ usable
[    0.000000]   BIOS-e820: [mem 0x000000000009fc00-0x000000000009ffff] u
  ↳ reserved
....
```

- Examine fields of the current task struct(supported by x86 and arm64 only):

```
(gdb) p $lx_current().pid
$1 = 4998
(gdb) p $lx_current().comm
$2 = "modprobe\000\000\000\000\000\000\000"
```

- Make use of the per-cpu function for the current or a specified CPU:

```
(gdb) p $lx_per_cpu("runqueues").nr_running
$3 = 1
(gdb) p $lx_per_cpu("runqueues", 2).nr_running
$4 = 0
```

- Dig into hrtimers using the container_of helper:

```
(gdb) set $next = $lx_per_cpu("hrtimer_bases").clock_base[0].active.next
(gdb) p *$container_of($next, "struct hrtimer", "node")
$5 = {
  node = {
    node = {
      __rb_parent_color = 18446612133355256072,
      rb_right = 0x0 <irq_stack_union>,
      rb_left = 0x0 <irq_stack_union>
    },
    expires = {
      tv64 = 1835268000000
    }
  },
  softexpires = {
    tv64 = 1835268000000
  },
  function = 0xffffffff81078232 <tick_sched_timer>,
```

```
base = 0xffff88003fd0d6f0,
state = 1,
start_pid = 0,
start_site = 0xffffffff81055c1f <hrtimer_start_range_ns+20>,
start_comm = "swapper/2\000\000\000\000\000\000"
}
```

13.4 List of commands and functions

The number of commands and convenience functions may evolve over the time, this is just a snapshot of the initial version:

```
(gdb) apropos lx
function lx_current -- Return current task
function lx_module -- Find module by name and return the module variable
function lx_per_cpu -- Return per-cpu variable
function lx_task_by_pid -- Find Linux task by PID and return the task_struct_variable
function lx_thread_info -- Calculate Linux thread_info from task variable
lx-dmesg -- Print Linux kernel log buffer
lx-lsmod -- List currently loaded modules
lx-symbols -- (Re-)load symbols of Linux kernel and currently loaded modules
```

Detailed help can be obtained via "help <command-name>" for commands and "help function <function-name>" for convenience functions.

USING KGDB, KDB AND THE KERNEL DEBUGGER INTERNALS

Author

Jason Wessel

14.1 Introduction

The kernel has two different debugger front ends (kdb and kgdb) which interface to the debug core. It is possible to use either of the debugger front ends and dynamically transition between them if you configure the kernel properly at compile and runtime.

Kdb is simplistic shell-style interface which you can use on a system console with a keyboard or serial console. You can use it to inspect memory, registers, process lists, dmesg, and even set breakpoints to stop in a certain location. Kdb is not a source level debugger, although you can set breakpoints and execute some basic kernel run control. Kdb is mainly aimed at doing some analysis to aid in development or diagnosing kernel problems. You can access some symbols by name in kernel built-ins or in kernel modules if the code was built with `CONFIG_KALLSYMS`.

Kgdb is intended to be used as a source level debugger for the Linux kernel. It is used along with gdb to debug a Linux kernel. The expectation is that gdb can be used to "break in" to the kernel to inspect memory, variables and look through call stack information similar to the way an application developer would use gdb to debug an application. It is possible to place breakpoints in kernel code and perform some limited execution stepping.

Two machines are required for using kgdb. One of these machines is a development machine and the other is the target machine. The kernel to be debugged runs on the target machine. The development machine runs an instance of gdb against the vmlinux file which contains the symbols (not a boot image such as bzImage, zImage, uImage...). In gdb the developer specifies the connection parameters and connects to kgdb. The type of connection a developer makes with gdb depends on the availability of kgdb I/O modules compiled as built-ins or loadable kernel modules in the test machine's kernel.

14.2 Compiling a kernel

- In order to enable compilation of kdb, you must first enable kgdb.
- The kgdb test compile options are described in the kgdb test suite chapter.

14.2.1 Kernel config options for kgdb

To enable CONFIG_KGDB you should look under *Kernel hacking → Kernel debugging* and select *KGDB: kernel debugger*.

While it is not a hard requirement that you have symbols in your vmlinux file, gdb tends not to be very useful without the symbolic data, so you will want to turn on CONFIG_DEBUG_INFO which is called *Compile the kernel with debug info* in the config menu.

It is advised, but not required, that you turn on the CONFIG_FRAME_POINTER kernel option which is called *Compile the kernel with frame pointers* in the config menu. This option inserts code into the compiled executable which saves the frame information in registers or on the stack at different points which allows a debugger such as gdb to more accurately construct stack back traces while debugging the kernel.

If the architecture that you are using supports the kernel option CONFIG_STRICT_KERNEL_RWX, you should consider turning it off. This option will prevent the use of software breakpoints because it marks certain regions of the kernel's memory space as read-only. If kgdb supports it for the architecture you are using, you can use hardware breakpoints if you desire to run with the CONFIG_STRICT_KERNEL_RWX option turned on, else you need to turn off this option.

Next you should choose one or more I/O drivers to interconnect debugging host and debugged target. Early boot debugging requires a KGDB I/O driver that supports early debugging and the driver must be built into the kernel directly. Kgdb I/O driver configuration takes place via kernel or module parameters which you can learn more about in the section that describes the parameter kgdboc.

Here is an example set of .config symbols to enable or disable for kgdb:

```
# CONFIG_STRICT_KERNEL_RWX is not set
CONFIG_FRAME_POINTER=y
CONFIG_KGDB=y
CONFIG_KGDB_SERIAL_CONSOLE=y
```

14.2.2 Kernel config options for kdb

Kdb is quite a bit more complex than the simple gdbstub sitting on top of the kernel's debug core. Kdb must implement a shell, and also adds some helper functions in other parts of the kernel, responsible for printing out interesting data such as what you would see if you ran lsmod, or ps. In order to build kdb into the kernel you follow the same steps as you would for kgdb.

The main config option for kdb is CONFIG_KGDB_KDB which is called *KGDB_KDB: include kdb frontend for kgdb* in the config menu. In theory you would have already also selected an I/O driver such as the CONFIG_KGDB_SERIAL_CONSOLE interface if you plan on using kdb on a serial port, when you were configuring kgdb.

If you want to use a PS/2-style keyboard with kdb, you would select `CONFIG_KDB_KEYBOARD` which is called *KGDB_KDB: keyboard as input device* in the config menu. The `CONFIG_KDB_KEYBOARD` option is not used for anything in the gdb interface to kgdb. The `CONFIG_KDB_KEYBOARD` option only works with kdb.

Here is an example set of `.config` symbols to enable/disable kdb:

```
# CONFIG_STRICT_KERNEL_RWX is not set
CONFIG_FRAME_POINTER=y
CONFIG_KGDB=y
CONFIG_KGDB_SERIAL_CONSOLE=y
CONFIG_KGDB_KDB=y
CONFIG_KDB_KEYBOARD=y
```

14.3 Kernel Debugger Boot Arguments

This section describes the various runtime kernel parameters that affect the configuration of the kernel debugger. The following chapter covers using kdb and kgdb as well as providing some examples of the configuration parameters.

14.3.1 Kernel parameter: `kgdboc`

The kgdboc driver was originally an abbreviation meant to stand for "kgdb over console". Today it is the primary mechanism to configure how to communicate from gdb to kgdb as well as the devices you want to use to interact with the kdb shell.

For kgdb/gdb, kgdboc is designed to work with a single serial port. It is intended to cover the circumstance where you want to use a serial console as your primary console as well as using it to perform kernel debugging. It is also possible to use kgdb on a serial port which is not designated as a system console. Kgdboc may be configured as a kernel built-in or a kernel loadable module. You can only make use of kgdbwait and early debugging if you build kgdboc into the kernel as a built-in.

Optionally you can elect to activate kms (Kernel Mode Setting) integration. When you use kms with kgdboc and you have a video driver that has atomic mode setting hooks, it is possible to enter the debugger on the graphics console. When the kernel execution is resumed, the previous graphics mode will be restored. This integration can serve as a useful tool to aid in diagnosing crashes or doing analysis of memory with kdb while allowing the full graphics console applications to run.

kgdboc arguments

Usage:

```
kgdboc=[kms][[,]kbd][[,]serial_device][,baud]
```

The order listed above must be observed if you use any of the optional configurations together.

Abbreviations:

- kms = Kernel Mode Setting

- kbd = Keyboard

You can configure kgdboc to use the keyboard, and/or a serial device depending on if you are using kdb and/or kgdb, in one of the following scenarios. The order listed above must be observed if you use any of the optional configurations together. Using kms + only gdb is generally not a useful combination.

Using loadable module or built-in

1. As a kernel built-in:

Use the kernel boot argument:

```
kgdboc=<tty-device>, [baud]
```

2. As a kernel loadable module:

Use the command:

```
modprobe kgdboc kgdboc=<tty-device>, [baud]
```

Here are two examples of how you might format the kgdboc string. The first is for an x86 target using the first serial port. The second example is for the ARM Versatile AB using the second serial port.

1. kgdboc=ttyS0,115200
2. kgdboc=ttyAMA1,115200

Configure kgdboc at runtime with sysfs

At run time you can enable or disable kgdboc by echoing a parameters into the sysfs. Here are two examples:

1. Enable kgdboc on ttyS0:

```
echo ttyS0 > /sys/module/kgdboc/parameters/kgdboc
```

2. Disable kgdboc:

```
echo "" > /sys/module/kgdboc/parameters/kgdboc
```

Note: You do not need to specify the baud if you are configuring the console on tty which is already configured or open.

More examples

You can configure kgdboc to use the keyboard, and/or a serial device depending on if you are using kdb and/or kgdb, in one of the following scenarios.

1. kdb and kgdb over only a serial port:

```
kgdboc=<serial_device>[,baud]
```

Example:

```
kgdboc=ttyS0,115200
```

2. kdb and kgdb with keyboard and a serial port:

```
kgdboc=kbd,<serial_device>[,baud]
```

Example:

```
kgdboc=kbd,ttyS0,115200
```

3. kdb with a keyboard:

```
kgdboc=kbd
```

4. kdb with kernel mode setting:

```
kgdboc=kms,kbd
```

5. kdb with kernel mode setting and kgdb over a serial port:

```
kgdboc=kms,kbd,ttyS0,115200
```

Note: Kgdboc does not support interrupting the target via the gdb remote protocol. You must manually send a SysRq-G unless you have a proxy that splits console output to a terminal program. A console proxy has a separate TCP port for the debugger and a separate TCP port for the "human" console. The proxy can take care of sending the SysRq-G for you.

When using kgdboc with no debugger proxy, you can end up connecting the debugger at one of two entry points. If an exception occurs after you have loaded kgdboc, a message should print on the console stating it is waiting for the debugger. In this case you disconnect your terminal program and then connect the debugger in its place. If you want to interrupt the target system and forcibly enter a debug session you have to issue a Sysrq sequence and then type the letter g. Then you disconnect the terminal session and connect gdb. Your options if you don't like this are to hack gdb to send the SysRq-G for you as well as on the initial connect, or to use a debugger proxy that allows an unmodified gdb to do the debugging.

14.3.2 Kernel parameter: kgdboc_earlycon

If you specify the kernel parameter `kgdboc_earlycon` and your serial driver registers a boot console that supports polling (doesn't need interrupts and implements a nonblocking read() function) kgdb will attempt to work using the boot console until it can transition to the regular tty driver specified by the `kgdboc` parameter.

Normally there is only one boot console (especially that implements the read() function) so just adding `kgdboc_earlycon` on its own is sufficient to make this work. If you have more than one boot console you can add the boot console's name to differentiate. Note that names that are registered through the boot console layer and the tty layer are not the same for the same port.

For instance, on one board to be explicit you might do:

```
kgdboc_earlycon=qcom_geni kgdboc=ttyMSM0
```

If the only boot console on the device was "qcom_geni", you could simplify:

```
kgdboc_earlycon kgdboc=ttyMSM0
```

14.3.3 Kernel parameter: kgdbwait

The Kernel command line option `kgdbwait` makes kgdb wait for a debugger connection during booting of a kernel. You can only use this option if you compiled a kgdb I/O driver into the kernel and you specified the I/O driver configuration as a kernel command line option. The `kgdbwait` parameter should always follow the configuration parameter for the kgdb I/O driver in the kernel command line else the I/O driver will not be configured prior to asking the kernel to use it to wait.

The kernel will stop and wait as early as the I/O driver and architecture allows when you use this option. If you build the kgdb I/O driver as a loadable kernel module `kgdbwait` will not do anything.

14.3.4 Kernel parameter: kgdbcon

The `kgdbcon` feature allows you to see `printk()` messages inside gdb while gdb is connected to the kernel. Kdb does not make use of the `kgdbcon` feature.

Kgdb supports using the gdb serial protocol to send console messages to the debugger when the debugger is connected and running. There are two ways to activate this feature.

1. Activate with the kernel command line option:

```
kgdbcon
```

2. Use sysfs before configuring an I/O driver:

```
echo 1 > /sys/module/kgdb/parameters/kgdb_use_con
```

Note: If you do this after you configure the kgdb I/O driver, the setting will not take effect until the next point the I/O is reconfigured.

Important: You cannot use kgdboc + kgdbcon on a tty that is an active system console. An example of incorrect usage is:

```
console=ttyS0,115200 kgdboc=ttyS0 kgdbcon
```

It is possible to use this option with kgdboc on a tty that is not a system console.

14.3.5 Run time parameter: kgdbreboot

The kgdbreboot feature allows you to change how the debugger deals with the reboot notification. You have 3 choices for the behavior. The default behavior is always set to 0.

1	echo -1 > /sys/module/debug_core/parameters/kgdbreboot	Ignore the reboot notification entirely.
2	echo 0 > /sys/module/debug_core/parameters/kgdbreboot	Send the detach message to any attached debugger client.
3	echo 1 > /sys/module/debug_core/parameters/kgdbreboot	Enter the debugger on reboot notify.

14.3.6 Kernel parameter: nokaslr

If the architecture that you are using enable KASLR by default, you should consider turning it off. KASLR randomizes the virtual address where the kernel image is mapped and confuse gdb which resolve kernel symbol address from symbol table of vmlinux.

14.4 Using kdb

14.4.1 Quick start for kdb on a serial port

This is a quick example of how to use kdb.

1. Configure kgdboc at boot using kernel parameters:

```
console=ttyS0,115200 kgdboc=ttyS0,115200 nokaslr
```

OR

Configure kgdboc after the kernel has booted; assuming you are using a serial port console:

```
echo ttyS0 > /sys/module/kgdboc/parameters/kgdboc
```

2. Enter the kernel debugger manually or by waiting for an oops or fault. There are several ways you can enter the kernel debugger manually; all involve using the SysRq-G, which means you must have enabled CONFIG_MAGIC_SYSRQ=y in your kernel config.

- When logged in as root or with a super user session you can run:

```
echo g > /proc/sysrq-trigger
```

- Example using minicom 2.2

Press: CTRL-A f g

- When you have telneted to a terminal server that supports sending a remote break

Press: CTRL-]

Type in: send break

Press: Enter g

3. From the kdb prompt you can run the help command to see a complete list of the commands that are available.

Some useful commands in kdb include:

lsmod	Shows where kernel modules are loaded
ps	Displays only the active processes
ps A	Shows all the processes
summary	Shows kernel version info and memory usage
bt	Get a backtrace of the current process using dump_stack()
dmesg	View the kernel syslog buffer
go	Continue the system

4. When you are done using kdb you need to consider rebooting the system or using the go command to resuming normal kernel execution. If you have paused the kernel for a lengthy period of time, applications that rely on timely networking or anything to do with real wall clock time could be adversely affected, so you should take this into consideration when using the kernel debugger.

14.4.2 Quick start for kdb using a keyboard connected console

This is a quick example of how to use kdb with a keyboard.

1. Configure kgdboc at boot using kernel parameters:

```
kgdboc=kbd
```

OR

Configure kgdboc after the kernel has booted:

```
echo kbd > /sys/module/kgdboc/parameters/kgdboc
```

2. Enter the kernel debugger manually or by waiting for an oops or fault. There are several ways you can enter the kernel debugger manually; all involve using the SysRq-G, which means you must have enabled CONFIG_MAGIC_SYSRQ=y in your kernel config.

- When logged in as root or with a super user session you can run:

```
echo g > /proc/sysrq-trigger
```

- Example using a laptop keyboard:

Press and hold down: Alt

Press and hold down: Fn

Press and release the key with the label: SysRq

Release: Fn

Press and release: g

Release: Alt

- Example using a PS/2 101-key keyboard

Press and hold down: Alt

Press and release the key with the label: SysRq

Press and release: g

Release: Alt

3. Now type in a kdb command such as help, dmesg, bt or go to continue kernel execution.

14.5 Using kgdb / gdb

In order to use kgdb you must activate it by passing configuration information to one of the kgdb I/O drivers. If you do not pass any configuration information kgdb will not do anything at all. Kgdb will only actively hook up to the kernel trap hooks if a kgdb I/O driver is loaded and configured. If you unconfigure a kgdb I/O driver, kgdb will unregister all the kernel hook points.

All kgdb I/O drivers can be reconfigured at run time, if CONFIG_SYSFS and CONFIG_MODULES are enabled, by echo'ing a new config string to /sys/module/<driver>/parameter/<option>. The driver can be unconfigured by passing an empty string. You cannot change the configuration while the debugger is attached. Make sure to detach the debugger with the detach command prior to trying to unconfigure a kgdb I/O driver.

14.5.1 Connecting with gdb to a serial port

1. Configure kgdboc

Configure kgdboc at boot using kernel parameters:

```
kgdboc=ttyS0,115200
```

OR

Configure kgdboc after the kernel has booted:

```
echo ttyS0 > /sys/module/kgdboc/parameters/kgdboc
```

2. Stop kernel execution (break into the debugger)

In order to connect to gdb via kgdboc, the kernel must first be stopped. There are several ways to stop the kernel which include using kgdbwait as a boot argument, via a SysRq-G, or running the kernel until it takes an exception where it waits for the debugger to attach.

- When logged in as root or with a super user session you can run:

```
echo g > /proc/sysrq-trigger
```

- Example using minicom 2.2

Press: CTRL-A f g

- When you have telneted to a terminal server that supports sending a remote break

Press: CTRL-]

Type in: send break

Press: Enter g

3. Connect from gdb

Example (using a directly connected port):

```
% gdb ./vmlinux
(gdb) set serial baud 115200
(gdb) target remote /dev/ttyS0
```

Example (kgdb to a terminal server on TCP port 2012):

```
% gdb ./vmlinux
(gdb) target remote 192.168.2.2:2012
```

Once connected, you can debug a kernel the way you would debug an application program.

If you are having problems connecting or something is going seriously wrong while debugging, it will most often be the case that you want to enable gdb to be verbose about its target communications. You do this prior to issuing the `target remote` command by typing in:

```
set debug remote 1
```

Remember if you continue in gdb, and need to "break in" again, you need to issue an other SysRq-G. It is easy to create a simple entry point by putting a breakpoint at `sys_sync` and then you can run `sync` from a shell or script to break into the debugger.

14.6 kgdb and kdb interoperability

It is possible to transition between kdb and kgdb dynamically. The debug core will remember which you used the last time and automatically start in the same mode.

14.6.1 Switching between kdb and kgdb

Switching from kgdb to kdb

There are two ways to switch from kgdb to kdb: you can use gdb to issue a maintenance packet, or you can blindly type the command \$3#33. Whenever the kernel debugger stops in kgdb mode it will print the message KGDB or \$3#33 for KDB. It is important to note that you have to type the sequence correctly in one pass. You cannot type a backspace or delete because kgdb will interpret that as part of the debug stream.

1. Change from kgdb to kdb by blindly typing:

```
$3#33
```

2. Change from kgdb to kdb with gdb:

```
maintenance packet 3
```

Note: Now you must kill gdb. Typically you press CTRL-Z and issue the command:

```
kill -9 %
```

Change from kdb to kgdb

There are two ways you can change from kdb to kgdb. You can manually enter kgdb mode by issuing the kgdb command from the kdb shell prompt, or you can connect gdb while the kdb shell prompt is active. The kdb shell looks for the typical first commands that gdb would issue with the gdb remote protocol and if it sees one of those commands it automatically changes into kgdb mode.

1. From kdb issue the command:

```
kgdb
```

Now disconnect your terminal program and connect gdb in its place

2. At the kdb prompt, disconnect the terminal program and connect gdb in its place.

14.6.2 Running kdb commands from gdb

It is possible to run a limited set of kdb commands from gdb, using the gdb monitor command. You don't want to execute any of the run control or breakpoint operations, because it can disrupt the state of the kernel debugger. You should be using gdb for breakpoints and run control operations if you have gdb connected. The more useful commands to run are things like lsmod, dmesg, ps or possibly some of the memory information commands. To see all the kdb commands you can run `monitor help`.

Example:

```
(gdb) monitor ps
1 idle process (state I) and
27 sleeping system daemon (state M) processes suppressed,
use 'ps A' to see all.
Task Addr      Pid  Parent [*] cpu State Thread      Command
0xc78291d0      1      0  0      0    S  0xc7829404  init
0xc7954150     942      1  0      0    S  0xc7954384  dropbear
0xc78789c0     944      1  0      0    S  0xc7878bf4  sh
(gdb)
```

14.7 kgdb Test Suite

When kgdb is enabled in the kernel config you can also elect to enable the config parameter `KGDB_TESTS`. Turning this on will enable a special kgdb I/O module which is designed to test the kgdb internal functions.

The kgdb tests are mainly intended for developers to test the kgdb internals as well as a tool for developing a new kgdb architecture specific implementation. These tests are not really for end users of the Linux kernel. The primary source of documentation would be to look in the `drivers/misc/kgdbts.c` file.

The kgdb test suite can also be configured at compile time to run the core set of tests by setting the kernel config parameter `KGDB_TESTS_ON_BOOT`. This particular option is aimed at automated regression testing and does not require modifying the kernel boot config arguments. If this is turned on, the kgdb test suite can be disabled by specifying `kgdbts=` as a kernel boot argument.

14.8 Kernel Debugger Internals

14.8.1 Architecture Specifics

The kernel debugger is organized into a number of components:

1. The debug core

The debug core is found in `kernel/debugger/debug_core.c`. It contains:

- A generic OS exception handler which includes sync'ing the processors into a stopped state on an multi-CPU system.
- The API to talk to the kgdb I/O drivers
- The API to make calls to the arch-specific kgdb implementation
- The logic to perform safe memory reads and writes to memory while using the debugger
- A full implementation for software breakpoints unless overridden by the arch
- The API to invoke either the kdb or kgdb frontend to the debug core.
- The structures and callback API for atomic kernel mode setting.

Note: kgdboc is where the kms callbacks are invoked.

2. kgdb arch-specific implementation

This implementation is generally found in `arch/*/kernel/kgdb.c`. As an example, `arch/x86/kernel/kgdb.c` contains the specifics to implement HW breakpoint as well as the initialization to dynamically register and unregister for the trap handlers on this architecture. The arch-specific portion implements:

- contains an arch-specific trap catcher which invokes `kgdb_handle_exception()` to start kgdb about doing its work
- translation to and from gdb specific packet format to struct `pt_regs`
- Registration and unregistration of architecture specific trap hooks
- Any special exception handling and cleanup
- NMI exception handling and cleanup
- (optional) HW breakpoints

3. gdbstub frontend (aka kgdb)

The gdbstub is located in `kernel/debug/gdbstub.c`. It contains:

- All the logic to implement the gdb serial protocol

4. kdb frontend

The kdb debugger shell is broken down into a number of components. The kdb core is located in `kernel/debug/kdb`. There are a number of helper functions in some of the other kernel components to make it possible for kdb to examine and report information about the kernel without taking locks that could cause a kernel deadlock. The kdb core contains implements the following functionality.

- A simple shell
- The kdb core command set
- A registration API to register additional kdb shell commands.
 - A good example of a self-contained kdb module is the `ftrace` command for dumping the ftrace buffer. See: `kernel/trace/trace_kdb.c`
 - For an example of how to dynamically register a new kdb command you can build the `kdb_hello.ko` kernel module from `samples/kdb/kdb_hello.c`. To build this example you can set `CONFIG_SAMPLES=y` and `CONFIG_SAMPLE_KDB=m` in your kernel config. Later run `modprobe kdb_hello` and the next time you enter the kdb shell, you can run the `hello` command.
- The implementation for `kdb_printf()` which emits messages directly to I/O drivers, bypassing the kernel log.
- SW / HW breakpoint management for the kdb shell

5. kgdb I/O driver

Each kgdb I/O driver has to provide an implementation for the following:

- configuration via built-in or module

- dynamic configuration and kgdb hook registration calls
- read and write character interface
- A cleanup handler for unconfiguring from the kgdb core
- (optional) Early debug methodology

Any given kgdb I/O driver has to operate very closely with the hardware and must do it in such a way that does not enable interrupts or change other parts of the system context without completely restoring them. The kgdb core will repeatedly "poll" a kgdb I/O driver for characters when it needs input. The I/O driver is expected to return immediately if there is no data available. Doing so allows for the future possibility to touch watchdog hardware in such a way as to have a target system not reset when these are enabled.

If you are intent on adding kgdb architecture specific support for a new architecture, the architecture should define HAVE_ARCH_KGDB in the architecture specific Kconfig file. This will enable kgdb for the architecture, and at that point you must create an architecture specific kgdb implementation.

There are a few flags which must be set on every architecture in their `asm/kgdb.h` file. These are:

- **NUMREGBYTES:**
The size in bytes of all of the registers, so that we can ensure they will all fit into a packet.
- **BUFSIZE:**
The size in bytes of the buffer GDB will read into. This must be larger than NUMREGBYTES.
- **CACHE_FLUSH_IS_SAFE:**
Set to 1 if it is always safe to call `flush_cache_range` or `flush_icache_range`. On some architectures, these functions may not be safe to call on SMP since we keep other CPUs in a holding pattern.

There are also the following functions for the common backend, found in `kernel/kgdb.c`, that must be supplied by the architecture-specific backend unless marked as (optional), in which case a default function maybe used if the architecture does not need to provide a specific implementation.

`int kgdb_skipexception(int exception, struct pt_regs *regs)`

(optional) exit kgdb_handle_exception early

Parameters

`int exception`

Exception vector number

`struct pt_regs *regs`

Current `struct pt_regs`.

On some architectures it is required to skip a breakpoint exception when it occurs after a breakpoint has been removed. This can be implemented in the architecture specific portion of kgdb.

`void kgdb_breakpoint(void)`

compiled in breakpoint

Parameters**void**

no arguments

Description

This will be implemented as a static inline per architecture. This function is called by the kgdb core to execute an architecture specific trap to cause kgdb to enter the exception processing.

int `kgdb_arch_init`(void)

Perform any architecture specific initialization.

Parameters**void**

no arguments

Description

This function will handle the initialization of any architecture specific callbacks.

void `kgdb_arch_exit`(void)

Perform any architecture specific uninitialization.

Parameters**void**

no arguments

Description

This function will handle the uninitialization of any architecture specific callbacks, for dynamic registration and unregistration.

void `pt_regs_to_gdb_regs`(unsigned long *gdb_regs, struct pt_regs *regs)

Convert ptrace regs to GDB regs

Parameters**unsigned long *gdb_regs**

A pointer to hold the registers in the order GDB wants.

struct pt_regs *regs

The struct pt_regs of the current process.

Convert the pt_regs in **regs** into the format for registers that GDB expects, stored in **gdb_regs**.

void `sleeping_thread_to_gdb_regs`(unsigned long *gdb_regs, struct task_struct *p)

Convert ptrace regs to GDB regs

Parameters**unsigned long *gdb_regs**

A pointer to hold the registers in the order GDB wants.

struct task_struct *p

The struct task_struct of the desired process.

Convert the register values of the sleeping process in **p** to the format that GDB expects. This function is called when kgdb does not have access to the struct `pt_regs` and therefore it should fill the gdb registers `gdb_regs` with what has been saved in struct `thread_struct` thread field during switch_to.

void **gdb_regs_to_pt_regs**(unsigned long *gdb_regs, struct pt_regs *regs)

 Convert GDB regs to ptrace regs.

Parameters

unsigned long *gdb_regs

 A pointer to hold the registers we've received from GDB.

struct pt_regs *regs

 A pointer to a struct `pt_regs` to hold these values in.

 Convert the GDB regs in `gdb_regs` into the `pt_regs`, and store them in `regs`.

int **kgdb_arch_handle_exception**(int vector, int signo, int err_code, char *remcom_in_buffer, char *remcom_out_buffer, struct pt_regs *regs)

Handle architecture specific GDB packets.

Parameters

int vector

 The error vector of the exception that happened.

int signo

 The signal number of the exception that happened.

int err_code

 The error code of the exception that happened.

char *remcom_in_buffer

 The buffer of the packet we have read.

char *remcom_out_buffer

 The buffer of BUFMAX bytes to write a packet into.

struct pt_regs *regs

 The struct `pt_regs` of the current process.

This function MUST handle the 'c' and 's' command packets, as well packets to set / remove a hardware breakpoint, if used. If there are additional packets which the hardware needs to handle, they are handled here. The code should return -1 if it wants to process more packets, and a 0 or 1 if it wants to exit from the kgdb callback.

void **kgdb_arch_handle_qxfer_pkt**(char *remcom_in_buffer, char *remcom_out_buffer)

Handle architecture specific GDB XML packets.

Parameters

char *remcom_in_buffer

 The buffer of the packet we have read.

char *remcom_out_buffer

 The buffer of BUFMAX bytes to write a packet into.

void kgdb_call_nmi_hook(void *ignored)
Call kgdb_nmallback() on the current CPU

Parameters**void *ignored**

This parameter is only here to match the prototype.

If you're using the default implementation of `kgdb_roundup_cpus()` this function will be called per CPU. If you don't implement `kgdb_call_nmi_hook()` a default will be used.

void kgdb_roundup_cpus(void)
Get other CPUs into a holding pattern

Parameters**void**

no arguments

Description

On SMP systems, we need to get the attention of the other CPUs and get them into a known state. This should do what is needed to get the other CPUs to call `kgdb_wait()`. Note that on some arches, the NMI approach is not used for rounding up all the CPUs. Normally those architectures can just not implement this and get the default.

On non-SMP systems, this is not called.

void kgdb_arch_set_pc(struct pt_regs *regs, unsigned long pc)
Generic call back to the program counter

Parameters**struct pt_regs *regs**

Current struct `pt_regs`.

unsigned long pc

The new value for the program counter

This function handles updating the program counter and requires an architecture specific implementation.

void kgdb_arch_late(void)
Perform any architecture specific initialization.

Parameters**void**

no arguments

Description

This function will handle the late initialization of any architecture specific callbacks. This is an optional function for handling things like late initialization of hw breakpoints. The default implementation does nothing.

struct kgdb_arch
Describe architecture specific values.

Definition:

```
struct kgdb_arch {
    unsigned char          gdb_bpt_instr[BREAK_INSTR_SIZE];
    unsigned long           flags;
    int (*set_breakpoint)(unsigned long, char *);
    int (*remove_breakpoint)(unsigned long, char *);
    int (*set_hw_breakpoint)(unsigned long, int, enum kgdb_bptype);
    int (*remove_hw_breakpoint)(unsigned long, int, enum kgdb_bptype);
    void (*disable_hw_break)(struct pt_regs *regs);
    void (*remove_all_hw_break)(void);
    void (*correct_hw_break)(void);
    void (*enable_nmi)(bool on);
};
```

Members**`gdb_bpt_instr`**

The instruction to trigger a breakpoint.

`flags`

Flags for the breakpoint, currently just `KGDB_HW_BREAKPOINT`.

`set_breakpoint`

Allow an architecture to specify how to set a software breakpoint.

`remove_breakpoint`

Allow an architecture to specify how to remove a software breakpoint.

`set_hw_breakpoint`

Allow an architecture to specify how to set a hardware breakpoint.

`remove_hw_breakpoint`

Allow an architecture to specify how to remove a hardware breakpoint.

`disable_hw_break`

Allow an architecture to specify how to disable hardware breakpoints for a single cpu.

`remove_all_hw_break`

Allow an architecture to specify how to remove all hardware breakpoints.

`correct_hw_break`

Allow an architecture to specify how to correct the hardware debug registers.

`enable_nmi`

Manage NMI-triggered entry to KGDB

`struct kgdb_io`

Describe the interface for an I/O driver to talk with KGDB.

Definition:

```
struct kgdb_io {
    const char          *name;
    int (*read_char) (void);
    void (*write_char) (u8);
    void (*flush) (void);
    int (*init) (void);
```

```

void (*deinit) (void);
void (*pre_exception) (void);
void (*post_exception) (void);
struct console      *cons;
};

```

Members

name

Name of the I/O driver.

read_char

Pointer to a function that will return one char.

write_char

Pointer to a function that will write one char.

flush

Pointer to a function that will flush any pending writes.

init

Pointer to a function that will initialize the device.

deinit

Pointer to a function that will deinit the device. Implies that this I/O driver is temporary and expects to be replaced. Called when an I/O driver is replaced or explicitly unregistered.

pre_exception

Pointer to a function that will do any prep work for the I/O driver.

post_exception

Pointer to a function that will do any cleanup work for the I/O driver.

cons

valid if the I/O device is a console; else NULL.

14.8.2 kgdboc internals

kgdboc and uarts

The kgdboc driver is actually a very thin driver that relies on the underlying low level to the hardware driver having "polling hooks" to which the tty driver is attached. In the initial implementation of kgdboc the serial_core was changed to expose a low level UART hook for doing polled mode reading and writing of a single character while in an atomic context. When kgdb makes an I/O request to the debugger, kgdboc invokes a callback in the serial core which in turn uses the callback in the UART driver.

When using kgdboc with a UART, the UART driver must implement two callbacks in the struct uart_ops. Example from drivers/8250.c:

```

#ifndef CONFIG_CONSOLE_POLL
    .poll_get_char = serial8250_get_poll_char,
    .poll_put_char = serial8250_put_poll_char,
#endif

```

Any implementation specifics around creating a polling driver use the `#ifdef CONFIG_CONSOLE_POLL`, as shown above. Keep in mind that polling hooks have to be implemented in such a way that they can be called from an atomic context and have to restore the state of the UART chip on return such that the system can return to normal when the debugger detaches. You need to be very careful with any kind of lock you consider, because failing here is most likely going to mean pressing the reset button.

kgdboc and keyboards

The kgdboc driver contains logic to configure communications with an attached keyboard. The keyboard infrastructure is only compiled into the kernel when `CONFIG_KDB_KEYBOARD=y` is set in the kernel configuration.

The core polled keyboard driver for PS/2 type keyboards is in `drivers/char/kdb_keyboard.c`. This driver is hooked into the debug core when kgdboc populates the callback in the array called `kdb_poll_funcs[]`. The `kdb_get_kbd_char()` is the top-level function which polls hardware for single character input.

kgdboc and kms

The kgdboc driver contains logic to request the graphics display to switch to a text context when you are using `kgdboc=kms,kbd`, provided that you have a video driver which has a frame buffer console and atomic kernel mode setting support.

Every time the kernel debugger is entered it calls `kgdboc_pre_exp_handler()` which in turn calls `con_debug_enter()` in the virtual console layer. On resuming kernel execution, the kernel debugger calls `kgdboc_post_exp_handler()` which in turn calls `con_debug_leave()`.

Any video driver that wants to be compatible with the kernel debugger and the atomic kms callbacks must implement the `mode_set_base_atomic`, `fb_debug_enter` and `fb_debug_leave` operations. For the `fb_debug_enter` and `fb_debug_leave` the option exists to use the generic drm fb helper functions or implement something custom for the hardware. The following example shows the initialization of the `.mode_set_base_atomic` operation in `drivers/gpu/drm/i915/intel_display.c`:

```
static const struct drm_crtc_helper_funcs intel_helper_funcs = {
[...]
    .mode_set_base_atomic = intel_pipe_set_base_atomic,
[...]
};
```

Here is an example of how the i915 driver initializes the `fb_debug_enter` and `fb_debug_leave` functions to use the generic drm helpers in `drivers/gpu/drm/i915/intel_fb.c`:

```
static struct fb_ops intelfb_ops = {
[...]
    .fb_debug_enter = drm_fb_helper_debug_enter,
    .fb_debug_leave = drm_fb_helper_debug_leave,
[...]
};
```

14.9 Credits

The following people have contributed to this document:

1. Amit Kale <amitkale@linsyssoft.com>
2. Tom Rini <trini@kernel.crashing.org>

In March 2008 this document was completely rewritten by:

- Jason Wessel <jason.wessel@windriver.com>

In Jan 2010 this document was updated to include kdb.

- Jason Wessel <jason.wessel@windriver.com>

LINUX KERNEL SELFTESTS

The kernel contains a set of "self tests" under the tools/testing/selftests/ directory. These are intended to be small tests to exercise individual code paths in the kernel. Tests are intended to be run after building, installing and booting a kernel.

Kselftest from mainline can be run on older stable kernels. Running tests from mainline offers the best coverage. Several test rings run mainline kselftest suite on stable releases. The reason is that when a new test gets added to test existing code to regression test a bug, we should be able to run that test on an older kernel. Hence, it is important to keep code that can still test an older kernel and make sure it skips the test gracefully on newer releases.

You can find additional information on Kselftest framework, how to write new tests using the framework on Kselftest wiki:

<https://kselftest.wiki.kernel.org/>

On some systems, hot-plug tests could hang forever waiting for cpu and memory to be ready to be offlined. A special hot-plug target is created to run the full range of hot-plug tests. In default mode, hot-plug tests run in safe mode with a limited scope. In limited mode, cpu-hotplug test is run on a single cpu as opposed to all hotplug capable cpus, and memory hotplug test is run on 2% of hotplug capable memory instead of 10%.

kselftest runs as a userspace process. Tests that can be written/run in userspace may wish to use the *Test Harness*. Tests that need to be run in kernel space may wish to use a *Test Module*.

15.1 Running the selftests (hotplug tests are run in limited mode)

To build the tests:

```
$ make headers  
$ make -C tools/testing/selftests
```

To run the tests:

```
$ make -C tools/testing/selftests run_tests
```

To build and run the tests with a single command, use:

```
$ make kselftest
```

Note that some tests will require root privileges.

Kselftest supports saving output files in a separate directory and then running tests. To locate output files in a separate directory two syntaxes are supported. In both cases the working directory must be the root of the kernel src. This is applicable to "Running a subset of selftests" section below.

To build, save output files in a separate directory with O=

```
$ make O=/tmp/kselftest kselftest
```

To build, save output files in a separate directory with KBUILD_OUTPUT

```
$ export KBUILD_OUTPUT=/tmp/kselftest; make kselftest
```

The O= assignment takes precedence over the KBUILD_OUTPUT environment variable.

The above commands by default run the tests and print full pass/fail report. Kselftest supports "summary" option to make it easier to understand the test results. Please find the detailed individual test results for each test in /tmp/testname file(s) when summary option is specified. This is applicable to "Running a subset of selftests" section below.

To run kselftest with summary option enabled

```
$ make summary=1 kselftest
```

15.2 Running a subset of selftests

You can use the "TARGETS" variable on the make command line to specify single test to run, or a list of tests to run.

To run only tests targeted for a single subsystem:

```
$ make -C tools/testing/selftests TARGETS=ptrace run_tests
```

You can specify multiple tests to build and run:

```
$ make TARGETS="size timers" kselftest
```

To build, save output files in a separate directory with O=

```
$ make O=/tmp/kselftest TARGETS="size timers" kselftest
```

To build, save output files in a separate directory with KBUILD_OUTPUT

```
$ export KBUILD_OUTPUT=/tmp/kselftest; make TARGETS="size timers" kselftest
```

Additionally you can use the "SKIP_TARGETS" variable on the make command line to specify one or more targets to exclude from the TARGETS list.

To run all tests but a single subsystem:

```
$ make -C tools/testing/selftests SKIP_TARGETS=ptrace run_tests
```

You can specify multiple tests to skip:

```
$ make SKIP_TARGETS="size timers" kselftest
```

You can also specify a restricted list of tests to run together with a dedicated skipplist:

```
$ make TARGETS="breakpoints size timers" SKIP_TARGETS=size kselftest
```

See the top-level tools/testing/selftests/Makefile for the list of all possible targets.

15.3 Running the full range hotplug selftests

To build the hotplug tests:

```
$ make -C tools/testing/selftests hotplug
```

To run the hotplug tests:

```
$ make -C tools/testing/selftests run_hotplug
```

Note that some tests will require root privileges.

15.4 Install selftests

You can use the "install" target of "make" (which calls the *kselftest_install.sh* tool) to install selftests in the default location (*tools/testing/selftests/kselftest_install*), or in a user specified location via the *INSTALL_PATH* "make" variable.

To install selftests in default location:

```
$ make -C tools/testing/selftests install
```

To install selftests in a user specified location:

```
$ make -C tools/testing/selftests install INSTALL_PATH=/some/other/path
```

15.5 Running installed selftests

Found in the install directory, as well as in the Kselftest tarball, is a script named *run_kselftest.sh* to run the tests.

You can simply do the following to run the installed Kselftests. Please note some tests will require root privileges:

```
$ cd kselftest_install
$ ./run_kselftest.sh
```

To see the list of available tests, the *-l* option can be used:

```
$ ./run_kselftest.sh -l
```

The `-c` option can be used to run all the tests from a test collection, or the `-t` option for specific single tests. Either can be used multiple times:

```
$ ./run_kselftest.sh -c size -c seccomp -t timers:posix_timers -tu  
-t timer:nanosleep
```

For other features see the script usage output, seen with the `-h` option.

15.6 Timeout for selftests

Selftests are designed to be quick and so a default timeout is used of 45 seconds for each test. Tests can override the default timeout by adding a settings file in their directory and set a timeout variable there to the configured a desired upper timeout for the test. Only a few tests override the timeout with a value higher than 45 seconds, selftests strives to keep it that way. Timeouts in selftests are not considered fatal because the system under which a test runs may change and this can also modify the expected time it takes to run a test. If you have control over the systems which will run the tests you can configure a test runner on those systems to use a greater or lower timeout on the command line as with the `-o` or the `--override-timeout` argument. For example to use 165 seconds instead one would use:

```
$ ./run_kselftest.sh --override-timeout 165
```

You can look at the TAP output to see if you ran into the timeout. Test runners which know a test must run under a specific time can then optionally treat these timeouts then as fatal.

15.7 Packaging selftests

In some cases packaging is desired, such as when tests need to run on a different system. To package selftests, run:

```
$ make -C tools/testing/selftests gen_tar
```

This generates a tarball in the `INSTALL_PATH/kselftest-packages` directory. By default, `.gz` format is used. The tar compression format can be overridden by specifying a `FORMAT` make variable. Any value recognized by `tar's auto-compress` option is supported, such as:

```
$ make -C tools/testing/selftests gen_tar FORMAT=.xz
```

`make gen_tar` invokes `make install` so you can use it to package a subset of tests by using variables specified in [Running a subset of selftests](#) section:

```
$ make -C tools/testing/selftests gen_tar TARGETS="size" FORMAT=.xz
```

15.8 Contributing new tests

In general, the rules for selftests are

- Do as much as you can if you're not root;
- Don't take too long;
- Don't break the build on any architecture, and
- Don't cause the top-level "make run_tests" to fail if your feature is unconfigured.

15.9 Contributing new tests (details)

- In your Makefile, use facilities from lib.mk by including it instead of reinventing the wheel. Specify flags and binaries generation flags on need basis before including lib.mk.

```
CFLAGS = $(KHDR_INCLUDES)
TEST_GEN_PROGS := close_range_test
include ../../lib.mk
```

- Use TEST_GEN_XXX if such binaries or files are generated during compiling.
TEST_PROGS, TEST_GEN_PROGS mean it is the executable tested by default.
TEST_CUSTOM_PROGS should be used by tests that require custom build rules and prevent common build rule use.
TEST_PROGS are for test shell scripts. Please ensure shell script has its exec bit set. Otherwise, lib.mk run_tests will generate a warning.
TEST_CUSTOM_PROGS and TEST_PROGS will be run by common run_tests.

TEST_PROGS_EXTENDED, TEST_GEN_PROGS_EXTENDED mean it is the executable which is not tested by default. TEST_FILES, TEST_GEN_FILES mean it is the file which is used by test.

- First use the headers inside the kernel source and/or git repo, and then the system headers. Headers for the kernel release as opposed to headers installed by the distro on the system should be the primary focus to be able to find regressions. Use KHDR_INCLUDES in Makefile to include headers from the kernel source.
- If a test needs specific kernel config options enabled, add a config file in the test directory to enable them.

e.g: tools/testing/selftests/android/config

- Create a .gitignore file inside test directory and add all generated objects in it.
- Add new test name in TARGETS in selftests/Makefile:

```
TARGETS += android
```

- All changes should pass:

```
kseltest-{all,install,clean,gen_tar}
kseltest-{all,install,clean,gen_tar} 0=abo_path
kseltest-{all,install,clean,gen_tar} 0=rel_path
make -C tools/testing/selftests {all,install,clean,gen_tar}
make -C tools/testing/selftests {all,install,clean,gen_tar} 0=abs_path
make -C tools/testing/selftests {all,install,clean,gen_tar} 0=rel_path
```

15.10 Test Module

Kselftest tests the kernel from userspace. Sometimes things need testing from within the kernel, one method of doing this is to create a test module. We can tie the module into the kselftest framework by using a shell script test runner. `kseltest/module.sh` is designed to facilitate this process. There is also a header file provided to assist writing kernel modules that are for use with kselftest:

- `tools/testing/selftests/kseltest_module.h`
- `tools/testing/selftests/kseltest/module.sh`

Note that test modules should taint the kernel with `TAINT_TEST`. This will happen automatically for modules which are in the `tools/testing/` directory, or for modules which use the `kseltest_module.h` header above. Otherwise, you'll need to add `MODULE_INFO(test, "Y")` to your module source. selftests which do not load modules typically should not taint the kernel, but in cases where a non-test module is loaded, `TEST_TAINT` can be applied from userspace by writing to `/proc/sys/kernel/tainted`.

15.10.1 How to use

Here we show the typical steps to create a test module and tie it into kselftest. We use kselftests for lib/ as an example.

1. Create the test module
2. Create the test script that will run (load/unload) the module e.g. `tools/testing/selftests/lib/printf.sh`
3. Add line to config file e.g. `tools/testing/selftests/lib/config`
4. Add test script to makefile e.g. `tools/testing/selftests/lib/Makefile`
5. Verify it works:

```
# Assumes you have booted a fresh build of this kernel tree
cd /path/to/linux/tree
make kseltest-merge
make modules
sudo make modules_install
make TARGETS=lib kseltest
```

15.10.2 Example Module

A bare bones test module might look like this:

```
// SPDX-License-Identifier: GPL-2.0+
#define pr_fmt(fmt) KBUILD_MODNAME ": " fmt

#include "../tools/testing/selftests/kselftest_module.h"

KSTM_MODULE_GLOBALS();

/*
 * Kernel module for testing the foobinator
 */

static int __init test_function()
{
    ...
}

static void __init selftest(void)
{
    KSTM_CHECK_ZERO(do_test_case("", 0));
}

KSTM_MODULE_LOADERS(test_foo);
MODULE_AUTHOR("John Developer <jd@fooman.org>");
MODULE_LICENSE("GPL");
MODULE_INFO(test, "Y");
```

15.10.3 Example test script

```
#!/bin/bash
# SPDX-License-Identifier: GPL-2.0+
$(dirname $0)../kselftest/module.sh "foo" test_foo
```

15.11 Test Harness

The kselftest_harness.h file contains useful helpers to build tests. The test harness is for userspace testing, for kernel space testing see [Test Module](#) above.

The tests from tools/testing/selftests/seccomp/seccomp_bpf.c can be used as example.

15.11.1 Example

```
#include "../kseltest_harness.h"

TEST(standalone_test) {
    do_some_stuff;
    EXPECT_GT(10, stuff) {
        stuff_state_t state;
        enumerate_stuff_state(&state);
        TH_LOG("expectation failed with state: %s", state.msg);
    }
    more_stuff;
    ASSERT_NE(some_stuff, NULL) TH_LOG("how did it happen?!");
    last_stuff;
    EXPECT_EQ(0, last_stuff);
}

Fixture(my_fixture) {
    mytype_t *data;
    int awesomeness_level;
};

Fixture_Setup(my_fixture) {
    self->data = mytype_new();
    ASSERT_NE(NULL, self->data);
}
Fixture_TearDown(my_fixture) {
    mytype_free(self->data);
}
TEST_F(my_fixture, data_is_good) {
    EXPECT_EQ(1, is_my_data_good(self->data));
}

TEST_HARNESS_MAIN
```

15.11.2 Helpers

TH_LOG

`TH_LOG (fmt, ...)`

Parameters

`fmt`

format string

`...`

optional arguments

Description

`TH_LOG(format, ...)`

Optional debug logging function available for use in tests. Logging may be enabled or disabled by defining TH_LOG_ENABLED. E.g., `#define TH_LOG_ENABLED 1`

If no definition is provided, logging is enabled by default.

If there is no way to print an error message for the process running the test (e.g. not allowed to write to stderr), it is still possible to get the ASSERT_* number for which the test failed. This behavior can be enabled by writing `_metadata->no_print = true;` before the check sequence that is unable to print. When an error occur, instead of printing an error message and calling `abort(3)`, the test process call `_exit(2)` with the assert number as argument, which is then printed by the parent process.

TEST

`TEST (test_name)`

Defines the test function and creates the registration stub

Parameters

test_name

test name

Description

```
TEST(name) { implementation }
```

Defines a test by name. Names must be unique and tests must not be run in parallel. The implementation containing block is a function and scoping should be treated as such. Returning early may be performed with a bare "return;" statement.

EXPECT_* and ASSERT_* are valid in a `TEST() { }` context.

TEST_SIGNAL

`TEST_SIGNAL (test_name, signal)`

Parameters

test_name

test name

signal

signal number

Description

```
TEST_SIGNAL(name, signal) { implementation }
```

Defines a test by name and the expected term signal. Names must be unique and tests must not be run in parallel. The implementation containing block is a function and scoping should be treated as such. Returning early may be performed with a bare "return;" statement.

EXPECT_* and ASSERT_* are valid in a `TEST() { }` context.

FIXTURE_DATA

`FIXTURE_DATA (datatype_name)`

Wraps the struct name so we have one less argument to pass around

Parameters

datatype_name

datatype name

Description

```
FIXTURE_DATA(datatype_name)
```

Almost always, you want just [FIXTURE\(\)](#) instead (see below). This call may be used when the type of the fixture data is needed. In general, this should not be needed unless the *self* is being passed to a helper directly.

Fixture

Fixture (fixture_name)

Called once per fixture to setup the data and register

Parameters

fixture_name

fixture name

Description

```
FIXTURE(fixture_name) {  
    type property1;  
    ...  
};
```

Defines the data provided to [TEST_F\(\)](#)-defined tests as *self*. It should be populated and cleaned up using [FIXTURE_SETUP\(\)](#) and [FIXTURE_TEARDOWN\(\)](#).

Fixture_Setup

Fixture_Setup (fixture_name)

Prepares the setup function for the fixture. *_metadata* is included so that `EXPECT_*`, `ASSERT_*` etc. work correctly.

Parameters

fixture_name

fixture name

Description

```
FIXTURE_SETUP(fixture_name) { implementation }
```

Populates the required "setup" function for a fixture. An instance of the datatype defined with [FIXTURE_DATA\(\)](#) will be exposed as *self* for the implementation.

`ASSERT_*` are valid for use in this context and will preempt the execution of any dependent fixture tests.

A bare "return;" statement may be used to return early.

Fixture_Teardown

```
Fixture_Teardown (fixture_name)
```

Parameters**fixture_name**

fixture name

Description

metadata is included so that EXPECT*, ASSERT_* etc. work correctly.

```
Fixture_Teardown(fixture_name) { implementation }
```

Populates the required "teardown" function for a fixture. An instance of the datatype defined with [Fixture_Data\(\)](#) will be exposed as *self* for the implementation to clean up.

A bare "return;" statement may be used to return early.

Fixture_Variant

```
Fixture_Variant (fixture_name)
```

Optionally called once per fixture to declare fixture variant

Parameters**fixture_name**

fixture name

Description

```
Fixture_Variant(fixture_name) {
    type property1;
    ...
};
```

Defines type of constant parameters provided to [Fixture_Setup\(\)](#), [Test_F\(\)](#) and [Fixture_Teardown](#) as *variant*. Variants allow the same tests to be run with different arguments.

Fixture_Variant_Add

```
Fixture_Variant_Add (fixture_name, variant_name)
```

Called once per fixture variant to setup and register the data

Parameters**fixture_name**

fixture name

variant_name

name of the parameter set

Description

```
Fixture_Variant_Add(fixture_name, variant_name) {
    .property1 = val1,
```

```
}; ...
```

Defines a variant of the test fixture, provided to `Fixture_Setup()` and `Test_F()` as *variant*. Tests of each fixture will be run once for each variant.

TEST_F

```
TEST_F (fixture_name, test_name)
```

Emits test registration and helpers for fixture-based test cases

Parameters

fixture_name

fixture name

test_name

test name

Description

```
TEST_F(fixture, name) { implementation }
```

Defines a test that depends on a fixture (e.g., is part of a test case). Very similar to `TEST()` except that *self* is the setup instance of fixture's datatype exposed for use by the implementation.

TEST_HARNESS_MAIN

```
TEST_HARNESS_MAIN ()
```

Simple wrapper to run the test harness

Parameters

Description

```
TEST_HARNESS_MAIN
```

Use once to append a main() to the test file.

15.11.3 Operators

Operators for use in `TEST()` and `TEST_F()`. `ASSERT_*` calls will stop test execution immediately. `EXPECT_*` calls will emit a failure warning, note it, and continue.

ASSERT_EQ

```
ASSERT_EQ (expected, seen)
```

Parameters

expected

expected value

seen

measured value

Description

ASSERT_EQ(expected, measured): expected == measured

ASSERT_NE

ASSERT_NE (expected, seen)

Parameters

expected

expected value

seen

measured value

Description

ASSERT_NE(expected, measured): expected != measured

ASSERT_LT

ASSERT_LT (expected, seen)

Parameters

expected

expected value

seen

measured value

Description

ASSERT_LT(expected, measured): expected < measured

ASSERT_LE

ASSERT_LE (expected, seen)

Parameters

expected

expected value

seen

measured value

Description

ASSERT_LE(expected, measured): expected <= measured

ASSERT_GT

ASSERT_GT (expected, seen)

Parameters

expected

expected value

seen

measured value

Description

ASSERT_GT(expected, measured): expected > measured

ASSERT_GE

ASSERT_GE (expected, seen)

Parameters

expected

expected value

seen

measured value

Description

ASSERT_GE(expected, measured): expected \geq measured

ASSERT_NULL

ASSERT_NULL (seen)

Parameters

seen

measured value

Description

ASSERT_NULL(measured): NULL == measured

ASSERT_TRUE

ASSERT_TRUE (seen)

Parameters

seen

measured value

Description

ASSERT_TRUE(measured): measured != 0

ASSERT_FALSE

ASSERT_FALSE (seen)

Parameters

seen

measured value

Description

ASSERT_FALSE(measured): measured == 0

ASSERT_STREQ

ASSERT_STREQ (expected, seen)

Parameters

expected

expected value

seen

measured value

Description

ASSERT_STREQ(expected, measured): !strcmp(expected, measured)

ASSERT_STRNE

ASSERT_STRNE (expected, seen)

Parameters**expected**

expected value

seen

measured value

Description

ASSERT_STRNE(expected, measured): strcmp(expected, measured)

EXPECT_EQ

EXPECT_EQ (expected, seen)

Parameters**expected**

expected value

seen

measured value

Description

EXPECT_EQ(expected, measured): expected == measured

EXPECT_NE

EXPECT_NE (expected, seen)

Parameters**expected**

expected value

seen

measured value

Description

EXPECT_NE(expected, measured): expected != measured

EXPECT_LT

EXPECT_LT (expected, seen)

Parameters

expected

expected value

seen

measured value

Description

EXPECT_LT(expected, measured): expected < measured

EXPECT_LE

EXPECT_LE (expected, seen)

Parameters

expected

expected value

seen

measured value

Description

EXPECT_LE(expected, measured): expected <= measured

EXPECT_GT

EXPECT_GT (expected, seen)

Parameters

expected

expected value

seen

measured value

Description

EXPECT_GT(expected, measured): expected > measured

EXPECT_GE

EXPECT_GE (expected, seen)

Parameters

expected

expected value

seen

measured value

Description

EXPECT_GE(expected, measured): expected >= measured

EXPECT_NULL

EXPECT_NULL (seen)

Parameters

seen

measured value

Description

EXPECT_NULL(measured): NULL == measured

EXPECT_TRUE

EXPECT_TRUE (seen)

Parameters**seen**

measured value

Description

EXPECT_TRUE(measured): 0 != measured

EXPECT_FALSE

EXPECT_FALSE (seen)

Parameters**seen**

measured value

Description

EXPECT_FALSE(measured): 0 == measured

EXPECT_STREQ

EXPECT_STREQ (expected, seen)

Parameters**expected**

expected value

seen

measured value

Description

EXPECT_STREQ(expected, measured): !strcmp(expected, measured)

EXPECT_STRNE

EXPECT_STRNE (expected, seen)

Parameters**expected**

expected value

seen

measured value

Description

EXPECT_STRNE(expected, measured): strcmp(expected, measured)

KUNIT - LINUX KERNEL UNIT TESTING

16.1 Getting Started

This page contains an overview of the kunit_tool and KUnit framework, teaching how to run existing tests and then how to write a simple test case, and covers common problems users face when using KUnit for the first time.

16.1.1 Installing Dependencies

KUnit has the same dependencies as the Linux kernel. As long as you can build the kernel, you can run KUnit.

16.1.2 Running tests with kunit_tool

kunit_tool is a Python script, which configures and builds a kernel, runs tests, and formats the test results. From the kernel repository, you can run kunit_tool:

```
./tools/testing/kunit/kunit.py run
```

Note: You may see the following error: "The source tree is not clean, please run 'make ARCH=um mrproper'"

This happens because internally kunit.py specifies .kunit (default option) as the build directory in the command make O=output/dir through the argument --build_dir. Hence, before starting an out-of-tree build, the source tree must be clean.

There is also the same caveat mentioned in the "Build directory for the kernel" section of the admin-guide, that is, its use, it must be used for all invocations of make. The good news is that it can indeed be solved by running make ARCH=um mrproper, just be aware that this will delete the current configuration and all generated files.

If everything worked correctly, you should see the following:

```
Configuring KUnit Kernel ...
Building KUnit Kernel ...
Starting KUnit Kernel ...
```

The tests will pass or fail.

Note: Because it is building a lot of sources for the first time, the Building KUnit Kernel step may take a while.

For detailed information on this wrapper, see: [Running tests with kunit_tool](#).

Selecting which tests to run

By default, `kunit_tool` runs all tests reachable with minimal configuration, that is, using default values for most of the `kconfig` options. However, you can select which tests to run by:

- [Customizing Kconfig](#) used to compile the kernel, or
- [Filtering tests by name](#) to select specifically which compiled tests to run.

Customizing Kconfig

A good starting point for the `.kunitconfig` is the KUnit default config. If you didn't run `kunit.py run` yet, you can generate it by running:

```
cd $PATH_TO_LINUX_REPO
tools/testing/kunit/kunit.py config
cat .kunit/.kunitconfig
```

Note: `.kunitconfig` lives in the `--build_dir` used by `kunit.py`, which is `.kunit` by default.

Before running the tests, `kunit_tool` ensures that all config options set in `.kunitconfig` are set in the kernel `.config`. It will warn you if you have not included dependencies for the options used.

There are many ways to customize the configurations:

- a. Edit `.kunit/.kunitconfig`. The file should contain the list of `kconfig` options required to run the desired tests, including their dependencies. You may want to remove `CONFIG_KUNIT_ALL_TESTS` from the `.kunitconfig` as it will enable a number of additional tests that you may not want. If you need to run on an architecture other than UML see [Running tests on QEMU](#).
- b. Enable additional `kconfig` options on top of `.kunit/.kunitconfig`. For example, to include the kernel's linked-list test you can run:

```
./tools/testing/kunit/kunit.py run \
    --kconfig_add CONFIG_LIST_KUNIT_TEST=y
```

- c. Provide the path of one or more `.kunitconfig` files from the tree. For example, to run only `FAT_FS` and `EXT4` tests you can run:

```
./tools/testing/kunit/kunit.py run \
    --kunitconfig ./fs/fat/.kunitconfig \
    --kunitconfig ./fs/ext4/.kunitconfig
```

- d. If you change the .kunitconfig, kunit.py will trigger a rebuild of the .config file. But you can edit the .config file directly or with tools like make menuconfig 0=.kunit. As long as its a superset of .kunitconfig, kunit.py won't overwrite your changes.

Note: To save a .kunitconfig after finding a satisfactory configuration:

```
make savedefconfig 0=.kunit
cp .kunit/defconfig .kunit/.kunitconfig
```

Filtering tests by name

If you want to be more specific than Kconfig can provide, it is also possible to select which tests to execute at boot-time by passing a glob filter (read instructions regarding the pattern in the manpage *glob(7)*). If there is a " ." (period) in the filter, it will be interpreted as a separator between the name of the test suite and the test case, otherwise, it will be interpreted as the name of the test suite. For example, let's assume we are using the default config:

- a. inform the name of a test suite, like "kunit_executor_test", to run every test case it contains:

```
./tools/testing/kunit/kunit.py run "kunit_executor_test"
```

- b. inform the name of a test case prefixed by its test suite, like "example.example_simple_test", to run specifically that test case:

```
./tools/testing/kunit/kunit.py run "example.example_simple_test"
```

- c. use wildcard characters (*?[]) to run any test case that matches the pattern, like "*.*64*" to run test cases containing "64" in the name inside any test suite:

```
./tools/testing/kunit/kunit.py run "*.*64*"
```

16.1.3 Running Tests without the KUnit Wrapper

If you do not want to use the KUnit Wrapper (for example: you want code under test to integrate with other systems, or use a different/ unsupported architecture or configuration), KUnit can be included in any kernel, and the results are read out and parsed manually.

Note: CONFIG_KUNIT should not be enabled in a production environment. Enabling KUnit disables Kernel Address-Space Layout Randomization (KASLR), and tests may affect the state of the kernel in ways not suitable for production.

Configuring the Kernel

To enable KUnit itself, you need to enable the CONFIG_KUNIT Kconfig option (under Kernel Hacking/Kernel Testing and Coverage in menuconfig). From there, you can enable any KUnit tests. They usually have config options ending in _KUNIT_TEST.

KUnit and KUnit tests can be compiled as modules. The tests in a module will run when the module is loaded.

Running Tests (without KUnit Wrapper)

Build and run your kernel. In the kernel log, the test output is printed out in the TAP format. This will only happen by default if KUnit/tests are built-in. Otherwise the module will need to be loaded.

Note: Some lines and/or data may get interspersed in the TAP output.

16.1.4 Writing Your First Test

In your kernel repository, let's add some code that we can test.

1. Create a file drivers/misc/example.h, which includes:

```
int misc_example_add(int left, int right);
```

2. Create a file drivers/misc/example.c, which includes:

```
#include <linux/errno.h>

#include "example.h"

int misc_example_add(int left, int right)
{
    return left + right;
}
```

3. Add the following lines to drivers/misc/Kconfig:

```
config MISC_EXAMPLE
    bool "My example"
```

4. Add the following lines to drivers/misc/Makefile:

```
obj-$(CONFIG_MISC_EXAMPLE) += example.o
```

Now we are ready to write the test cases.

1. Add the below test case in drivers/misc/example_test.c:

```
#include <kunit/test.h>
#include "example.h"

/* Define the test cases. */

static void misc_example_add_test_basic(struct kunit *test)
{
    KUNIT_EXPECT_EQ(test, 1, misc_example_add(1, 0));
    KUNIT_EXPECT_EQ(test, 2, misc_example_add(1, 1));
    KUNIT_EXPECT_EQ(test, 0, misc_example_add(-1, 1));
    KUNIT_EXPECT_EQ(test, INT_MAX, misc_example_add(0, INT_MAX));
    KUNIT_EXPECT_EQ(test, -1, misc_example_add(INT_MAX, INT_MIN));
}

static void misc_example_test_failure(struct kunit *test)
{
    KUNIT_FAIL(test, "This test never passes.");
}

static struct kunit_case misc_example_test_cases[] = {
    KUNIT_CASE(misc_example_add_test_basic),
    KUNIT_CASE(misc_example_test_failure),
    {}
};

static struct kunit_suite misc_example_test_suite = {
    .name = "misc-example",
    .test_cases = misc_example_test_cases,
};
kunit_test_suite(misc_example_test_suite);

MODULE_LICENSE("GPL");
```

2. Add the following lines to drivers/misc/Kconfig:

```
config MISC EXAMPLE TEST
    tristate "Test for my example" if !KUNIT_ALL_TESTS
    depends on MISC EXAMPLE && KUNIT
    default KUNIT_ALL_TESTS
```

Note: If your test does not support being built as a loadable module (which is discouraged), replace tristate by bool, and depend on KUNIT=y instead of KUNIT.

3. Add the following lines to drivers/misc/Makefile:

```
obj-$(CONFIG_MISC EXAMPLE TEST) += example_test.o
```

4. Add the following lines to .kunit/.kunitconfig:

```
CONFIG_MISC EXAMPLE=y
CONFIG_MISC EXAMPLE_TEST=y
```

5. Run the test:

```
./tools/testing/kunit/kunit.py run
```

You should see the following failure:

```
...
[16:08:57] [PASSED] misc-example:misc_example_add_test_basic
[16:08:57] [FAILED] misc-example:misc_example_test_failure
[16:08:57] EXPECTATION FAILED at drivers/misc/example-test.c:17
[16:08:57]     This test never passes.
...
```

Congrats! You just wrote your first KUnit test.

16.1.5 Next Steps

If you're interested in using some of the more advanced features of kunit.py, take a look at [Running tests with kunit_tool](#)

If you'd like to run tests without using kunit.py, check out [Run Tests without kunit_tool](#)

For more information on writing KUnit tests (including some common techniques for testing different things), see [Writing Tests](#)

16.2 KUnit Architecture

The KUnit architecture is divided into two parts:

- [In-Kernel Testing Framework](#)
- [kunit_tool \(Command-line Test Harness\)](#)

16.2.1 In-Kernel Testing Framework

The kernel testing library supports KUnit tests written in C using KUnit. These KUnit tests are kernel code. KUnit performs the following tasks:

- Organizes tests
- Reports test results
- Provides test utilities

Test Cases

The test case is the fundamental unit in KUnit. KUnit test cases are organised into suites. A KUnit test case is a function with type signature `void (*)(struct kunit *test)`. These test case functions are wrapped in a struct called `struct kunit_case`.

Each KUnit test case receives a `struct kunit` context object that tracks a running test. The KUnit assertion macros and other KUnit utilities use the `struct kunit` context object. As an exception, there are two fields:

- `->priv`: The setup functions can use it to store arbitrary test user data.
- `->param_value`: It contains the parameter value which can be retrieved in the parameterized tests.

Test Suites

A KUnit suite includes a collection of test cases. The KUnit suites are represented by the `struct kunit_suite`. For example:

```
static struct kunit_case example_test_cases[] = {
    KUNIT_CASE(example_test_foo),
    KUNIT_CASE(example_test_bar),
    KUNIT_CASE(example_test_baz),
    {}
};

static struct kunit_suite example_test_suite = {
    .name = "example",
    .init = example_test_init,
    .exit = example_test_exit,
    .test_cases = example_test_cases,
};
kunit_test_suite(example_test_suite);
```

In the above example, the test suite `example_test_suite`, runs the test cases `example_test_foo`, `example_test_bar`, and `example_test_baz`. Before running the test, the `example_test_init` is called and after running the test, `example_test_exit` is called. The `kunit_test_suite(example_test_suite)` registers the test suite with the KUnit test framework.

Executor

The KUnit executor can list and run built-in KUnit tests on boot. The Test suites are stored in a linker section called `.kunit_test_suites`. For the code, see `KUNIT_TABLE()` macro definition in `include/asm-generic/vmlinux.lds.h`. The linker section consists of an array of pointers to `struct kunit_suite`, and is populated by the `kunit_test_suites()` macro. The KUnit executor iterates over the linker section array in order to run all the tests that are compiled into the kernel.

On the kernel boot, the KUnit executor uses the start and end addresses of this section to iterate over and run all tests. For the implementation of the executor, see `lib/kunit/executor.c`. When

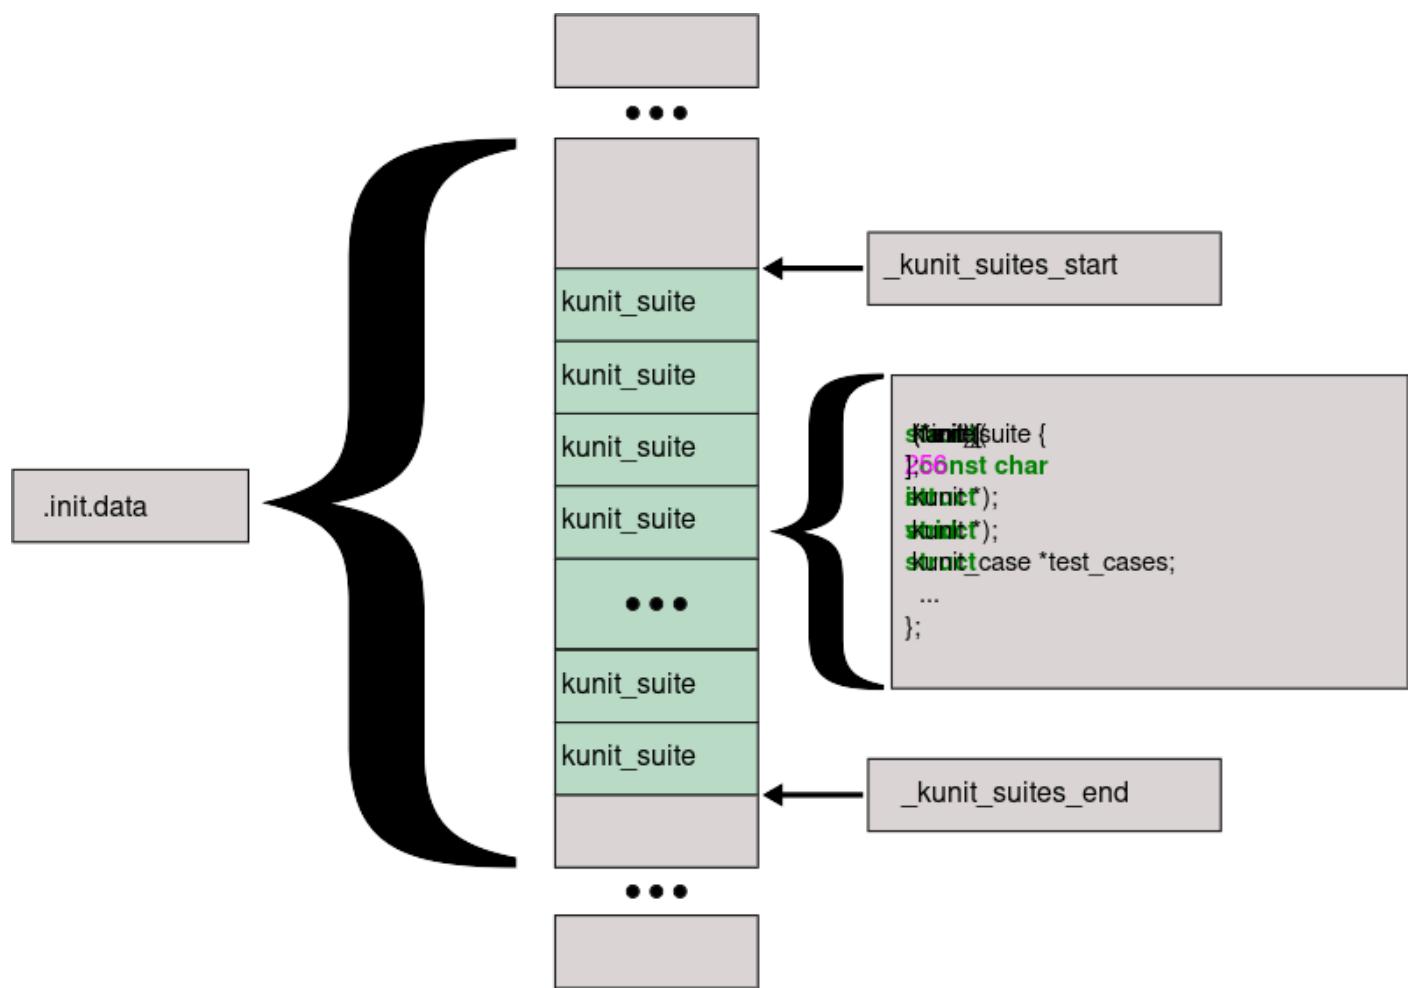


Fig. 1: KUnit Suite Memory Diagram

built as a module, the `kunit_test_suites()` macro defines a `module_init()` function, which runs all the tests in the compilation unit instead of utilizing the executor.

In KUnit tests, some error classes do not affect other tests or parts of the kernel, each KUnit case executes in a separate thread context. See the `kunit_try_catch_run()` function in `lib/kunit/try-catch.c`.

Assertion Macros

KUnit tests verify state using expectations/assertions. All expectations/assertions are formatted as: `KUNIT_{EXPECT|ASSERT}_{<op>}[_MSG](kunit, property[, message])`

- `{EXPECT|ASSERT}` determines whether the check is an assertion or an expectation. In the event of a failure, the testing flow differs as follows:

- For expectations, the test is marked as failed and the failure is logged.
- Failing assertions, on the other hand, result in the test case being terminated immediately.

- * Assertions call the function: `void __noreturn __kunit_abort(struct kunit *)`.
- * `__kunit_abort` calls the function: `void __noreturn kunit_try_catch_throw(struct kunit_try_catch *try_catch)`.
- * `kunit_try_catch_throw` calls the function: `void kthread_complete_and_exit(struct completion *, long) __noreturn;` and terminates the special thread context.

- `<op>` denotes a check with options: TRUE (supplied property has the boolean value "true"), EQ (two supplied properties are equal), NOT_ERR_OR_NULL (supplied pointer is not null and does not contain an "err" value).
- `[_MSG]` prints a custom message on failure.

Test Result Reporting

KUnit prints the test results in KTAP format. KTAP is based on TAP14, see [The Kernel Test Anything Protocol \(KTAP\), version 1](#). KTAP works with KUnit and Kselftest. The KUnit executor prints KTAP results to dmesg, and debugfs (if configured).

Parameterized Tests

Each KUnit parameterized test is associated with a collection of parameters. The test is invoked multiple times, once for each parameter value and the parameter is stored in the `param_value` field. The test case includes a `KUNIT_CASE_PARAM()` macro that accepts a generator function. The generator function is passed the previous parameter and returns the next parameter. It also includes a macro for generating array-based common-case generators.

16.2.2 kunit_tool (Command-line Test Harness)

`kunit_tool` is a Python script, found in `tools/testing/kunit/kunit.py`. It is used to configure, build, execute, parse test results and run all of the previous commands in correct order (i.e., configure, build, execute and parse). You have two options for running KUnit tests: either build the kernel with KUnit enabled and manually parse the results (see [Run Tests without kunit_tool](#)) or use `kunit_tool` (see [Running tests with kunit_tool](#)).

- `configure` command generates the kernel `.config` from a `.kunitconfig` file (and any architecture-specific options). The Python scripts available in `qemu_configs` folder (for example, `tools/testing/kunit/qemu_configs/powerpc.py`) contains additional configuration options for specific architectures. It parses both the existing `.config` and the `.kunitconfig` files to ensure that `.config` is a superset of `.kunitconfig`. If not, it will combine the two and run `make olddefconfig` to regenerate the `.config` file. It then checks to see if `.config` has become a superset. This verifies that all the Kconfig dependencies are correctly specified in the file `.kunitconfig`. The `kunit_config.py` script contains the code for parsing Kconfigs. The code which runs `make olddefconfig` is part of the `kunit_kernel.py` script. You can invoke this command through: `./tools/testing/kunit/kunit.py config` and generate a `.config` file.
- `build` runs `make` on the kernel tree with required options (depends on the architecture and some options, for example: `build_dir`) and reports any errors. To build a KUnit kernel from the current `.config`, you can use the `build` argument: `./tools/testing/kunit/kunit.py build`.
- `exec` command executes kernel results either directly (using User-mode Linux configuration), or through an emulator such as QEMU. It reads results from the log using standard output (`stdout`), and passes them to `parse` to be parsed. If you already have built a kernel with built-in KUnit tests, you can run the kernel and display the test results with the `exec` argument: `./tools/testing/kunit/kunit.py exec`.
- `parse` extracts the KTAP output from a kernel log, parses the test results, and prints a summary. For failed tests, any diagnostic output will be included.

16.3 Running tests with kunit_tool

We can either run KUnit tests using `kunit_tool` or can run tests manually, and then use `kunit_tool` to parse the results. To run tests manually, see: [Run Tests without kunit_tool](#). As long as we can build the kernel, we can run KUnit.

`kunit_tool` is a Python script which configures and builds a kernel, runs tests, and formats the test results.

Run command:

```
./tools/testing/kunit/kunit.py run
```

We should see the following:

```
Configuring KUnit Kernel ...
Building KUnit kernel...
Starting KUnit kernel...
```

We may want to use the following options:

```
./tools/testing/kunit/kunit.py run --timeout=30 --jobs=`nproc` --all`
```

- `--timeout` sets a maximum amount of time for tests to run.
- `--jobs` sets the number of threads to build the kernel.

`kunit_tool` will generate a `.kunitconfig` with a default configuration, if no other `.kunitconfig` file exists (in the build directory). In addition, it verifies that the generated `.config` file contains the `CONFIG` options in the `.kunitconfig`. It is also possible to pass a separate `.kunitconfig` fragment to `kunit_tool`. This is useful if we have several different groups of tests we want to run independently, or if we want to use pre-defined test configs for certain subsystems.

To use a different `.kunitconfig` file (such as one provided to test a particular subsystem), pass it as an option:

```
./tools/testing/kunit/kunit.py run --kunitconfig=fs/ext4/.kunitconfig
```

To view `kunit_tool` flags (optional command-line arguments), run:

```
./tools/testing/kunit/kunit.py run --help
```

16.3.1 Creating a `.kunitconfig` file

If we want to run a specific set of tests (rather than those listed in the KUnit `defconfig`), we can provide Kconfig options in the `.kunitconfig` file. For default `.kunitconfig`, see: <https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/tree/tools/testing/kunit/configs/default.config>. A `.kunitconfig` is a `minconfig` (a `.config` generated by running `make savedefconfig`), used for running a specific set of tests. This file contains the regular Kernel configs with specific test targets. The `.kunitconfig` also contains any other config options required by the tests (For example: dependencies for features under tests, configs that enable/disable certain code blocks, arch configs and so on).

To create a `.kunitconfig`, using the KUnit `defconfig`:

```
cd $PATH_TO_LINUX_REPO
cp tools/testing/kunit/configs/default.config .kunit/.kunitconfig
```

We can then add any other Kconfig options. For example:

```
CONFIG_LIST_KUNIT_TEST=y
```

`kunit_tool` ensures that all config options in `.kunitconfig` are set in the kernel `.config` before running the tests. It warns if we have not included the options dependencies.

Note: Removing something from the `.kunitconfig` will not rebuild the `.config` file. The configuration is only updated if the `.kunitconfig` is not a subset of `.config`. This means that we can use other tools (For example: `make menuconfig`) to adjust other config options. The build dir needs to be set for `make menuconfig` to work, therefore by default use `make O=.kunit menuconfig`.

16.3.2 Configuring, building, and running tests

If we want to make manual changes to the KUnit build process, we can run part of the KUnit build process independently. When running `kunit_tool`, from a `.kunitconfig`, we can generate a `.config` by using the `config` argument:

```
./tools/testing/kunit/kunit.py config
```

To build a KUnit kernel from the current `.config`, we can use the `build` argument:

```
./tools/testing/kunit/kunit.py build
```

If we already have built UML kernel with built-in KUnit tests, we can run the kernel, and display the test results with the `exec` argument:

```
./tools/testing/kunit/kunit.py exec
```

The `run` command discussed in section: **Running tests with `kunit_tool`**, is equivalent to running the above three commands in sequence.

16.3.3 Parsing test results

KUnit tests output displays results in TAP (Test Anything Protocol) format. When running tests, `kunit_tool` parses this output and prints a summary. To see the raw test results in TAP format, we can pass the `--raw_output` argument:

```
./tools/testing/kunit/kunit.py run --raw_output
```

If we have KUnit results in the raw TAP format, we can parse them and print the human-readable summary with the `parse` command for `kunit_tool`. This accepts a filename for an argument, or will read from standard input.

```
# Reading from a file
./tools/testing/kunit/kunit.py parse /var/log/dmesg
# Reading from stdin
dmesg | ./tools/testing/kunit/kunit.py parse
```

16.3.4 Filtering tests

By passing a bash style glob filter to the `exec` or `run` commands, we can run a subset of the tests built into a kernel . For example: if we only want to run KUnit resource tests, use:

```
./tools/testing/kunit/kunit.py run 'kunit-resource*'
```

This uses the standard glob format with wildcard characters.

16.3.5 Running tests on QEMU

kunit_tool supports running tests on qemu as well as via UML. To run tests on qemu, by default it requires two flags:

- `--arch`: Selects a configs collection (Kconfig, qemu config options and so on), that allow KUnit tests to be run on the specified architecture in a minimal way. The architecture argument is same as the option name passed to the `ARCH` variable used by Kbuild. Not all architectures currently support this flag, but we can use `--qemu_config` to handle it. If `um` is passed (or this flag is ignored), the tests will run via UML. Non-UML architectures, for example: i386, x86_64, arm and so on; run on qemu.
- `--cross_compile`: Specifies the Kbuild toolchain. It passes the same argument as passed to the `CROSS_COMPILE` variable used by Kbuild. As a reminder, this will be the prefix for the toolchain binaries such as GCC. For example:
 - `sparc64-linux-gnu` if we have the sparc toolchain installed on our system.
 - `$HOME/toolchains/microblaze/gcc-9.2.0-nolibc/microblaze-linux/bin/microblaze-linux` if we have downloaded the microblaze toolchain from the 0-day website to a directory in our home directory called `toolchains`.

This means that for most architectures, running under qemu is as simple as:

```
./tools/testing/kunit/kunit.py run --arch=x86_64
```

When cross-compiling, we'll likely need to specify a different toolchain, for example:

```
./tools/testing/kunit/kunit.py run \
  --arch=s390 \
  --cross_compile=s390x-linux-gnu-
```

If we want to run KUnit tests on an architecture not supported by the `--arch` flag, or want to run KUnit tests on qemu using a non-default configuration; then we can write our own ``QemuConfig``. These QemuConfigs are written in Python. They have an import line `from..qemu_config import QemuArchParams` at the top of the file. The file must contain a variable called `QEMU_ARCH` that has an instance of `QemuArchParams` assigned to it. See example in: `tools/testing/kunit/qemu_configs/x86_64.py`.

Once we have a `QemuConfig`, we can pass it into `kunit_tool`, using the `--qemu_config` flag. When used, this flag replaces the `--arch` flag. For example: using `tools/testing/kunit/qemu_configs/x86_64.py`, the invocation appear as

```
./tools/testing/kunit/kunit.py run \
  --timeout=60 \
  --jobs=12 \
  --qemu_config=./tools/testing/kunit/qemu_configs/x86_64.py
```

16.3.6 Running command-line arguments

kunit_tool has a number of other command-line arguments which can be useful for our test environment. Below are the most commonly used command line arguments:

- `--help`: Lists all available options. To list common options, place `--help` before the command. To list options specific to that command, place `--help` after the command.

Note: Different commands (config, build, run, etc) have different supported options.

- `--build_dir`: Specifies kunit_tool build directory. It includes the `.kunitconfig`, `.config` files and compiled kernel.
- `--make_options`: Specifies additional options to pass to make, when compiling a kernel (using build or run commands). For example: to enable compiler warnings, we can pass `--make_options W=1`.
- `--alltests`: Enable a predefined set of options in order to build as many tests as possible.

Note: The list of enabled options can be found in `tools/testing/kunit/configs/all_tests.config`.

If you only want to enable all tests with otherwise satisfied dependencies, instead add `CONFIG_KUNIT_ALL_TESTS=y` to your `.kunitconfig`.

- `--kunitconfig`: Specifies the path or the directory of the `.kunitconfig` file. For example:
 - `lib/kunit/.kunitconfig` can be the path of the file.
 - `lib/kunit` can be the directory in which the file is located.

This file is used to build and run with a predefined set of tests and their dependencies. For example, to run tests for a given subsystem.

- `--kconfig_add`: Specifies additional configuration options to be appended to the `.kunitconfig` file. For example:

```
./tools/testing/kunit/kunit.py run --kconfig_add CONFIG_KASAN=y
```

- `--arch`: Runs tests on the specified architecture. The architecture argument is same as the Kbuild ARCH environment variable. For example, i386, x86_64, arm, um, etc. Non-UML architectures run on qemu. Default is `um`.
- `--cross_compile`: Specifies the Kbuild toolchain. It passes the same argument as passed to the CROSS_COMPILE variable used by Kbuild. This will be the prefix for the toolchain binaries such as GCC. For example:
 - `sparc64-linux-gnu-` if we have the sparc toolchain installed on our system.
 - `$HOME/toolchains/microblaze/gcc-9.2.0-nolibc/microblaze-linux/bin/microblaze-linux` if we have downloaded the microblaze toolchain from the 0-day website to a specified path in our home directory called toolchains.
- `--qemu_config`: Specifies the path to a file containing a custom qemu architecture definition. This should be a python file containing a `QemuArchParams` object.

- `--qemu_args`: Specifies additional qemu arguments, for example, `-smp 8`.
- `--jobs`: Specifies the number of jobs (commands) to run simultaneously. By default, this is set to the number of cores on your system.
- `--timeout`: Specifies the maximum number of seconds allowed for all tests to run. This does not include the time taken to build the tests.
- `--kernel_args`: Specifies additional kernel command-line arguments. May be repeated.
- `--run_isolated`: If set, boots the kernel for each individual suite/test. This is useful for debugging a non-hermetic test, one that might pass/fail based on what ran before it.
- `--raw_output`: If set, generates unformatted output from kernel. Possible options are:
 - `all`: To view the full kernel output, use `--raw_output=all`.
 - `kunit`: This is the default option and filters to KUnit output. Use `--raw_output` or `--raw_output=kunit`.
- `--json`: If set, stores the test results in a JSON format and prints to `stdout` or saves to a file if a filename is specified.
- `--filter`: Specifies filters on test attributes, for example, `speed!=slow`. Multiple filters can be used by wrapping input in quotes and separating filters by commas. Example: `--filter "speed>slow, module=example"`.
- `--filter_action`: If set to `skip`, filtered tests will be shown as skipped in the output rather than showing no output.
- `--list_tests`: If set, lists all tests that will be run.
- `--list_tests_attr`: If set, lists all tests that will be run and all of their attributes.

16.4 Run Tests without kunit_tool

If we do not want to use `kunit_tool` (For example: we want to integrate with other systems, or run tests on real hardware), we can include KUnit in any kernel, read out results, and parse manually.

Note: KUnit is not designed for use in a production system. It is possible that tests may reduce the stability or security of the system.

16.4.1 Configure the Kernel

KUnit tests can run without `kunit_tool`. This can be useful, if:

- We have an existing kernel configuration to test.
- Need to run on real hardware (or using an emulator/VM `kunit_tool` does not support).
- Wish to integrate with some existing testing systems.

KUnit is configured with the `CONFIG_KUNIT` option, and individual tests can also be built by enabling their config options in our `.config`. KUnit tests usually (but don't always) have config

options ending in `_KUNIT_TEST`. Most tests can either be built as a module, or be built into the kernel.

Note: We can enable the `KUNIT_ALL_TESTS` config option to automatically enable all tests with satisfied dependencies. This is a good way of quickly testing everything applicable to the current config.

Once we have built our kernel (and/or modules), it is simple to run the tests. If the tests are built-in, they will run automatically on the kernel boot. The results will be written to the kernel log (`dmesg`) in TAP format.

If the tests are built as modules, they will run when the module is loaded.

```
# modprobe example-test
```

The results will appear in TAP format in `dmesg`.

16.4.2 debugfs

KUnit can be accessed from userspace via the debugfs filesystem (See more information about debugfs at Documentation/filesystems/debugfs.rst).

If `CONFIG_KUNIT_DEBUGFS` is enabled, the KUnit debugfs filesystem is mounted at `/sys/kernel/debug/kunit`. You can use this filesystem to perform the following actions.

16.4.3 Retrieve Test Results

You can use debugfs to retrieve KUnit test results. The test results are accessible from the debugfs filesystem in the following read-only file:

```
/sys/kernel/debug/kunit/<test_suite>/results
```

The test results are printed in a KTAP document. Note this document is separate to the kernel log and thus, may have different test suite numbering.

16.4.4 Run Tests After Kernel Has Booted

You can use the debugfs filesystem to trigger built-in tests to run after boot. To run the test suite, you can use the following command to write to the `/sys/kernel/debug/kunit/<test_suite>/run` file:

```
echo "any string" > /sys/kernel/debugfs/kunit/<test_suite>/run
```

As a result, the test suite runs and the results are printed to the kernel log.

However, this feature is not available with KUnit suites that use init data, because init data may have been discarded after the kernel boots. KUnit suites that use init data should be defined using the `kunit_test_init_section_suites()` macro.

Also, you cannot use this feature to run tests concurrently. Instead a test will wait to run until other tests have completed or failed.

Note: For test authors, to use this feature, tests will need to correctly initialise and/or clean up any data, so the test runs correctly a second time.

16.5 Writing Tests

16.5.1 Test Cases

The fundamental unit in KUnit is the test case. A test case is a function with the signature `void (*)(struct kunit *test)`. It calls the function under test and then sets *expectations* for what should happen. For example:

```
void example_test_success(struct kunit *test)
{
}

void example_test_failure(struct kunit *test)
{
    KUNIT_FAIL(test, "This test never passes.");
}
```

In the above example, `example_test_success` always passes because it does nothing; no expectations are set, and therefore all expectations pass. On the other hand `example_test_failure` always fails because it calls `KUNIT_FAIL`, which is a special expectation that logs a message and causes the test case to fail.

Expectations

An *expectation* specifies that we expect a piece of code to do something in a test. An expectation is called like a function. A test is made by setting expectations about the behavior of a piece of code under test. When one or more expectations fail, the test case fails and information about the failure is logged. For example:

```
void add_test_basic(struct kunit *test)
{
    KUNIT_EXPECT_EQ(test, 1, add(1, 0));
    KUNIT_EXPECT_EQ(test, 2, add(1, 1));
}
```

In the above example, `add_test_basic` makes a number of assertions about the behavior of a function called `add`. The first parameter is always of type `struct kunit *`, which contains information about the current test context. The second parameter, in this case, is what the value is expected to be. The last value is what the value actually is. If `add` passes all of these expectations, the test case, `add_test_basic` will pass; if any one of these expectations fails, the test case will fail.

A test case *fails* when any expectation is violated; however, the test will continue to run, and try other expectations until the test case ends or is otherwise terminated. This is as opposed to *assertions* which are discussed later.

To learn about more KUnit expectations, see [Test API](#).

Note: A single test case should be short, easy to understand, and focused on a single behavior.

For example, if we want to rigorously test the add function above, create additional tests cases which would test each property that an add function should have as shown below:

```
void add_test_basic(struct kunit *test)
{
    KUNIT_EXPECT_EQ(test, 1, add(1, 0));
    KUNIT_EXPECT_EQ(test, 2, add(1, 1));
}

void add_test_negative(struct kunit *test)
{
    KUNIT_EXPECT_EQ(test, 0, add(-1, 1));
}

void add_test_max(struct kunit *test)
{
    KUNIT_EXPECT_EQ(test, INT_MAX, add(0, INT_MAX));
    KUNIT_EXPECT_EQ(test, -1, add(INT_MAX, INT_MIN));
}

void add_test_overflow(struct kunit *test)
{
    KUNIT_EXPECT_EQ(test, INT_MIN, add(INT_MAX, 1));
}
```

Assertions

An assertion is like an expectation, except that the assertion immediately terminates the test case if the condition is not satisfied. For example:

```
static void test_sort(struct kunit *test)
{
    int *a, i, r = 1;
    a = kunit_kmalloc_array(test, TEST_LEN, sizeof(*a), GFP_KERNEL);
    KUNIT_ASSERT_NOT_ERR_OR_NULL(test, a);
    for (i = 0; i < TEST_LEN; i++) {
        r = (r * 725861) % 6599;
        a[i] = r;
    }
    sort(a, TEST_LEN, sizeof(*a), cmpint, NULL);
    for (i = 0; i < TEST_LEN-1; i++)
        KUNIT_EXPECT_LT(test, a[i], a[i + 1]);
}
```

In this example, we need to be able to allocate an array to test the `sort()` function. So we use `KUNIT_ASSERT_NOT_ERR_OR_NULL()` to abort the test if there's an allocation error.

Note: In other test frameworks, ASSERT macros are often implemented by calling `return` so they only work from the test function. In KUnit, we stop the current kthread on failure, so you can call them from anywhere.

Note: Warning: There is an exception to the above rule. You shouldn't use assertions in the suite's `exit()` function, or in the `free` function for a resource. These run when a test is shutting down, and an assertion here prevents further cleanup code from running, potentially leading to a memory leak.

16.5.2 Customizing error messages

Each of the `KUNIT_EXPECT` and `KUNIT_ASSERT` macros have a `_MSG` variant. These take a format string and arguments to provide additional context to the automatically generated error messages.

```
char some_str[41];
generate_sha1_hex_string(some_str);

/* Before. Not easy to tell why the test failed. */
KUNIT_EXPECT_EQ(test, strlen(some_str), 40);

/* After. Now we see the offending string. */
KUNIT_EXPECT_EQ_MSG(test, strlen(some_str), 40, "some_str='%s'", some_str);
```

Alternatively, one can take full control over the error message by using `KUNIT_FAIL()`, e.g.

```
/* Before */
KUNIT_EXPECT_EQ(test, some_setup_function(), 0);

/* After: full control over the failure message. */
if (some_setup_function())
    KUNIT_FAIL(test, "Failed to setup thing for testing");
```

Test Suites

We need many test cases covering all the unit's behaviors. It is common to have many similar tests. In order to reduce duplication in these closely related tests, most unit testing frameworks (including KUnit) provide the concept of a *test suite*. A test suite is a collection of test cases for a unit of code with optional setup and teardown functions that run before/after the whole suite and/or every test case.

Note: A test case will only run if it is associated with a test suite.

For example:

```
static struct kunit_case example_test_cases[] = {
    KUNIT_CASE(example_test_foo),
    KUNIT_CASE(example_test_bar),
    KUNIT_CASE(example_test_baz),
    {}
};

static struct kunit_suite example_test_suite = {
    .name = "example",
    .init = example_test_init,
    .exit = example_test_exit,
    .suite_init = example_suite_init,
    .suite_exit = example_suite_exit,
    .test_cases = example_test_cases,
};
kunit_test_suite(example_test_suite);
```

In the above example, the test suite `example_test_suite` would first run `example_suite_init`, then run the test cases `example_test_foo`, `example_test_bar`, and `example_test_baz`. Each would have `example_test_init` called immediately before it and `example_test_exit` called immediately after it. Finally, `example_suite_exit` would be called after everything else. `kunit_test_suite(example_test_suite)` registers the test suite with the KUnit test framework.

Note: The `exit` and `suite_exit` functions will run even if `init` or `suite_init` fail. Make sure that they can handle any inconsistent state which may result from `init` or `suite_init` encountering errors or exiting early.

`kunit_test_suite(...)` is a macro which tells the linker to put the specified test suite in a special linker section so that it can be run by KUnit either after `late_init`, or when the test module is loaded (if the test was built as a module).

For more information, see [Test API](#).

16.5.3 Writing Tests For Other Architectures

It is better to write tests that run on UML to tests that only run under a particular architecture. It is better to write tests that run under QEMU or another easy to obtain (and monetarily free) software environment to a specific piece of hardware.

Nevertheless, there are still valid reasons to write a test that is architecture or hardware specific. For example, we might want to test code that really belongs in `arch/some-arch/*`. Even so, try to write the test so that it does not depend on physical hardware. Some of our test cases may not need hardware, only few tests actually require the hardware to test it. When hardware is not available, instead of disabling tests, we can skip them.

Now that we have narrowed down exactly what bits are hardware specific, the actual procedure for writing and running the tests is same as writing normal KUnit tests.

Important: We may have to reset hardware state. If this is not possible, we may only be able

to run one test case per invocation.

16.6 Common Patterns

16.6.1 Isolating Behavior

Unit testing limits the amount of code under test to a single unit. It controls what code gets run when the unit under test calls a function. Where a function is exposed as part of an API such that the definition of that function can be changed without affecting the rest of the code base. In the kernel, this comes from two constructs: classes, which are structs that contain function pointers provided by the implementer, and architecture-specific functions, which have definitions selected at compile time.

Classes

Classes are not a construct that is built into the C programming language; however, it is an easily derived concept. Accordingly, in most cases, every project that does not use a standardized object oriented library (like GNOME's GObject) has their own slightly different way of doing object oriented programming; the Linux kernel is no exception.

The central concept in kernel object oriented programming is the class. In the kernel, a *class* is a struct that contains function pointers. This creates a contract between *implementers* and *users* since it forces them to use the same function signature without having to call the function directly. To be a class, the function pointers must specify that a pointer to the class, known as a *class handle*, be one of the parameters. Thus the member functions (also known as *methods*) have access to member variables (also known as *fields*) allowing the same implementation to have multiple *instances*.

A class can be *overridden* by *child classes* by embedding the *parent class* in the child class. Then when the child class *method* is called, the child implementation knows that the pointer passed to it is of a parent contained within the child. Thus, the child can compute the pointer to itself because the pointer to the parent is always a fixed offset from the pointer to the child. This offset is the offset of the parent contained in the child struct. For example:

```
struct shape {
    int (*area)(struct shape *this);
};

struct rectangle {
    struct shape parent;
    int length;
    int width;
};

int rectangle_area(struct shape *this)
{
    struct rectangle *self = container_of(this, struct rectangle, parent);

    return self->length * self->width;
```

```
};

void rectangle_new(struct rectangle *self, int length, int width)
{
    self->parent.area = rectangle_area;
    self->length = length;
    self->width = width;
}
```

In this example, computing the pointer to the child from the pointer to the parent is done by `container_of`.

Faking Classes

In order to unit test a piece of code that calls a method in a class, the behavior of the method must be controllable, otherwise the test ceases to be a unit test and becomes an integration test.

A fake class implements a piece of code that is different than what runs in a production instance, but behaves identical from the standpoint of the callers. This is done to replace a dependency that is hard to deal with, or is slow. For example, implementing a fake EEPROM that stores the "contents" in an internal buffer. Assume we have a class that represents an EEPROM:

```
struct eeprom {
    ssize_t (*read)(struct eeprom *this, size_t offset, char *buffer, size_
    ↵t count);
    ssize_t (*write)(struct eeprom *this, size_t offset, const char_
    ↵*buffer, size_t count);
};
```

And we want to test code that buffers writes to the EEPROM:

```
struct eeprom_buffer {
    ssize_t (*write)(struct eeprom_buffer *this, const char *buffer, size_
    ↵t count);
    int flush(struct eeprom_buffer *this);
    size_t flush_count; /* Flushes when buffer exceeds flush_count. */
};

struct eeprom_buffer *new_eeprom_buffer(struct eeprom *eeprom);
void destroy_eeprom_buffer(struct eeprom *eeprom);
```

We can test this code by *faking out* the underlying EEPROM:

```
struct fake_eeprom {
    struct eeprom parent;
    char contents[FAKE_EEPROM_CONTENTS_SIZE];
};

ssize_t fake_eeprom_read(struct eeprom *parent, size_t offset, char *buffer, u
    ↵size_t count)
```

```

{
    struct fake_eeprom *this = container_of(parent, struct fake_eeprom,
parent);

    count = min(count, FAKE_EEPROM_CONTENTS_SIZE - offset);
    memcpy(buffer, this->contents + offset, count);

    return count;
}

ssize_t fake_eeprom_write(struct eeprom *parent, size_t offset, const char*
buffer, size_t count)
{
    struct fake_eeprom *this = container_of(parent, struct fake_eeprom,
parent);

    count = min(count, FAKE_EEPROM_CONTENTS_SIZE - offset);
    memcpy(this->contents + offset, buffer, count);

    return count;
}

void fake_eeprom_init(struct fake_eeprom *this)
{
    this->parent.read = fake_eeprom_read;
    this->parent.write = fake_eeprom_write;
    memset(this->contents, 0, FAKE_EEPROM_CONTENTS_SIZE);
}

```

We can now use it to test `struct eeprom_buffer`:

```

struct eeprom_buffer_test {
    struct fake_eeprom *fake_eeprom;
    struct eeprom_buffer *eeprom_buffer;
};

static void eeprom_buffer_test_does_not_write_until_flush(struct kunit *test)
{
    struct eeprom_buffer_test *ctx = test->priv;
    struct eeprom_buffer *eeprom_buffer = ctx->eeprom_buffer;
    struct fake_eeprom *fake_eeprom = ctx->fake_eeprom;
    char buffer[] = {0xff};

    eeprom_buffer->flush_count = SIZE_MAX;

    eeprom_buffer->write(eeprom_buffer, buffer, 1);
    KUNIT_EXPECT_EQ(test, fake_eeprom->contents[0], 0);

    eeprom_buffer->write(eeprom_buffer, buffer, 1);
    KUNIT_EXPECT_EQ(test, fake_eeprom->contents[1], 0);
}

```

```

        eeprom_buffer->flush(eeprom_buffer);
        KUNIT_EXPECT_EQ(test, fake_eeprom->contents[0], 0xff);
        KUNIT_EXPECT_EQ(test, fake_eeprom->contents[1], 0xff);
    }

static void eeprom_buffer_test_flushes_after_flush_count_met(struct kunit *_test)
{
    struct eeprom_buffer_test *ctx = test->priv;
    struct eeprom_buffer *eeprom_buffer = ctx->eeprom_buffer;
    struct fake_eeprom *fake_eeprom = ctx->fake_eeprom;
    char buffer[] = {0xff};

    eeprom_buffer->flush_count = 2;

    eeprom_buffer->write(eeprom_buffer, buffer, 1);
    KUNIT_EXPECT_EQ(test, fake_eeprom->contents[0], 0);

    eeprom_buffer->write(eeprom_buffer, buffer, 1);
    KUNIT_EXPECT_EQ(test, fake_eeprom->contents[0], 0xff);
    KUNIT_EXPECT_EQ(test, fake_eeprom->contents[1], 0xff);
}

static void eeprom_buffer_test_flushes_increments_of_flush_count(struct kunit *_test)
{
    struct eeprom_buffer_test *ctx = test->priv;
    struct eeprom_buffer *eeprom_buffer = ctx->eeprom_buffer;
    struct fake_eeprom *fake_eeprom = ctx->fake_eeprom;
    char buffer[] = {0xff, 0xff};

    eeprom_buffer->flush_count = 2;

    eeprom_buffer->write(eeprom_buffer, buffer, 1);
    KUNIT_EXPECT_EQ(test, fake_eeprom->contents[0], 0);

    eeprom_buffer->write(eeprom_buffer, buffer, 2);
    KUNIT_EXPECT_EQ(test, fake_eeprom->contents[0], 0xff);
    KUNIT_EXPECT_EQ(test, fake_eeprom->contents[1], 0xff);
    /* Should have only flushed the first two bytes. */
    KUNIT_EXPECT_EQ(test, fake_eeprom->contents[2], 0);
}

static int eeprom_buffer_test_init(struct kunit *test)
{
    struct eeprom_buffer_test *ctx;

    ctx = kunit_kzalloc(test, sizeof(*ctx), GFP_KERNEL);
    KUNIT_ASSERT_NOT_ERR_OR_NULL(test, ctx);
}

```

```

        ctx->fake_eeprom = kunit_kzalloc(test, sizeof(*ctx->fake_eeprom), GFP_
↪KERNEL);
        KUNIT_ASSERT_NOT_ERR_OR_NULL(test, ctx->fake_eeprom);
        fake_eeprom_init(ctx->fake_eeprom);

        ctx->eprom_buffer = new_eeprom_buffer(&ctx->fake_eeprom->parent);
        KUNIT_ASSERT_NOT_ERR_OR_NULL(test, ctx->eprom_buffer);

        test->priv = ctx;

        return 0;
}

static void eeprom_buffer_test_exit(struct kunit *test)
{
    struct eeprom_buffer_test *ctx = test->priv;

    destroy_eeprom_buffer(ctx->eprom_buffer);
}

```

16.6.2 Testing Against Multiple Inputs

Testing just a few inputs is not enough to ensure that the code works correctly, for example: testing a hash function.

We can write a helper macro or function. The function is called for each input. For example, to test `shalsum(1)`, we can write:

```

#define TEST_SHA1(in, want) \
    shalsum(in, out); \
    KUNIT_EXPECT_STREQ_MSG(test, out, want, "shalsum(%s)", in);

char out[40];
TEST_SHA1("hello world", "2aae6c35c94fcfb415dbe95f408b9ce91ee846ed");
TEST_SHA1("hello world!", "430ce34d020724ed75a196dfc2ad67c77772d169");

```

Note the use of the `_MSG` version of `KUNIT_EXPECT_STREQ` to print a more detailed error and make the assertions clearer within the helper macros.

The `_MSG` variants are useful when the same expectation is called multiple times (in a loop or helper function) and thus the line number is not enough to identify what failed, as shown below.

In complicated cases, we recommend using a *table-driven test* compared to the helper macro variation, for example:

```

int i;
char out[40];

struct sha1_test_case {
    const char *str;
    const char *sha1;
}

```

```

};

struct sha1_test_case cases[] = {
{
    .str = "hello world",
    .sha1 = "2aae6c35c94fcfb415dbe95f408b9ce91ee846ed",
},
{
    .str = "hello world!",
    .sha1 = "430ce34d020724ed75a196dfc2ad67c77772d169",
},
};

for (i = 0; i < ARRAY_SIZE(cases); ++i) {
    shalsum(cases[i].str, out);
    KUNIT_EXPECT_STREQ_MSG(test, out, cases[i].sha1,
                           "shalsum(%s)", cases[i].str);
}

```

There is more boilerplate code involved, but it can:

- be more readable when there are multiple inputs/outputs (due to field names).
 - For example, see `fs/ext4/inode-test.c`.
- reduce duplication if test cases are shared across multiple tests.
 - For example: if we want to test `sha256sum`, we could add a `sha256` field and reuse `cases`.
- be converted to a “parameterized test”.

Parameterized Testing

The table-driven testing pattern is common enough that KUnit has special support for it.

By reusing the same `cases` array from above, we can write the test as a “parameterized test” with the following.

```

// This is copy-pasted from above.
struct sha1_test_case {
    const char *str;
    const char *sha1;
};

const struct sha1_test_case cases[] = {
{
    .str = "hello world",
    .sha1 = "2aae6c35c94fcfb415dbe95f408b9ce91ee846ed",
},
{
    .str = "hello world!",
    .sha1 = "430ce34d020724ed75a196dfc2ad67c77772d169",
},
};

```

```

// Creates `shal_gen_params()` to iterate over `cases` while using
// the struct member `str` for the case description.
KUNIT_ARRAY_PARAM_DESC(shal, cases, str);

// Looks no different from a normal test.
static void shal_test(struct kunit *test)
{
    // This function can just contain the body of the for-loop.
    // The former `cases[i]` is accessible under test->param_value.
    char out[40];
    struct shal_test_case *test_param = (struct shal_test_case *) (test->
→param_value);

    shalsum(test_param->str, out);
    KUNIT_EXPECT_STREQ_MSG(test, out, test_param->shal,
                           "shalsum(%s)", test_param->str);
}

// Instead of KUNIT_CASE, we use KUNIT_CASE_PARAM and pass in the
// function declared by KUNIT_ARRAY_PARAM or KUNIT_ARRAY_PARAM_DESC.
static struct kunit_case shal_test_cases[] = {
    KUNIT_CASE_PARAM(shal_test, shal_gen_params),
    {}
};

```

16.6.3 Allocating Memory

Where you might use kzalloc, you can instead use kunit_kzalloc as KUnit will then ensure that the memory is freed once the test completes.

This is useful because it lets us use the KUNIT_ASSERT_EQ macros to exit early from a test without having to worry about remembering to call kfree. For example:

```

void example_test_allocation(struct kunit *test)
{
    char *buffer = kunit_kzalloc(test, 16, GFP_KERNEL);
    /* Ensure allocation succeeded. */
    KUNIT_ASSERT_NOT_ERR_OR_NULL(test, buffer);

    KUNIT_ASSERT_STREQ(test, buffer, "");
}

```

16.6.4 Registering Cleanup Actions

If you need to perform some cleanup beyond simple use of `kunit_kzalloc`, you can register a custom “deferred action”, which is a cleanup function run when the test exits (whether cleanly, or via a failed assertion).

Actions are simple functions with no return value, and a single `void*` context argument, and fulfill the same role as “cleanup” functions in Python and Go tests, “defer” statements in languages which support them, and (in some cases) destructors in RAII languages.

These are very useful for unregistering things from global lists, closing files or other resources, or freeing resources.

For example:

```
static void cleanup_device(void *ctx)
{
    struct device *dev = (struct device *)ctx;

    device_unregister(dev);
}

void example_device_test(struct kunit *test)
{
    struct my_device dev;

    device_register(&dev);

    kunit_add_action(test, &cleanup_device, &dev);
}
```

Note that, for functions like `device_unregister` which only accept a single pointer-sized argument, it's possible to automatically generate a wrapper with the `KUNIT_DEFINE_ACTION_WRAPPER()` macro, for example:

```
KUNIT_DEFINE_ACTION_WRAPPER(device_unregister, device_unregister_wrapper,
    ~struct device *);
kunit_add_action(test, &device_unregister_wrapper, &dev);
```

You should do this in preference to manually casting to the `kunit_action_t` type, as casting function pointers will break Control Flow Integrity (CFI).

`kunit_add_action` can fail if, for example, the system is out of memory. You can use `kunit_add_action_or_reset` instead which runs the action immediately if it cannot be deferred.

If you need more control over when the cleanup function is called, you can trigger it early using `kunit_release_action`, or cancel it entirely with `kunit_remove_action`.

16.6.5 Testing Static Functions

If we do not want to expose functions or variables for testing, one option is to conditionally export the used symbol. For example:

```
/* In my_file.c */

VISIBLE_IF_KUNIT int do_interesting_thing();
EXPORT_SYMBOL_IF_KUNIT(do_interesting_thing);

/* In my_file.h */

#if IS_ENABLED(CONFIG_KUNIT)
    int do_interesting_thing(void);
#endif
```

Alternatively, you could conditionally #include the test file at the end of your .c file. For example:

```
/* In my_file.c */

static int do_interesting_thing();

#ifndef CONFIG_MY_KUNIT_TEST
#include "my_kunit_test.c"
#endif
```

16.6.6 Injecting Test-Only Code

Similar to as shown above, we can add test-specific logic. For example:

```
/* In my_file.h */

#ifndef CONFIG_MY_KUNIT_TEST
/* Defined in my_kunit_test.c */
void test_only_hook(void);
#else
void test_only_hook(void) { }
#endif
```

This test-only code can be made more useful by accessing the current kunit_test as shown in next section: *Accessing The Current Test*.

16.6.7 Accessing The Current Test

In some cases, we need to call test-only code from outside the test file. This is helpful, for example, when providing a fake implementation of a function, or to fail any current test from within an error handler. We can do this via the `kunit_test` field in `task_struct`, which we can access using the `kunit_get_current_test()` function in `kunit/test-bug.h`.

`kunit_get_current_test()` is safe to call even if KUnit is not enabled. If KUnit is not enabled, or if no test is running in the current task, it will return `NULL`. This compiles down to either a no-op or a static key check, so will have a negligible performance impact when no test is running.

The example below uses this to implement a "mock" implementation of a function, `foo`:

```
#include <kunit/test-bug.h> /* for kunit_get_current_test */

struct test_data {
    int foo_result;
    int want_foo_called_with;
};

static int fake_foo(int arg)
{
    struct kunit *test = kunit_get_current_test();
    struct test_data *test_data = test->priv;

    KUNIT_EXPECT_EQ(test, test_data->want_foo_called_with, arg);
    return test_data->foo_result;
}

static void example_simple_test(struct kunit *test)
{
    /* Assume priv (private, a member used to pass test data from
     * the init function) is allocated in the suite's .init */
    struct test_data *test_data = test->priv;

    test_data->foo_result = 42;
    test_data->want_foo_called_with = 1;

    /* In a real test, we'd probably pass a pointer to fake_foo somewhere
     * like an ops struct, etc. instead of calling it directly. */
    KUNIT_EXPECT_EQ(test, fake_foo(1), 42);
}
```

In this example, we are using the `priv` member of `struct kunit` as a way of passing data to the test from the `init` function. In general `priv` is pointer that can be used for any user data. This is preferred over static variables, as it avoids concurrency issues.

Had we wanted something more flexible, we could have used a named `kunit_resource`. Each test can have multiple resources which have string names providing the same flexibility as a `priv` member, but also, for example, allowing helper functions to create resources without conflicting with each other. It is also possible to define a clean up function for each resource, making it easy to avoid resource leaks. For more information, see [Resource API](#).

16.6.8 Failing The Current Test

If we want to fail the current test, we can use `kunit_fail_current_test(fmt, args...)` which is defined in `<kunit/test-bug.h>` and does not require pulling in `<kunit/test.h>`. For example, we have an option to enable some extra debug checks on some data structures as shown below:

```
#include <kunit/test-bug.h>

#ifndef CONFIG_EXTRA_DEBUG_CHECKS
static void validate_my_data(struct data *data)
{
    if (is_valid(data))
        return;

    kunit_fail_current_test("data %p is invalid", data);

    /* Normal, non-KUnit, error reporting code here. */
}
#else
static void my_debug_function(void) { }
#endif
```

`kunit_fail_current_test()` is safe to call even if KUnit is not enabled. If KUnit is not enabled, or if no test is running in the current task, it will do nothing. This compiles down to either a no-op or a static key check, so will have a negligible performance impact when no test is running.

16.6.9 Managing Fake Devices and Drivers

When testing drivers or code which interacts with drivers, many functions will require a `struct device` or `struct device_driver`. In many cases, setting up a real device is not required to test any given function, so a fake device can be used instead.

KUnit provides helper functions to create and manage these fake devices, which are internally of type `struct kunit_device`, and are attached to a special `kunit_bus`. These devices support managed device resources (`devres`), as described in `Documentation/driver-api/driver-model/devres.rst`

To create a KUnit-managed `struct device_driver`, use `kunit_driver_create()`, which will create a driver with the given name, on the `kunit_bus`. This driver will automatically be destroyed when the corresponding test finishes, but can also be manually destroyed with `driver_unregister()`.

To create a fake device, use the `kunit_device_register()`, which will create and register a device, using a new KUnit-managed driver created with `kunit_driver_create()`. To provide a specific, non-KUnit-managed driver, use `kunit_device_register_with_driver()` instead. Like with managed drivers, KUnit-managed fake devices are automatically cleaned up when the test finishes, but can be manually cleaned up early with `kunit_device_unregister()`.

The KUnit devices should be used in preference to `root_device_register()`, and instead of `platform_device_register()` in cases where the device is not otherwise a platform device.

For example:

```
#include <kunit/device.h>

static void test_my_device(struct kunit *test)
{
    struct device *fake_device;
    const char *dev_managed_string;

    // Create a fake device.
    fake_device = kunit_device_register(test, "my_device");
    KUNIT_ASSERT_NOT_ERR_OR_NULL(test, fake_device)

    // Pass it to functions which need a device.
    dev_managed_string = devm_kstrdup(fake_device, "Hello, World!");

    // Everything is cleaned up automatically when the test ends.
}
```

16.7 API Reference

16.7.1 Test API

This file documents all of the standard testing API.

enum kunit_status

Type of result for a test or test suite

Constants

KUNIT_SUCCESS

Denotes the test suite has not failed nor been skipped

KUNIT_FAILURE

Denotes the test has failed.

KUNIT_SKIPPED

Denotes the test has been skipped.

struct kunit_case

represents an individual test case.

Definition:

```
struct kunit_case {
    void (*run_case)(struct kunit *test);
    const char *name;
    const void* (*generate_params)(const void *prev, char *desc);
    struct kunit_attributes attr;
};
```

Members

run_case

the function representing the actual test case.

name

the name of the test case.

generate_params

the generator function for parameterized tests.

attr

the attributes associated with the test

Description

A test case is a function with the signature, `void (*)(struct kunit *)` that makes expectations and assertions (see `KUNIT_EXPECT_TRUE()` and `KUNIT_ASSERT_TRUE()`) about code under test. Each test case is associated with a `struct kunit_suite` and will be run after the suite's init function and followed by the suite's exit function.

A test case should be static and should only be created with the `KUNIT_CASE()` macro; additionally, every array of test cases should be terminated with an empty test case.

```
void add_test_basic(struct kunit *test)
{
    KUNIT_EXPECT_EQ(test, 1, add(1, 0));
    KUNIT_EXPECT_EQ(test, 2, add(1, 1));
    KUNIT_EXPECT_EQ(test, 0, add(-1, 1));
    KUNIT_EXPECT_EQ(test, INT_MAX, add(0, INT_MAX));
    KUNIT_EXPECT_EQ(test, -1, add(INT_MAX, INT_MIN));
}

static struct kunit_case example_test_cases[] = {
    KUNIT_CASE(add_test_basic),
    {}
};
```

Example**KUNIT_CASE**

`KUNIT_CASE (test_name)`

A helper for creating a `struct kunit_case`

Parameters**test_name**

a reference to a test case function.

Description

Takes a symbol for a function representing a test case and creates a `struct kunit_case` object from it. See the documentation for `struct kunit_case` for an example on how to use it.

KUNIT_CASE_ATTR

`KUNIT_CASE_ATTR (test_name, attributes)`

A helper for creating a `struct kunit_case` with attributes

Parameters

test_name

a reference to a test case function.

attributes

a reference to a struct kunit_attributes object containing test attributes

KUNIT_CASE_SLOW

KUNIT_CASE_SLOW (test_name)

A helper for creating a *struct kunit_case* with the slow attribute

Parameters**test_name**

a reference to a test case function.

KUNIT_CASE_PARAM

KUNIT_CASE_PARAM (test_name, gen_params)

A helper for creation a parameterized *struct kunit_case*

Parameters**test_name**

a reference to a test case function.

gen_params

a reference to a parameter generator function.

Description

The generator function:

```
const void* gen_params(const void *prev, char *desc)
```

is used to lazily generate a series of arbitrarily typed values that fit into a void*. The argument **prev** is the previously returned value, which should be used to derive the next value; **prev** is set to NULL on the initial generator call. When no more values are available, the generator must return NULL. Optionally write a string into **desc** (size of KUNIT_PARAM_DESC_SIZE) describing the parameter.

KUNIT_CASE_PARAM_ATTR

KUNIT_CASE_PARAM_ATTR (test_name, gen_params, attributes)

A helper for creating a parameterized *struct kunit_case* with attributes

Parameters**test_name**

a reference to a test case function.

gen_params

a reference to a parameter generator function.

attributes

a reference to a struct kunit_attributes object containing test attributes

struct kunit_suite

describes a related collection of [struct kunit_case](#)

Definition:

```
struct kunit_suite {
    const char name[256];
    int (*suite_init)(struct kunit_suite *suite);
    void (*suite_exit)(struct kunit_suite *suite);
    int (*init)(struct kunit *test);
    void (*exit)(struct kunit *test);
    struct kunit_case *test_cases;
    struct kunit_attributes attr;
};
```

Members**name**

the name of the test. Purely informational.

suite_init

called once per test suite before the test cases.

suite_exit

called once per test suite after all test cases.

init

called before every test case.

exit

called after every test case.

test_cases

a null terminated array of test cases.

attr

the attributes associated with the test suite

Description

A kunit_suite is a collection of related [struct kunit_case](#)'s, such that **init** is called before every test case and **exit** is called after every test case, similar to the notion of a *test fixture* or a *test class* in other unit testing frameworks like JUnit or Googletest.

Note that **exit** and **suite_exit** will run even if **init** or **suite_init** fail: make sure they can handle any inconsistent state which may result.

Every [struct kunit_case](#) must be associated with a kunit_suite for KUnit to run it.

struct kunit

represents a running instance of a test.

Definition:

```
struct kunit {
    void *priv;
};
```

Members

priv

for user to store arbitrary data. Commonly used to pass data created in the init function (see *struct kunit_suite*).

Description

Used to store information about the current context under which the test is running. Most of this data is private and should only be accessed indirectly via public functions; the one exception is **priv** which can be used by the test writer to store arbitrary data.

kunit_test_suites

`kunit_test_suites (__suites...)`

used to register one or more *struct kunit_suite* with KUnit.

Parameters

__suites...

a statically allocated list of *struct kunit_suite*.

Description

Registers **suites** with the test framework. This is done by placing the array of *struct kunit_suite* * in the .kunit_test_suites ELF section.

When builtin, KUnit tests are all run via the executor at boot, and when built as a module, they run on module load.

kunit_test_init_section_suites

`kunit_test_init_section_suites (__suites...)`

used to register one or more *struct kunit_suite* containing init functions or init data.

Parameters

__suites...

a statically allocated list of *struct kunit_suite*.

Description

This functions similar to *kunit_test_suites()* except that it compiles the list of suites during init phase.

This macro also suffixes the array and suite declarations it makes with _probe; so that modpost suppresses warnings about referencing init data for symbols named in this manner.

Also, do not mark the suite or test case structs with __initdata because they will be used after the init phase with debugfs.

Note

these init tests are not able to be run after boot so there is no "run" debugfs file generated for these tests.

`void *kunit_kmalloc_array(struct kunit *test, size_t n, size_t size, gfp_t gfp)`

Like kmalloc_array() except the allocation is *test managed*.

Parameters

struct kunit *test
 The test context object.

size_t n
 number of elements.

size_t size
 The size in bytes of the desired memory.

gfp_t gfp
 flags passed to underlying kmalloc().

Description

Just like `kmalloc_array(...)`, except the allocation is managed by the test case and is automatically cleaned up after the test case concludes. See [`kunit_add_action\(\)`](#) for more information.

Note that some internal context data is also allocated with GFP_KERNEL, regardless of the gfp passed in.

`void *kunit_kmalloc(struct kunit *test, size_t size, gfp_t gfp)`
 Like kmalloc() except the allocation is *test managed*.

Parameters

struct kunit *test
 The test context object.

size_t size
 The size in bytes of the desired memory.

gfp_t gfp
 flags passed to underlying kmalloc().

Description

See kmalloc() and [`kunit_kmalloc_array\(\)`](#) for more information.

Note that some internal context data is also allocated with GFP_KERNEL, regardless of the gfp passed in.

`void kunit_kfree(struct kunit *test, const void *ptr)`
 Like kfree except for allocations managed by KUnit.

Parameters

struct kunit *test
 The test case to which the resource belongs.

const void *ptr
 The memory allocation to free.

`void *kunit_kzalloc(struct kunit *test, size_t size, gfp_t gfp)`
 Just like [`kunit_kmalloc\(\)`](#), but zeroes the allocation.

Parameters

struct kunit *test
 The test context object.

size_t size
 The size in bytes of the desired memory.

gfp_t gfp

flags passed to underlying kmalloc().

Description

See kzalloc() and [*kunit_kmalloc_array\(\)*](#) for more information.

void *kunit_kcalloc(struct kunit *test, size_t n, size_t size, gfp_t gfp)

Just like [*kunit_kmalloc_array\(\)*](#), but zeroes the allocation.

Parameters

struct kunit *test

The test context object.

size_t n

number of elements.

size_t size

The size in bytes of the desired memory.

gfp_t gfp

flags passed to underlying kmalloc().

Description

See kcalloc() and [*kunit_kmalloc_array\(\)*](#) for more information.

kunit_mark_skipped

kunit_mark_skipped (test_or_suite, fmt, ...)

Marks **test_or_suite** as skipped

Parameters

test_or_suite

The test context object.

fmt

A printk() style format string.

...

variable arguments

Description

Marks the test as skipped. **fmt** is given output as the test status comment, typically the reason the test was skipped.

Test execution continues after [*kunit_mark_skipped\(\)*](#) is called.

kunit_skip

kunit_skip (test_or_suite, fmt, ...)

Marks **test_or_suite** as skipped

Parameters

test_or_suite

The test context object.

fmt

A printk() style format string.

...

variable arguments

Description

Skips the test. **fmt** is given output as the test status comment, typically the reason the test was skipped.

Test execution is halted after [*kunit_skip\(\)*](#) is called.

kunit_info

```
kunit_info (test, fmt, ...)
```

Prints an INFO level message associated with **test**.

Parameters**test**

The test context object.

fmt

A printk() style format string.

...

variable arguments

Description

Prints an info level message associated with the test suite being run. Takes a variable number of format parameters just like printk().

kunit_warn

```
kunit_warn (test, fmt, ...)
```

Prints a WARN level message associated with **test**.

Parameters**test**

The test context object.

fmt

A printk() style format string.

...

variable arguments

Description

Prints a warning level message.

kunit_err

```
kunit_err (test, fmt, ...)
```

Prints an ERROR level message associated with **test**.

Parameters

test

The test context object.

fmt

A printk() style format string.

...

variable arguments

Description

Prints an error level message.

KUNIT_SUCCEED**KUNIT_SUCCEED (test)**

A no-op expectation. Only exists for code clarity.

Parameters**test**

The test context object.

Description

The opposite of [KUNIT_FAIL\(\)](#), it is an expectation that cannot fail. In other words, it does nothing and only exists for code clarity. See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_FAIL**KUNIT_FAIL (test, fmt, ...)**

Always causes a test to fail when evaluated.

Parameters**test**

The test context object.

fmt

an informational message to be printed when the assertion is made.

...

string format arguments.

Description

The opposite of [KUNIT_SUCCEED\(\)](#), it is an expectation that always fails. In other words, it always results in a failed expectation, and consequently always causes the test case to fail when evaluated. See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_EXPECT_TRUE**KUNIT_EXPECT_TRUE (test, condition)**

Causes a test failure when the expression is not true.

Parameters**test**

The test context object.

condition

an arbitrary boolean expression. The test fails when this does not evaluate to true.

Description

This and expectations of the form `KUNIT_EXPECT_*` will cause the test case to fail when the specified condition is not met; however, it will not prevent the test case from continuing to run; this is otherwise known as an *expectation failure*.

KUNIT_EXPECT_FALSE

`KUNIT_EXPECT_FALSE (test, condition)`

Makes a test failure when the expression is not false.

Parameters**test**

The test context object.

condition

an arbitrary boolean expression. The test fails when this does not evaluate to false.

Description

Sets an expectation that **condition** evaluates to false. See `KUNIT_EXPECT_TRUE()` for more information.

KUNIT_EXPECT_EQ

`KUNIT_EXPECT_EQ (test, left, right)`

Sets an expectation that **left** and **right** are equal.

Parameters**test**

The test context object.

left

an arbitrary expression that evaluates to a primitive C type.

right

an arbitrary expression that evaluates to a primitive C type.

Description

Sets an expectation that the values that **left** and **right** evaluate to are equal. This is semantically equivalent to `KUNIT_EXPECT_TRUE(test, (left) == (right))`. See `KUNIT_EXPECT_TRUE()` for more information.

KUNIT_EXPECT_PTR_EQ

`KUNIT_EXPECT_PTR_EQ (test, left, right)`

Expects that pointers **left** and **right** are equal.

Parameters**test**

The test context object.

left

an arbitrary expression that evaluates to a pointer.

right

an arbitrary expression that evaluates to a pointer.

Description

Sets an expectation that the values that **left** and **right** evaluate to are equal. This is semantically equivalent to `KUNIT_EXPECT_TRUE(test, (left) == (right))`. See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_EXPECT_NE

`KUNIT_EXPECT_NE (test, left, right)`

An expectation that **left** and **right** are not equal.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a primitive C type.

right

an arbitrary expression that evaluates to a primitive C type.

Description

Sets an expectation that the values that **left** and **right** evaluate to are not equal. This is semantically equivalent to `KUNIT_EXPECT_TRUE(test, (left) != (right))`. See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_EXPECT_PTR_NE

`KUNIT_EXPECT_PTR_NE (test, left, right)`

Expects that pointers **left** and **right** are not equal.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a pointer.

right

an arbitrary expression that evaluates to a pointer.

Description

Sets an expectation that the values that **left** and **right** evaluate to are not equal. This is semantically equivalent to `KUNIT_EXPECT_TRUE(test, (left) != (right))`. See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_EXPECT_LT

`KUNIT_EXPECT_LT (test, left, right)`

An expectation that **left** is less than **right**.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a primitive C type.

right

an arbitrary expression that evaluates to a primitive C type.

Description

Sets an expectation that the value that **left** evaluates to is less than the value that **right** evaluates to. This is semantically equivalent to KUNIT_EXPECT_TRUE(**test**, (**left**) < (**right**)). See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_EXPECT_LE

`KUNIT_EXPECT_LE (test, left, right)`

An expectation that **left** is less than or equal to **right**.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a primitive C type.

right

an arbitrary expression that evaluates to a primitive C type.

Description

Sets an expectation that the value that **left** evaluates to is less than or equal to the value that **right** evaluates to. Semantically this is equivalent to KUNIT_EXPECT_TRUE(**test**, (**left**) <= (**right**)). See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_EXPECT_GT

`KUNIT_EXPECT_GT (test, left, right)`

An expectation that **left** is greater than **right**.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a primitive C type.

right

an arbitrary expression that evaluates to a primitive C type.

Description

Sets an expectation that the value that **left** evaluates to is greater than the value that **right** evaluates to. This is semantically equivalent to KUNIT_EXPECT_TRUE(**test**, (**left**) > (**right**)). See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_EXPECT_GE

`KUNIT_EXPECT_GE (test, left, right)`

Expects that **left** is greater than or equal to **right**.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a primitive C type.

right

an arbitrary expression that evaluates to a primitive C type.

Description

Sets an expectation that the value that **left** evaluates to is greater than the value that **right** evaluates to. This is semantically equivalent to `KUNIT_EXPECT_TRUE(test, (left) >= (right))`. See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_EXPECT_STREQ

`KUNIT_EXPECT_STREQ (test, left, right)`

Expects that strings **left** and **right** are equal.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a null terminated string.

right

an arbitrary expression that evaluates to a null terminated string.

Description

Sets an expectation that the values that **left** and **right** evaluate to are equal. This is semantically equivalent to `KUNIT_EXPECT_TRUE(test, !strcmp((left), (right)))`. See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_EXPECT_STRNEQ

`KUNIT_EXPECT_STRNEQ (test, left, right)`

Expects that strings **left** and **right** are not equal.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a null terminated string.

right

an arbitrary expression that evaluates to a null terminated string.

Description

Sets an expectation that the values that **left** and **right** evaluate to are not equal. This is semantically equivalent to KUNIT_EXPECT_TRUE(test, strcmp((left), (right))). See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_EXPECT_MEMEQ

`KUNIT_EXPECT_MEMEQ (test, left, right, size)`

Expects that the first **size** bytes of **left** and **right** are equal.

Parameters

test

The test context object.

left

An arbitrary expression that evaluates to the specified size.

right

An arbitrary expression that evaluates to the specified size.

size

Number of bytes compared.

Description

Sets an expectation that the values that **left** and **right** evaluate to are equal. This is semantically equivalent to KUNIT_EXPECT_TRUE(test, !memcmp((left), (right), (size))). See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

Although this expectation works for any memory block, it is not recommended for comparing more structured data, such as structs. This expectation is recommended for comparing, for example, data arrays.

KUNIT_EXPECT_MEMNEQ

`KUNIT_EXPECT_MEMNEQ (test, left, right, size)`

Expects that the first **size** bytes of **left** and **right** are not equal.

Parameters

test

The test context object.

left

An arbitrary expression that evaluates to the specified size.

right

An arbitrary expression that evaluates to the specified size.

size

Number of bytes compared.

Description

Sets an expectation that the values that **left** and **right** evaluate to are not equal. This is semantically equivalent to KUNIT_EXPECT_TRUE(test, memcmp((left), (right), (size))). See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

Although this expectation works for any memory block, it is not recommended for comparing more structured data, such as structs. This expectation is recommended for comparing, for example, data arrays.

KUNIT_EXPECT_NULL

`KUNIT_EXPECT_NULL (test, ptr)`

Expects that **ptr** is null.

Parameters

test

The test context object.

ptr

an arbitrary pointer.

Description

Sets an expectation that the value that **ptr** evaluates to is null. This is semantically equivalent to `KUNIT_EXPECT_PTR_EQ(test, ptr, NULL)`. See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_EXPECT_NOT_NULL

`KUNIT_EXPECT_NOT_NULL (test, ptr)`

Expects that **ptr** is not null.

Parameters

test

The test context object.

ptr

an arbitrary pointer.

Description

Sets an expectation that the value that **ptr** evaluates to is not null. This is semantically equivalent to `KUNIT_EXPECT_PTR_NE(test, ptr, NULL)`. See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_EXPECT_NOT_ERR_OR_NULL

`KUNIT_EXPECT_NOT_ERR_OR_NULL (test, ptr)`

Expects that **ptr** is not null and not err.

Parameters

test

The test context object.

ptr

an arbitrary pointer.

Description

Sets an expectation that the value that **ptr** evaluates to is not null and not an errno stored in a pointer. This is semantically equivalent to `KUNIT_EXPECT_TRUE(test, !IS_ERR_OR_NULL(ptr))`. See [KUNIT_EXPECT_TRUE\(\)](#) for more information.

KUNIT_ASSERT_TRUE

`KUNIT_ASSERT_TRUE (test, condition)`

Sets an assertion that **condition** is true.

Parameters

test

The test context object.

condition

an arbitrary boolean expression. The test fails and aborts when this does not evaluate to true.

Description

This and assertions of the form `KUNIT_ASSERT_*` will cause the test case to fail *and immediately abort* when the specified condition is not met. Unlike an expectation failure, it will prevent the test case from continuing to run; this is otherwise known as an *assertion failure*.

KUNIT_ASSERT_FALSE

`KUNIT_ASSERT_FALSE (test, condition)`

Sets an assertion that **condition** is false.

Parameters

test

The test context object.

condition

an arbitrary boolean expression.

Description

Sets an assertion that the value that **condition** evaluates to is false. This is the same as `KUNIT_EXPECT_FALSE()`, except it causes an assertion failure (see `KUNIT_ASSERT_TRUE()`) when the assertion is not met.

KUNIT_ASSERT_EQ

`KUNIT_ASSERT_EQ (test, left, right)`

Sets an assertion that **left** and **right** are equal.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a primitive C type.

right

an arbitrary expression that evaluates to a primitive C type.

Description

Sets an assertion that the values that **left** and **right** evaluate to are equal. This is the same as `KUNIT_EXPECT_EQ()`, except it causes an assertion failure (see `KUNIT_ASSERT_TRUE()`) when the assertion is not met.

KUNIT_ASSERT_PTR_EQ

KUNIT_ASSERT_PTR_EQ (*test*, *left*, *right*)

Asserts that pointers **left** and **right** are equal.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a pointer.

right

an arbitrary expression that evaluates to a pointer.

Description

Sets an assertion that the values that **left** and **right** evaluate to are equal. This is the same as [**KUNIT_EXPECT_EQ\(\)**](#), except it causes an assertion failure (see [**KUNIT_ASSERT_TRUE\(\)**](#)) when the assertion is not met.

KUNIT_ASSERT_NE

KUNIT_ASSERT_NE (*test*, *left*, *right*)

An assertion that **left** and **right** are not equal.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a primitive C type.

right

an arbitrary expression that evaluates to a primitive C type.

Description

Sets an assertion that the values that **left** and **right** evaluate to are not equal. This is the same as [**KUNIT_EXPECT_NE\(\)**](#), except it causes an assertion failure (see [**KUNIT_ASSERT_TRUE\(\)**](#)) when the assertion is not met.

KUNIT_ASSERT_PTR_NE

KUNIT_ASSERT_PTR_NE (*test*, *left*, *right*)

Asserts that pointers **left** and **right** are not equal. [**KUNIT_ASSERT_PTR_EQ\(\)**](#) - Asserts that pointers **left** and **right** are equal.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a pointer.

right

an arbitrary expression that evaluates to a pointer.

Description

Sets an assertion that the values that **left** and **right** evaluate to are not equal. This is the same as [`KUNIT_EXPECT_NE\(\)`](#), except it causes an assertion failure (see [`KUNIT_ASSERT_TRUE\(\)`](#)) when the assertion is not met.

`KUNIT_ASSERT_LT`

`KUNIT_ASSERT_LT (test, left, right)`

An assertion that **left** is less than **right**.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a primitive C type.

right

an arbitrary expression that evaluates to a primitive C type.

Description

Sets an assertion that the value that **left** evaluates to is less than the value that **right** evaluates to. This is the same as [`KUNIT_EXPECT_LT\(\)`](#), except it causes an assertion failure (see [`KUNIT_ASSERT_TRUE\(\)`](#)) when the assertion is not met.

`KUNIT_ASSERT_LE`

`KUNIT_ASSERT_LE (test, left, right)`

An assertion that **left** is less than or equal to **right**.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a primitive C type.

right

an arbitrary expression that evaluates to a primitive C type.

Description

Sets an assertion that the value that **left** evaluates to is less than or equal to the value that **right** evaluates to. This is the same as [`KUNIT_EXPECT_LE\(\)`](#), except it causes an assertion failure (see [`KUNIT_ASSERT_TRUE\(\)`](#)) when the assertion is not met.

`KUNIT_ASSERT_GT`

`KUNIT_ASSERT_GT (test, left, right)`

An assertion that **left** is greater than **right**.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a primitive C type.

right

an arbitrary expression that evaluates to a primitive C type.

Description

Sets an assertion that the value that **left** evaluates to is greater than the value that **right** evaluates to. This is the same as [`KUNIT_EXPECT_GT\(\)`](#), except it causes an assertion failure (see [`KUNIT_ASSERT_TRUE\(\)`](#)) when the assertion is not met.

KUNIT_ASSERT_GE

`KUNIT_ASSERT_GE (test, left, right)`

Assertion that **left** is greater than or equal to **right**.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a primitive C type.

right

an arbitrary expression that evaluates to a primitive C type.

Description

Sets an assertion that the value that **left** evaluates to is greater than the value that **right** evaluates to. This is the same as [`KUNIT_EXPECT_GE\(\)`](#), except it causes an assertion failure (see [`KUNIT_ASSERT_TRUE\(\)`](#)) when the assertion is not met.

KUNIT_ASSERT_STREQ

`KUNIT_ASSERT_STREQ (test, left, right)`

An assertion that strings **left** and **right** are equal.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a null terminated string.

right

an arbitrary expression that evaluates to a null terminated string.

Description

Sets an assertion that the values that **left** and **right** evaluate to are equal. This is the same as [`KUNIT_EXPECT_STREQ\(\)`](#), except it causes an assertion failure (see [`KUNIT_ASSERT_TRUE\(\)`](#)) when the assertion is not met.

KUNIT_ASSERT_STRNEQ

`KUNIT_ASSERT_STRNEQ (test, left, right)`

Expects that strings **left** and **right** are not equal.

Parameters

test

The test context object.

left

an arbitrary expression that evaluates to a null terminated string.

right

an arbitrary expression that evaluates to a null terminated string.

Description

Sets an expectation that the values that **left** and **right** evaluate to are not equal. This is semantically equivalent to `KUNIT_ASSERT_TRUE(test, strcmp((left), (right)))`. See [KUNIT_ASSERT_TRUE\(\)](#) for more information.

KUNIT_ASSERT_NULL

`KUNIT_ASSERT_NULL (test, ptr)`

Asserts that pointers **ptr** is null.

Parameters

test

The test context object.

ptr

an arbitrary pointer.

Description

Sets an assertion that the values that **ptr** evaluates to is null. This is the same as [KUNIT_EXPECT_NULL\(\)](#), except it causes an assertion failure (see [KUNIT_ASSERT_TRUE\(\)](#)) when the assertion is not met.

KUNIT_ASSERT_NOT_NULL

`KUNIT_ASSERT_NOT_NULL (test, ptr)`

Asserts that pointers **ptr** is not null.

Parameters

test

The test context object.

ptr

an arbitrary pointer.

Description

Sets an assertion that the values that **ptr** evaluates to is not null. This is the same as [KUNIT_EXPECT_NOT_NULL\(\)](#), except it causes an assertion failure (see [KUNIT_ASSERT_TRUE\(\)](#)) when the assertion is not met.

KUNIT_ASSERT_NOT_ERR_OR_NULL

`KUNIT_ASSERT_NOT_ERR_OR_NULL (test, ptr)`

Assertion that **ptr** is not null and not err.

Parameters

test

The test context object.

ptr

an arbitrary pointer.

Description

Sets an assertion that the value that **ptr** evaluates to is not null and not an errno stored in a pointer. This is the same as [`KUNIT_EXPECT_NOT_ERR_OR_NULL\(\)`](#), except it causes an assertion failure (see [`KUNIT_ASSERT_TRUE\(\)`](#)) when the assertion is not met.

KUNIT_ARRAY_PARAM

`KUNIT_ARRAY_PARAM (name, array, get_desc)`

Define test parameter generator from an array.

Parameters

name

prefix for the test parameter generator function.

array

array of test parameters.

get_desc

function to convert param to description; NULL to use default

Description

Define function **name_gen_params** which uses **array** to generate parameters.

KUNIT_ARRAY_PARAM_DESC

`KUNIT_ARRAY_PARAM_DESC (name, array, desc_member)`

Define test parameter generator from an array.

Parameters

name

prefix for the test parameter generator function.

array

array of test parameters.

desc_member

structure member from array element to use as description

Description

Define function **name_gen_params** which uses **array** to generate parameters.

16.7.2 Resource API

This file documents the KUnit resource API.

Most users won't need to use this API directly, power users can use it to store state on a per-test basis, register custom cleanup actions, and more.

struct kunit_resource
 represents a *test managed resource*

Definition:

```
struct kunit_resource {
    void *data;
    const char *name;
    kunit_resource_free_t free;
};
```

Members

data
 for the user to store arbitrary data.

name
 optional name

free
 a user supplied function to free the resource.

Description

Represents a *test managed resource*, a resource which will automatically be cleaned up at the end of a test case. This cleanup is performed by the 'free' function. The [struct kunit_resource](#) itself is freed automatically with kfree() if it was allocated by KUnit (e.g., by [kunit_alloc_resource\(\)](#)), but must be freed by the user otherwise.

Resources are reference counted so if a resource is retrieved via [kunit_alloc_and_get_resource\(\)](#) or [kunit_find_resource\(\)](#), we need to call [kunit_put_resource\(\)](#) to reduce the resource reference count when finished with it. Note that [kunit_alloc_resource\(\)](#) does not require a [kunit_resource_put\(\)](#) because it does not retrieve the resource itself.

```
struct kunit_kmalloc_params {
    size_t size;
    gfp_t gfp;
};

static int kunit_kmalloc_init(struct kunit_resource *res, void *context)
{
    struct kunit_kmalloc_params *params = context;
    res->data = kmalloc(params->size, params->gfp);

    if (!res->data)
        return -ENOMEM;

    return 0;
```

```

}

static void kunit_kmalloc_free(struct kunit_resource *res)
{
    kfree(res->data);
}

void *kunit_kmalloc(struct kunit *test, size_t size, gfp_t gfp)
{
    struct kunit_malloc_params params;

    params.size = size;
    params.gfp = gfp;

    return kunit_alloc_resource(test, kunit_malloc_init,
                                kunit_malloc_free, gfp, &params);
}

```

Resources can also be named, with lookup/removal done on a name basis also. `kunit_add_named_resource()`, `kunit_find_named_resource()` and `kunit_destroy_named_resource()`. Resource names must be unique within the test instance.

Example

`void kunit_get_resource(struct kunit_resource *res)`

Hold resource for use. Should not need to be used by most users as we automatically get resources retrieved by `kunit_find_resource*()`.

Parameters

struct kunit_resource *res
resource

`void kunit_put_resource(struct kunit_resource *res)`

When caller is done with retrieved resource, `kunit_put_resource()` should be called to drop reference count. The resource list maintains a reference count on resources, so if no users are utilizing a resource and it is removed from the resource list, it will be freed via the associated free function (if any). Only needs to be used if we `alloc_and_get()` or `find()` resource.

Parameters

struct kunit_resource *res
resource

`int __kunit_add_resource(struct kunit *test, kunit_resource_init_t init, kunit_resource_free_t free, struct kunit_resource *res, void *data)`

Internal helper to add a resource.

Parameters

struct kunit *test
The test context object.

kunit_resource_init_t init
a user-supplied function to initialize the result (if needed). If none is supplied, the resource

data value is simply set to **data**. If an init function is supplied, **data** is passed to it instead.

kunit_resource_free_t free

a user-supplied function to free the resource (if needed).

struct kunit_resource *res

The resource.

void *data

value to pass to init function or set in resource data field.

Description

res->should_kfree is not initialised.

```
int kunit_add_resource(struct kunit *test, kunit_resource_init_t init, kunit_resource_free_t
                      free, struct kunit_resource *res, void *data)
```

Add a *test managed resource*.

Parameters

struct kunit *test

The test context object.

kunit_resource_init_t init

a user-supplied function to initialize the result (if needed). If none is supplied, the resource data value is simply set to **data**. If an init function is supplied, **data** is passed to it instead.

kunit_resource_free_t free

a user-supplied function to free the resource (if needed).

struct kunit_resource *res

The resource.

void *data

value to pass to init function or set in resource data field.

```
int kunit_add_named_resource(struct kunit *test, kunit_resource_init_t init,
                             kunit_resource_free_t free, struct kunit_resource *res, const
                             char *name, void *data)
```

Add a named *test managed resource*.

Parameters

struct kunit *test

The test context object.

kunit_resource_init_t init

a user-supplied function to initialize the resource data, if needed.

kunit_resource_free_t free

a user-supplied function to free the resource data, if needed.

struct kunit_resource *res

The resource.

const char *name

name to be set for resource.

void *data

value to pass to init function or set in resource data field.

```
struct kunit_resource *kunit_alloc_and_get_resource(struct kunit *test,  
                                              kunit_resource_init_t init,  
                                              kunit_resource_free_t free, gfp_t  
                                              internal_gfp, void *context)
```

Allocates and returns a *test managed resource*.

Parameters

struct kunit *test

The test context object.

kunit_resource_init_t init

a user supplied function to initialize the resource.

kunit_resource_free_t free

a user supplied function to free the resource (if needed).

gfp_t internal_gfp

gfp to use for internal allocations, if unsure, use GFP_KERNEL

void *context

for the user to pass in arbitrary data to the init function.

Description

Allocates a *test managed resource*, a resource which will automatically be cleaned up at the end of a test case. See *struct kunit_resource* for an example.

This is effectively identical to kunit_alloc_resource, but returns the *struct kunit_resource* pointer, not just the 'data' pointer. It therefore also increments the resource's refcount, so *kunit_put_resource()* should be called when you've finished with it.

Note

KUnit needs to allocate memory for a kunit_resource object. You must specify an **internal_gfp** that is compatible with the use context of your resource.

```
void *kunit_alloc_resource(struct kunit *test, kunit_resource_init_t init,  
                           kunit_resource_free_t free, gfp_t internal_gfp, void *context)
```

Allocates a *test managed resource*.

Parameters

struct kunit *test

The test context object.

kunit_resource_init_t init

a user supplied function to initialize the resource.

kunit_resource_free_t free

a user supplied function to free the resource (if needed).

gfp_t internal_gfp

gfp to use for internal allocations, if unsure, use GFP_KERNEL

void *context

for the user to pass in arbitrary data to the init function.

Description

Allocates a *test managed resource*, a resource which will automatically be cleaned up at the end of a test case. See `struct kunit_resource` for an example.

Note

KUnit needs to allocate memory for a `kunit_resource` object. You must specify an `internal_gfp` that is compatible with the use context of your resource.

```
bool kunit_resource_name_match(struct kunit *test, struct kunit_resource *res, void *match_name)
```

Match a resource with the same name.

Parameters

struct kunit *test

Test case to which the resource belongs.

struct kunit_resource *res

The resource.

void *match_name

The name to match against.

```
struct kunit_resource *kunit_find_resource(struct kunit *test, kunit_resource_match_t match, void *match_data)
```

Find a resource using match function/data.

Parameters

struct kunit *test

Test case to which the resource belongs.

kunit_resource_match_t match

match function to be applied to resources/match data.

void *match_data

data to be used in matching.

```
struct kunit_resource *kunit_find_named_resource(struct kunit *test, const char *name)
```

Find a resource using match name.

Parameters

struct kunit *test

Test case to which the resource belongs.

const char *name

match name.

```
int kunit_destroy_resource(struct kunit *test, kunit_resource_match_t match, void *match_data)
```

Find a `kunit_resource` and destroy it.

Parameters

struct kunit *test

Test case to which the resource belongs.

kunit_resource_match_t match

Match function. Returns whether a given resource matches `match_data`.

void *match_data
Data passed into **match**.

Return

0 if kunit_resource is found and freed, -ENOENT if not found.

void kunit_remove_resource(struct kunit *test, struct kunit_resource *res)
remove resource from resource list associated with test.

Parameters

struct kunit *test
The test context object.

struct kunit_resource *res
The resource to be removed.

Description

Note that the resource will not be immediately freed since it is likely the caller has a reference to it via alloc_and_get() or find(); in this case a final call to [kunit_put_resource\(\)](#) is required.

KUNIT_DEFINE_ACTION_WRAPPER

KUNIT_DEFINE_ACTION_WRAPPER (wrapper, orig, arg_type)

Wrap a function for use as a deferred action.

Parameters

wrapper
The name of the new wrapper function define.

orig
The original function to wrap.

arg_type
The type of the argument accepted by **orig**.

Description

Defines a wrapper for a function which accepts a single, pointer-sized argument. This wrapper can then be passed to [kunit_add_action\(\)](#) and similar. This should be used in preference to casting a function directly to kunit_action_t, as casting function pointers will break control flow integrity (CFI), leading to crashes.

int kunit_add_action(struct kunit *test, kunit_action_t *action, void *ctx)

Call a function when the test ends.

Parameters

struct kunit *test
Test case to associate the action with.

kunit_action_t *action
The function to run on test exit

void *ctx
Data passed into **func**

Description

Defer the execution of a function until the test exits, either normally or due to a failure. `ctx` is passed as additional context. All functions registered with `kunit_add_action()` will execute in the opposite order to that they were registered in.

This is useful for cleaning up allocated memory and resources, as these functions are called even if the test aborts early due to, e.g., a failed assertion.

See also: `devm_add_action()` for the devres equivalent.

Return

0 on success, an error if the action could not be deferred.

`int kunit_add_action_or_reset(struct kunit *test, kunit_action_t *action, void *ctx)`

Call a function when the test ends.

Parameters

struct kunit *test

Test case to associate the action with.

kunit_action_t *action

The function to run on test exit

void *ctx

Data passed into `func`

Description

Defer the execution of a function until the test exits, either normally or due to a failure. `ctx` is passed as additional context. All functions registered with `kunit_add_action()` will execute in the opposite order to that they were registered in.

This is useful for cleaning up allocated memory and resources, as these functions are called even if the test aborts early due to, e.g., a failed assertion.

If the action cannot be created (e.g., due to the system being out of memory), then `action(ctx)` will be called immediately, and an error will be returned.

See also: `devm_add_action_or_reset()` for the devres equivalent.

Return

0 on success, an error if the action could not be deferred.

`void kunit_remove_action(struct kunit *test, kunit_action_t *action, void *ctx)`

Cancel a matching deferred action.

Parameters

struct kunit *test

Test case the action is associated with.

kunit_action_t *action

The deferred function to cancel.

void *ctx

The context passed to the deferred function to trigger.

Description

Prevent an action deferred via `kunit_add_action()` from executing when the test terminates. If the function/context pair was deferred multiple times, only the most recent one will be cancelled.

See also: `devm_remove_action()` for the devres equivalent.

```
void kunit_release_action(struct kunit *test, kunit_action_t *action, void *ctx)
```

Run a matching action call immediately.

Parameters

struct kunit *test

Test case the action is associated with.

kunit_action_t *action

The deferred function to trigger.

void *ctx

The context passed to the deferred function to trigger.

Description

Execute a function deferred via `kunit_add_action()` immediately, rather than when the test ends.

If the function/context pair was deferred multiple times, it will only be executed once here. The most recent deferral will no longer execute when the test ends.

`kunit_release_action(test, func, ctx);` is equivalent to `func(ctx); kunit_remove_action(test, func, ctx);`

See also: `devm_release_action()` for the devres equivalent.

Managed Devices

Functions for using KUnit-managed struct device and struct device_driver. Include `kunit/device.h` to use these.

```
struct device_driver *kunit_driver_create(struct kunit *test, const char *name)
```

Create a struct device_driver attached to the kunit_bus

Parameters

struct kunit *test

The test context object.

const char *name

The name to give the created driver.

Description

Creates a struct device_driver attached to the kunit_bus, with the name **name**. This driver will automatically be cleaned up on test exit.

Return

a stub struct device_driver, managed by KUnit, with the name **name**.

```
struct device *kunit_device_register(struct kunit *test, const char *name)
```

Create a struct device for use in KUnit tests

Parameters

struct kunit *test

The test context object.

const char *name

The name to give the created device.

Description

Creates a struct kunit_device (which is a struct device) with the given name, and a corresponding driver. The device and driver will be cleaned up on test exit, or when kunit_device_unregister is called. See also kunit_device_register_with_driver, if you wish to provide your own struct device_driver.

Return

a pointer to a struct device which will be cleaned up when the test exits, or an error pointer if the device could not be allocated or registered.

```
struct device *kunit_device_register_with_driver(struct kunit *test, const char *name,
                                              const struct device_driver *drv)
```

Create a struct device for use in KUnit tests

Parameters

struct kunit *test

The test context object.

const char *name

The name to give the created device.

const struct device_driver *drv

The struct device_driver to associate with the device.

Description

Creates a struct kunit_device (which is a struct device) with the given name, and driver. The device will be cleaned up on test exit, or when kunit_device_unregister is called. See also kunit_device_register, if you wish KUnit to create and manage a driver for you.

Return

a pointer to a struct device which will be cleaned up when the test exits, or an error pointer if the device could not be allocated or registered.

```
void kunit_device_unregister(struct kunit *test, struct device *dev)
```

Unregister a KUnit-managed device

Parameters

struct kunit *test

The test context object which created the device

struct device *dev

The device.

Description

Unregisters and destroys a struct device which was created with kunit_device_register or kunit_device_register_with_driver. If KUnit created a driver, cleans it up as well.

16.7.3 Function Redirection API

Overview

When writing unit tests, it's important to be able to isolate the code being tested from other parts of the kernel. This ensures the reliability of the test (it won't be affected by external factors), reduces dependencies on specific hardware or config options (making the test easier to run), and protects the stability of the rest of the system (making it less likely for test-specific state to interfere with the rest of the system).

While for some code (typically generic data structures, helpers, and other "pure functions") this is trivial, for others (like device drivers, filesystems, core subsystems) the code is heavily coupled with other parts of the kernel.

This coupling is often due to global state in some way: be it a global list of devices, the filesystem, or some hardware state. Tests need to either carefully manage, isolate, and restore state, or they can avoid it altogether by replacing access to and mutation of this state with a "fake" or "mock" variant.

By refactoring access to such state, such as by introducing a layer of indirection which can use or emulate a separate set of test state. However, such refactoring comes with its own costs (and undertaking significant refactoring before being able to write tests is suboptimal).

A simpler way to intercept and replace some of the function calls is to use function redirection via static stubs.

Static Stubs

Static stubs are a way of redirecting calls to one function (the "real" function) to another function (the "replacement" function).

It works by adding a macro to the "real" function which checks to see if a test is running, and if a replacement function is available. If so, that function is called in place of the original.

Using static stubs is pretty straightforward:

1. Add the `KUNIT_STATIC_STUB_REDIRECT()` macro to the start of the "real" function.

This should be the first statement in the function, after any variable declarations. `KUNIT_STATIC_STUB_REDIRECT()` takes the name of the function, followed by all of the arguments passed to the real function.

For example:

```
void send_data_to_hardware(const char *str)
{
    KUNIT_STATIC_STUB_REDIRECT(send_data_to_hardware, str);
    /* real implementation */
}
```

2. Write one or more replacement functions.

These functions should have the same function signature as the real function. In the event they need to access or modify test-specific state, they can use `kunit_get_current_test()` to get a `struct kunit` pointer. This can then be passed to the expectation/assertion macros, or used to look up KUnit resources.

For example:

```
void fake_send_data_to_hardware(const char *str)
{
    struct kunit *test = kunit_get_current_test();
    KUNIT_EXPECT_STREQ(test, str, "Hello World!");
}
```

3. Activate the static stub from your test.

From within a test, the redirection can be enabled with `kunit_activate_static_stub()`, which accepts a `struct kunit` pointer, the real function, and the replacement function. You can call this several times with different replacement functions to swap out implementations of the function.

In our example, this would be

```
kunit_activate_static_stub(test,
                           send_data_to_hardware,
                           fake_send_data_to_hardware);
```

4. Call (perhaps indirectly) the real function.

Once the redirection is activated, any call to the real function will call the replacement function instead. Such calls may be buried deep in the implementation of another function, but must occur from the test's kthread.

For example:

```
send_data_to_hardware("Hello World!"); /* Succeeds */
send_data_to_hardware("Something else"); /* Fails the test. */
```

5. (Optionally) disable the stub.

When you no longer need it, disable the redirection (and hence resume the original behaviour of the 'real' function) using `kunit_deactivate_static_stub()`. Otherwise, it will be automatically disabled when the test exits.

For example:

```
kunit_deactivate_static_stub(test, send_data_to_hardware);
```

It's also possible to use these replacement functions to test to see if a function is called at all, for example:

```
void send_data_to_hardware(const char *str)
{
    KUNIT_STATIC_STUB_REDIRECT(send_data_to_hardware, str);
    /* real implementation */
```

```

}

/* In test file */
int times_called = 0;
void fake_send_data_to_hardware(const char *str)
{
    times_called++;
}
...
/* In the test case, redirect calls for the duration of the test */
kunit_activate_static_stub(test, send_data_to_hardware, fake_send_data_to_
    hardware);

send_data_to_hardware("hello");
KUNIT_EXPECT_EQ(test, times_called, 1);

/* Can also deactivate the stub early, if wanted */
kunit_deactivate_static_stub(test, send_data_to_hardware);

send_data_to_hardware("hello again");
KUNIT_EXPECT_EQ(test, times_called, 1);

```

API Reference

KUNIT_STATIC_STUB_REDIRECT

`KUNIT_STATIC_STUB_REDIRECT (real_fn_name, args...)`

call a replacement 'static stub' if one exists

Parameters

`real_fn_name`

The name of this function (as an identifier, not a string)

`args...`

All of the arguments passed to this function

Description

This is a function prologue which is used to allow calls to the current function to be redirected by a KUnit test. KUnit tests can call `kunit_activate_static_stub()` to pass a replacement function in. The replacement function will be called by `KUNIT_STATIC_STUB_REDIRECT()`, which will then return from the function. If the caller is not in a KUnit context, the function will continue execution as normal.

```

int real_func(int n)
{
    KUNIT_STATIC_STUB_REDIRECT(real_func, n);
    return 0;
}

int replacement_func(int n)

```

```
{
    return 42;
}

void example_test(struct kunit *test)
{
    kunit_activate_static_stub(test, real_func, replacement_func);
    KUNIT_EXPECT_EQ(test, real_func(1), 42);
}
```

Example

kunit_activate_static_stub

kunit_activate_static_stub (test, real_fn_addr, replacement_addr)

replace a function using static stubs.

Parameters

test

A pointer to the '*struct kunit*' test context for the current test.

real_fn_addr

The address of the function to replace.

replacement_addr

The address of the function to replace it with.

Description

When activated, calls to real_fn_addr from within this test (even if called indirectly) will instead call replacement_addr. The function pointed to by real_fn_addr must begin with the static stub prologue in *KUNIT_STATIC_STUB_REDIRECT()* for this to work. real_fn_addr and replacement_addr must have the same type.

The redirection can be disabled again with *kunit_deactivate_static_stub()*.

void kunit_deactivate_static_stub(*struct kunit* *test, void *real_fn_addr)

disable a function redirection

Parameters

***struct kunit* *test**

A pointer to the '*struct kunit*' test context for the current test.

void *real_fn_addr

The address of the function to no-longer redirect

Description

Deactivates a redirection configured with *kunit_activate_static_stub()*. After this function returns, calls to real_fn_addr() will execute the original real_fn, not any previously-configured replacement.

This page documents the KUnit kernel testing API. It is divided into the following sections:

Test API

- Documents all of the standard testing API

Resource API

- Documents the KUnit resource API

Function Redirection API

- Documents the KUnit Function Redirection API

16.8 Test Style and Nomenclature

To make finding, writing, and using KUnit tests as simple as possible, it is strongly encouraged that they are named and written according to the guidelines below. While it is possible to write KUnit tests which do not follow these rules, they may break some tooling, may conflict with other tests, and may not be run automatically by testing systems.

It is recommended that you only deviate from these guidelines when:

1. Porting tests to KUnit which are already known with an existing name.
2. Writing tests which would cause serious problems if automatically run. For example, non-deterministically producing false positives or negatives, or taking a long time to run.

16.8.1 Subsystems, Suites, and Tests

To make tests easy to find, they are grouped into suites and subsystems. A test suite is a group of tests which test a related area of the kernel. A subsystem is a set of test suites which test different parts of a kernel subsystem or a driver.

Subsystems

Every test suite must belong to a subsystem. A subsystem is a collection of one or more KUnit test suites which test the same driver or part of the kernel. A test subsystem should match a single kernel module. If the code being tested cannot be compiled as a module, in many cases the subsystem should correspond to a directory in the source tree or an entry in the MAINTAINERS file. If unsure, follow the conventions set by tests in similar areas.

Test subsystems should be named after the code being tested, either after the module (wherever possible), or after the directory or files being tested. Test subsystems should be named to avoid ambiguity where necessary.

If a test subsystem name has multiple components, they should be separated by underscores. *Do not* include "test" or "kunit" directly in the subsystem name unless we are actually testing other tests or the kunit framework itself. For example, subsystems could be called:

ext4

Matches the module and filesystem name.

apparmor

Matches the module name and LSM name.

kasan

Common name for the tool, prominent part of the path mm/kasan

snd_hda_codec_hdmi

Has several components (snd, hda, codec, hdmi) separated by underscores. Matches the module name.

Avoid names as shown in examples below:

linear-ranges

Names should use underscores, not dashes, to separate words. Prefer `linear_ranges`.

qos-kunit-test

This name should use underscores, and not have "kunit-test" as a suffix. `qos` is also ambiguous as a subsystem name, because several parts of the kernel have a `qos` subsystem. `power_qos` would be a better name.

pc_parallel_port

The corresponding module name is `parport_pc`, so this subsystem should also be named `parport_pc`.

Note: The KUnit API and tools do not explicitly know about subsystems. They are a way of categorizing test suites and naming modules which provides a simple, consistent way for humans to find and run tests. This may change in the future.

Suites

KUnit tests are grouped into test suites, which cover a specific area of functionality being tested. Test suites can have shared initialization and shutdown code which is run for all tests in the suite. Not all subsystems need to be split into multiple test suites (for example, simple drivers).

Test suites are named after the subsystem they are part of. If a subsystem contains several suites, the specific area under test should be appended to the subsystem name, separated by an underscore.

In the event that there are multiple types of test using KUnit within a subsystem (for example, both unit tests and integration tests), they should be put into separate suites, with the type of test as the last element in the suite name. Unless these tests are actually present, avoid using `_test`, `_unittest` or similar in the suite name.

The full test suite name (including the subsystem name) should be specified as the `.name` member of the `kunit_suite` struct, and forms the base for the module name. For example, test suites could include:

ext4_inode

Part of the ext4 subsystem, testing the inode area.

kunit_try_catch

Part of the kunit implementation itself, testing the `try_catch` area.

apparmor_property_entry

Part of the apparmor subsystem, testing the `property_entry` area.

kasan

The kasan subsystem has only one suite, so the suite name is the same as the subsystem name.

Avoid names, for example:

ext4_ext4_inode

There is no reason to state the subsystem twice.

property_entry

The suite name is ambiguous without the subsystem name.

kasan_integration_test

Because there is only one suite in the kasan subsystem, the suite should just be called as kasan. Do not redundantly add integration_test. It should be a separate test suite. For example, if the unit tests are added, then that suite could be named as kasan_unittest or similar.

Test Cases

Individual tests consist of a single function which tests a constrained codepath, property, or function. In the test output, an individual test's results will show up as subtests of the suite's results.

Tests should be named after what they are testing. This is often the name of the function being tested, with a description of the input or codepath being tested. As tests are C functions, they should be named and written in accordance with the kernel coding style.

Note: As tests are themselves functions, their names cannot conflict with other C identifiers in the kernel. This may require some creative naming. It is a good idea to make your test functions *static* to avoid polluting the global namespace.

Example test names include:

unpack_u32_with_null_name

Tests the unpack_u32 function when a NULL name is passed in.

test_list_splice

Tests the list_splice macro. It has the prefix test_ to avoid a name conflict with the macro itself.

Should it be necessary to refer to a test outside the context of its test suite, the *fully-qualified* name of a test should be the suite name followed by the test name, separated by a colon (i.e. suite:test).

16.8.2 Test Kconfig Entries

Every test suite should be tied to a Kconfig entry.

This Kconfig entry must:

- be named CONFIG_<name>_KUNIT_TEST: where <name> is the name of the test suite.
- be listed either alongside the config entries for the driver/subsystem being tested, or be under [Kernel Hacking]->[Kernel Testing and Coverage]
- depend on CONFIG_KUNIT.
- be visible only if CONFIG_KUNIT_ALL_TESTS is not enabled.
- have a default value of CONFIG_KUNIT_ALL_TESTS.

- have a brief description of KUnit in the help text.

If we are not able to meet above conditions (for example, the test is unable to be built as a module), Kconfig entries for tests should be tristate.

For example, a Kconfig entry might look like:

```
config FOO_KUNIT_TEST
    tristate "Kunit test for foo" if !KUNIT_ALL_TESTS
    depends on KUNIT
    default KUNIT_ALL_TESTS
    help
        This builds unit tests for foo.

        For more information on KUnit and unit tests in general,
        please refer to the KUnit documentation in Documentation/dev-tools/
        ↵kunit/.

    If unsure, say N.
```

16.8.3 Test File and Module Names

KUnit tests can often be compiled as a module. These modules should be named after the test suite, followed by `_test`. If this is likely to conflict with non-KUnit tests, the suffix `_kunit` can also be used.

The easiest way of achieving this is to name the file containing the test suite `<suite>_test.c` (or, as above, `<suite>_kunit.c`). This file should be placed next to the code under test.

If the suite name contains some or all of the name of the test's parent directory, it may make sense to modify the source filename to reduce redundancy. For example, a `foo_firmware` suite could be in the `foo/firmware_test.c` file.

16.9 Frequently Asked Questions

16.9.1 How is this different from Autotest, kselftest, and so on?

KUnit is a unit testing framework. Autotest, kselftest (and some others) are not.

A `unit test` is supposed to test a single unit of code in isolation and hence the name *unit test*. A unit test should be the finest granularity of testing and should allow all possible code paths to be tested in the code under test. This is only possible if the code under test is small and does not have any external dependencies outside of the test's control like hardware.

There are no testing frameworks currently available for the kernel that do not require installing the kernel on a test machine or in a virtual machine. All testing frameworks require tests to be written in userspace and run on the kernel under test. This is true for Autotest, kselftest, and some others, disqualifying any of them from being considered unit testing frameworks.

16.9.2 Does KUnit support running on architectures other than UML?

Yes, mostly.

For the most part, the KUnit core framework (what we use to write the tests) can compile to any architecture. It compiles like just another part of the kernel and runs when the kernel boots, or when built as a module, when the module is loaded. However, there is infrastructure, like the KUnit Wrapper (`tools/testing/kunit/kunit.py`) that might not support some architectures (see [Running tests on QEMU](#)).

In short, yes, you can run KUnit on other architectures, but it might require more work than using KUnit on UML.

For more information, see [Writing Tests For Other Architectures](#).

16.9.3 What is the difference between a unit test and other kinds of tests?

Most existing tests for the Linux kernel would be categorized as an integration test, or an end-to-end test.

- A unit test is supposed to test a single unit of code in isolation. A unit test should be the finest granularity of testing and, as such, allows all possible code paths to be tested in the code under test. This is only possible if the code under test is small and does not have any external dependencies outside of the test's control like hardware.
- An integration test tests the interaction between a minimal set of components, usually just two or three. For example, someone might write an integration test to test the interaction between a driver and a piece of hardware, or to test the interaction between the userspace libraries the kernel provides and the kernel itself. However, one of these tests would probably not test the entire kernel along with hardware interactions and interactions with the userspace.
- An end-to-end test usually tests the entire system from the perspective of the code under test. For example, someone might write an end-to-end test for the kernel by installing a production configuration of the kernel on production hardware with a production userspace and then trying to exercise some behavior that depends on interactions between the hardware, the kernel, and userspace.

16.9.4 KUnit is not working, what should I do?

Unfortunately, there are a number of things which can break, but here are some things to try.

1. Run `./tools/testing/kunit/kunit.py run` with the `--raw_output` parameter. This might show details or error messages hidden by the `kunit_tool` parser.
2. Instead of running `kunit.py run`, try running `kunit.py config`, `kunit.py build`, and `kunit.py exec` independently. This can help track down where an issue is occurring. (If you think the parser is at fault, you can run it manually against `stdin` or a file with `kunit.py parse`.)
3. Running the UML kernel directly can often reveal issues or error messages, `kunit_tool` ignores. This should be as simple as running `./vmlinux` after building the UML kernel (for example, by using `kunit.py build`). Note that UML has some unusual requirements (such as the host having a `tmpfs` filesystem mounted), and has had issues in the past when

built statically and the host has KASLR enabled. (On older host kernels, you may need to run `setarch `uname -m` -R ./vmlinux` to disable KASLR.)

4. Make sure the kernel .config has `CONFIG_KUNIT=y` and at least one test (e.g. `CONFIG_KUNIT_EXAMPLE_TEST=y`). `kunit_tool` will keep its .config around, so you can see what config was used after running `kunit.py run`. It also preserves any config changes you might make, so you can enable/disable things with `make ARCH=um menuconfig` or similar, and then re-run `kunit_tool`.
5. Try to run `make ARCH=um defconfig` before running `kunit.py run`. This may help clean up any residual config items which could be causing problems.
6. Finally, try running KUnit outside UML. KUnit and KUnit tests can be built into any kernel, or can be built as a module and loaded at runtime. Doing so should allow you to determine if UML is causing the issue you're seeing. When tests are built-in, they will execute when the kernel boots, and modules will automatically execute associated tests when loaded. Test results can be collected from `/sys/kernel/debug/kunit/<test suite>/results`, and can be parsed with `kunit.py parse`. For more details, see [Running tests on QEMU](#).

If none of the above tricks help, you are always welcome to email any issues to kunit-dev@googlegroups.com.

16.10 Tips For Running KUnit Tests

16.10.1 Using `kunit.py run` ("kunit tool")

Running from any directory

It can be handy to create a bash function like:

```
function run_kunit() {
    ( cd "$(git rev-parse --show-toplevel)" && ./tools/testing/kunit/kunit.py run "$@"
}
```

Note: Early versions of `kunit.py` (before 5.6) didn't work unless run from the kernel root, hence the use of a subshell and `cd`.

Running a subset of tests

`kunit.py run` accepts an optional glob argument to filter tests. The format is "`<suite_glob>[.test_glob]`".

Say that we wanted to run the `sysctl` tests, we could do so via:

```
$ echo -e 'CONFIG_KUNIT=y\nCONFIG_KUNIT_ALL_TESTS=y' > .kunit/.kunitconfig
$ ./tools/testing/kunit/kunit.py run 'sysctl*'
```

We can filter down to just the "write" tests via:

```
$ echo -e 'CONFIG_KUNIT=y\nCONFIG_KUNIT_ALL_TESTS=y' > .kunit/.kunitconfig
$ ./tools/testing/kunit/kunit.py run 'sysctl.*write*'
```

We're paying the cost of building more tests than we need this way, but it's easier than fiddling with `.kunitconfig` files or commenting out `kunit_suite`'s.

However, if we wanted to define a set of tests in a less ad hoc way, the next tip is useful.

Defining a set of tests

`kunit.py run` (along with `build`, and `config`) supports a `--kunitconfig` flag. So if you have a set of tests that you want to run on a regular basis (especially if they have other dependencies), you can create a specific `.kunitconfig` for them.

E.g. `kunit` has one for its tests:

```
$ ./tools/testing/kunit/kunit.py run --kunitconfig=lib/kunit/.kunitconfig
```

Alternatively, if you're following the convention of naming your file `.kunitconfig`, you can just pass in the dir, e.g.

```
$ ./tools/testing/kunit/kunit.py run --kunitconfig=lib/kunit
```

Note: This is a relatively new feature (5.12+) so we don't have any conventions yet about on what files should be checked in versus just kept around locally. It's up to you and your maintainer to decide if a config is useful enough to submit (and therefore have to maintain).

Note: Having `.kunitconfig` fragments in a parent and child directory is iffy. There's discussion about adding an "import" statement in these files to make it possible to have a top-level config run tests from all child directories. But that would mean `.kunitconfig` files are no longer just simple `.config` fragments.

One alternative would be to have `kunit` tool recursively combine configs automagically, but tests could theoretically depend on incompatible options, so handling that would be tricky.

Setting kernel cmdline parameters

You can use `--kernel_args` to pass arbitrary kernel arguments, e.g.

```
$ ./tools/testing/kunit/kunit.py run --kernel_args=param=42 --kernel_
→args=param2=false
```

Generating code coverage reports under UML

Note: TODO(brendanhiggins@google.com): There are various issues with UML and versions of gcc 7 and up. You're likely to run into missing .gcda files or compile errors.

This is different from the "normal" way of getting coverage information that is documented in [Using gcov with the Linux kernel](#).

Instead of enabling CONFIG_GCOV_KERNEL=y, we can set these options:

```
CONFIG_DEBUG_KERNEL=y
CONFIG_DEBUG_INFO=y
CONFIG_DEBUG_INFO_DWARF_TOOLCHAIN_DEFAULT=y
CONFIG_GCOV=y
```

Putting it together into a copy-pastable sequence of commands:

```
# Append coverage options to the current config
$ ./tools/testing/kunit/kunit.py run --kunitconfig=.kunit/ --kunitconfig=tools/
  ↳ testing/kunit/configs/coverage_uml.config
# Extract the coverage information from the build dir (.kunit/)
$ lcov -t "my_kunit_tests" -o coverage.info -c -d .kunit/

# From here on, it's the same process as with CONFIG_GCOV_KERNEL=y
# E.g. can generate an HTML report in a tmp dir like so:
$ genhtml -o /tmp/coverage_html coverage.info
```

If your installed version of gcc doesn't work, you can tweak the steps:

```
$ ./tools/testing/kunit/kunit.py run --make_options=CC=/usr/bin/gcc-6
$ lcov -t "my_kunit_tests" -o coverage.info -c -d .kunit/ --gcov-tool=/usr/bin/
  ↳ gcov-6
```

Alternatively, LLVM-based toolchains can also be used:

```
# Build with LLVM and append coverage options to the current config
$ ./tools/testing/kunit/kunit.py run --make_options LLVM=1 --kunitconfig=.
  ↳ kunit/ --kunitconfig=tools/testing/kunit/configs/coverage_uml.config
$ llvm-profdump merge -sparse default.profraw -o default.propdata
$ llvm-cov export --format=lcov .kunit/vmlinux -instr-profile default.propdata
  ↳> coverage.info
# The coverage.info file is in lcov-compatible format and it can be used to e.
  ↳ g. generate HTML report
$ genhtml -o /tmp/coverage_html coverage.info
```

16.10.2 Running tests manually

Running tests without using `kunit.py run` is also an important use case. Currently it's your only option if you want to test on architectures other than UML.

As running the tests under UML is fairly straightforward (configure and compile the kernel, run the `./linux` binary), this section will focus on testing non-UML architectures.

Running built-in tests

When setting tests to `=y`, the tests will run as part of boot and print results to `dmesg` in TAP format. So you just need to add your tests to your `.config`, build and boot your kernel as normal.

So if we compiled our kernel with:

```
CONFIG_KUNIT=y
CONFIG_KUNIT_EXAMPLE_TEST=y
```

Then we'd see output like this in `dmesg` signaling the test ran and passed:

```
TAP version 14
1..1
# Subtest: example
1..1
# example_simple_test: initializing
ok 1 - example_simple_test
ok 1 - example
```

Running tests as modules

Depending on the tests, you can build them as loadable modules.

For example, we'd change the config options from before to

```
CONFIG_KUNIT=y
CONFIG_KUNIT_EXAMPLE_TEST=m
```

Then after booting into our kernel, we can run the test via

```
$ modprobe kunit-example-test
```

This will then cause it to print TAP output to `stdout`.

Note: The `modprobe` will *not* have a non-zero exit code if any test failed (as of 5.13). But `kunit.py parse` would, see below.

Note: You can set `CONFIG_KUNIT=m` as well, however, some features will not work and thus some tests might break. Ideally tests would specify they depend on `KUNIT=y` in their Kconfig's,

but this is an edge case most test authors won't think about. As of 5.13, the only difference is that `current->kunit_test` will not exist.

Pretty-printing results

You can use `kunit.py parse` to parse dmesg for test output and print out results in the same familiar format that `kunit.py run` does.

```
$ ./tools/testing/kunit/kunit.py parse /var/log/dmesg
```

Retrieving per suite results

Regardless of how you're running your tests, you can enable `CONFIG_KUNIT_DEBUGFS` to expose per-suite TAP-formatted results:

```
CONFIG_KUNIT=y
CONFIG_KUNIT_EXAMPLE_TEST=m
CONFIG_KUNIT_DEBUGFS=y
```

The results for each suite will be exposed under `/sys/kernel/debug/kunit/<suite>/results`. So using our example config:

```
$ modprobe kunit-example-test > /dev/null
$ cat /sys/kernel/debug/kunit/example/results
... < TAP output > ...

# After removing the module, the corresponding files will go away
$ modprobe -r kunit-example-test
$ cat /sys/kernel/debug/kunit/example/results
/sys/kernel/debug/kunit/example/results: No such file or directory
```

Generating code coverage reports

See [Using gcov with the Linux kernel](#) for details on how to do this.

The only vaguely KUnit-specific advice here is that you probably want to build your tests as modules. That way you can isolate the coverage from tests from other code executed during boot, e.g.

```
# Reset coverage counters before running the test.
$ echo 0 > /sys/kernel/debug/gcov/reset
$ modprobe kunit-example-test
```

16.10.3 Test Attributes and Filtering

Test suites and cases can be marked with test attributes, such as speed of test. These attributes will later be printed in test output and can be used to filter test execution.

Marking Test Attributes

Tests are marked with an attribute by including a `kunit_attributes` object in the test definition.

Test cases can be marked using the `KUNIT_CASE_ATTR(test_name, attributes)` macro to define the test case instead of `KUNIT_CASE(test_name)`.

```
static const struct kunit_attributes example_attr = {
    .speed = KUNIT VERY_SLOW,
};

static struct kunit_case example_test_cases[] = {
    KUNIT_CASE_ATTR(example_test, example_attr),
};
```

Note: To mark a test case as slow, you can also use `KUNIT_CASE_SLOW(test_name)`. This is a helpful macro as the slow attribute is the most commonly used.

Test suites can be marked with an attribute by setting the "attr" field in the suite definition.

```
static const struct kunit_attributes example_attr = {
    .speed = KUNIT VERY_SLOW,
};

static struct kunit_suite example_test_suite = {
    ...,
    .attr = example_attr,
};
```

Note: Not all attributes need to be set in a `kunit_attributes` object. Unset attributes will remain uninitialized and act as though the attribute is set to 0 or NULL. Thus, if an attribute is set to 0, it is treated as unset. These unset attributes will not be reported and may act as a default value for filtering purposes.

Reporting Attributes

When a user runs tests, attributes will be present in the raw kernel output (in KTAP format). Note that attributes will be hidden by default in kunit.py output for all passing tests but the raw kernel output can be accessed using the --raw_output flag. This is an example of how test attributes for test cases will be formatted in kernel output:

```
# example_test.speed: slow
ok 1 example_test
```

This is an example of how test attributes for test suites will be formatted in kernel output:

```
KTAP version 2
# Subtest: example_suite
# module: kunit_example_test
1..3
...
ok 1 example_suite
```

Additionally, users can output a full attribute report of tests with their attributes, using the command line flag --list_tests_attr:

```
kunit.py run "example" --list_tests_attr
```

Note: This report can be accessed when running KUnit manually by passing in the module_param kunit.action=list_attr.

Filtering

Users can filter tests using the --filter command line flag when running tests. As an example:

```
kunit.py run --filter speed=slow
```

You can also use the following operations on filters: "<", ">", "<=", ">=", "!=" , and "=" . Example:

```
kunit.py run --filter "speed>slow"
```

This example will run all tests with speeds faster than slow. Note that the characters < and > are often interpreted by the shell, so they may need to be quoted or escaped, as above.

Additionally, you can use multiple filters at once. Simply separate filters using commas. Example:

```
kunit.py run --filter "speed>slow, module=kunit_example_test"
```

Note: You can use this filtering feature when running KUnit manually by passing the filter as a module param: kunit.filter="speed>slow, speed<=normal".

Filtered tests will not run or show up in the test output. You can use the `--filter_action=skip` flag to skip filtered tests instead. These tests will be shown in the test output in the test but will not run. To use this feature when running KUnit manually, use the module param `kunit.filter_action=skip`.

Rules of Filtering Procedure

Since both suites and test cases can have attributes, there may be conflicts between attributes during filtering. The process of filtering follows these rules:

- Filtering always operates at a per-test level.
- If a test has an attribute set, then the test's value is filtered on.
- Otherwise, the value falls back to the suite's value.
- If neither are set, the attribute has a global "default" value, which is used.

List of Current Attributes

`speed`

This attribute indicates the speed of a test's execution (how slow or fast the test is).

This attribute is saved as an enum with the following categories: "normal", "slow", or "very_slow". The assumed default speed for tests is "normal". This indicates that the test takes a relatively trivial amount of time (less than 1 second), regardless of the machine it is running on. Any test slower than this could be marked as "slow" or "very_slow".

The macro `KUNIT_CASE_SLOW(test_name)` can be easily used to set the speed of a test case to "slow".

`module`

This attribute indicates the name of the module associated with the test.

This attribute is automatically saved as a string and is printed for each suite. Tests can also be filtered using this attribute.

`is_init`

This attribute indicates whether the test uses init data or functions.

This attribute is automatically saved as a boolean and tests can also be filtered using this attribute.

This section details the kernel unit testing framework.

16.11 Introduction

KUnit (Kernel unit testing framework) provides a common framework for unit tests within the Linux kernel. Using KUnit, you can define groups of test cases called test suites. The tests either run on kernel boot if built-in, or load as a module. KUnit automatically flags and reports failed test cases in the kernel log. The test results appear in [KTAP \(Kernel - Test Anything Protocol\) format](#). It is inspired by JUnit, Python's unittest.mock, and GoogleTest/GoogleMock (C++ unit testing framework).

KUnit tests are part of the kernel, written in the C (programming) language, and test parts of the Kernel implementation (example: a C language function). Excluding build time, from invocation to completion, KUnit can run around 100 tests in less than 10 seconds. KUnit can test any kernel component, for example: file system, system calls, memory management, device drivers and so on.

KUnit follows the white-box testing approach. The test has access to internal system functionality. KUnit runs in kernel space and is not restricted to things exposed to user-space.

In addition, KUnit has `kunit_tool`, a script (`tools/testing/kunit/kunit.py`) that configures the Linux kernel, runs KUnit tests under QEMU or UML (User Mode Linux), parses the test results and displays them in a user friendly manner.

16.11.1 Features

- Provides a framework for writing unit tests.
- Runs tests on any kernel architecture.
- Runs a test in milliseconds.

16.11.2 Prerequisites

- Any Linux kernel compatible hardware.
- For Kernel under test, Linux kernel version 5.5 or greater.

16.12 Unit Testing

A unit test tests a single unit of code in isolation. A unit test is the finest granularity of testing and allows all possible code paths to be tested in the code under test. This is possible if the code under test is small and does not have any external dependencies outside of the test's control like hardware.

16.12.1 Write Unit Tests

To write good unit tests, there is a simple but powerful pattern: Arrange-Act-Assert. This is a great way to structure test cases and defines an order of operations.

- Arrange inputs and targets: At the start of the test, arrange the data that allows a function to work. Example: initialize a statement or object.
- Act on the target behavior: Call your function/code under test.
- Assert expected outcome: Verify that the result (or resulting state) is as expected.

16.12.2 Unit Testing Advantages

- Increases testing speed and development in the long run.
- Detects bugs at initial stage and therefore decreases bug fix cost compared to acceptance testing.
- Improves code quality.
- Encourages writing testable code.

Read also [*What is the difference between a unit test and other kinds of tests?*](#).

16.13 How do I use it?

You can find a step-by-step guide to writing and running KUnit tests in [*Getting Started*](#)

Alternatively, feel free to look through the rest of the KUnit documentation, or to experiment with tools/testing/kunit/kunit.py and the example test under lib/kunit/kunit-example-test.c

Happy testing!

THE KERNEL TEST ANYTHING PROTOCOL (KTAP), VERSION 1

TAP, or the Test Anything Protocol is a format for specifying test results used by a number of projects. Its website and specification are found at this [link](#). The Linux Kernel largely uses TAP output for test results. However, Kernel testing frameworks have special needs for test results which don't align with the original TAP specification. Thus, a "Kernel TAP" (KTAP) format is specified to extend and alter TAP to support these use-cases. This specification describes the generally accepted format of KTAP as it is currently used in the kernel.

KTAP test results describe a series of tests (which may be nested: i.e., test can have subtests), each of which can contain both diagnostic data -- e.g., log lines -- and a final result. The test structure and results are machine-readable, whereas the diagnostic data is unstructured and is there to aid human debugging.

KTAP output is built from four different types of lines: - Version lines - Plan lines - Test case result lines - Diagnostic lines

In general, valid KTAP output should also form valid TAP output, but some information, in particular nested test results, may be lost. Also note that there is a stagnant draft specification for TAP14, KTAP diverges from this in a couple of places (notably the "Subtest" header), which are described where relevant later in this document.

17.1 Version lines

All KTAP-formatted results begin with a "version line" which specifies which version of the (K)TAP standard the result is compliant with.

For example: - "KTAP version 1" - "TAP version 13" - "TAP version 14"

Note that, in KTAP, subtests also begin with a version line, which denotes the start of the nested test results. This differs from TAP14, which uses a separate "Subtest" line.

While, going forward, "KTAP version 1" should be used by compliant tests, it is expected that most parsers and other tooling will accept the other versions listed here for compatibility with existing tests and frameworks.

17.2 Plan lines

A test plan provides the number of tests (or subtests) in the KTAP output.

Plan lines must follow the format of "1..N" where N is the number of tests or subtests. Plan lines follow version lines to indicate the number of nested tests.

While there are cases where the number of tests is not known in advance -- in which case the test plan may be omitted -- it is strongly recommended one is present where possible.

17.3 Test case result lines

Test case result lines indicate the final status of a test. They are required and must have the format:

```
<result> <number> [<description>][ # [<directive>] [<diagnostic data>]]
```

The result can be either "ok", which indicates the test case passed, or "not ok", which indicates that the test case failed.

<number> represents the number of the test being performed. The first test must have the number 1 and the number then must increase by 1 for each additional subtest within the same test at the same nesting level.

The description is a description of the test, generally the name of the test, and can be any string of characters other than # or a newline. The description is optional, but recommended.

The directive and any diagnostic data is optional. If either are present, they must follow a hash sign, "#".

A directive is a keyword that indicates a different outcome for a test other than passed and failed. The directive is optional, and consists of a single keyword preceding the diagnostic data. In the event that a parser encounters a directive it doesn't support, it should fall back to the "ok" / "not ok" result.

Currently accepted directives are:

- "SKIP", which indicates a test was skipped (note the result of the test case result line can be either "ok" or "not ok" if the SKIP directive is used)
- "TODO", which indicates that a test is not expected to pass at the moment, e.g. because the feature it is testing is known to be broken. While this directive is inherited from TAP, its use in the kernel is discouraged.
- "XFAIL", which indicates that a test is expected to fail. This is similar to "TODO", above, and is used by some kselftest tests.
- "TIMEOUT", which indicates a test has timed out (note the result of the test case result line should be "not ok" if the TIMEOUT directive is used)
- "ERROR", which indicates that the execution of a test has failed due to a specific error that is included in the diagnostic data. (note the result of the test case result line should be "not ok" if the ERROR directive is used)

The diagnostic data is a plain-text field which contains any additional details about why this result was produced. This is typically an error message for ERROR or failed tests, or a description of missing dependencies for a SKIP result.

The diagnostic data field is optional, and results which have neither a directive nor any diagnostic data do not need to include the "#" field separator.

Example result lines include:

```
ok 1 test_case_name
```

The test "test_case_name" passed.

```
not ok 1 test_case_name
```

The test "test_case_name" failed.

```
ok 1 test # SKIP necessary dependency unavailable
```

The test "test" was SKIPPED with the diagnostic message "necessary dependency unavailable".

```
not ok 1 test # TIMEOUT 30 seconds
```

The test "test" timed out, with diagnostic data "30 seconds".

```
ok 5 check return code # rcode=0
```

The test "check return code" passed, with additional diagnostic data "rcode=0"

17.4 Diagnostic lines

If tests wish to output any further information, they should do so using "diagnostic lines". Diagnostic lines are optional, freeform text, and are often used to describe what is being tested and any intermediate results in more detail than the final result and diagnostic data line provides.

Diagnostic lines are formatted as "# <diagnostic_description>", where the description can be any string. Diagnostic lines can be anywhere in the test output. As a rule, diagnostic lines regarding a test are directly before the test result line for that test.

Note that most tools will treat unknown lines (see below) as diagnostic lines, even if they do not start with a "#": this is to capture any other useful kernel output which may help debug the test. It is nevertheless recommended that tests always prefix any diagnostic output they have with a "#" character.

17.5 Unknown lines

There may be lines within KTAP output that do not follow the format of one of the four formats for lines described above. This is allowed, however, they will not influence the status of the tests.

This is an important difference from TAP. Kernel tests may print messages to the system console or a log file. Both of these destinations may contain messages either from unrelated kernel or userspace activity, or kernel messages from non-test code that is invoked by the test. The kernel code invoked by the test likely is not aware that a test is in progress and thus can not print the message as a diagnostic message.

17.6 Nested tests

In KTAP, tests can be nested. This is done by having a test include within its output an entire set of KTAP-formatted results. This can be used to categorize and group related tests, or to split out different results from the same test.

The "parent" test's result should consist of all of its subtests' results, starting with another KTAP version line and test plan, and end with the overall result. If one of the subtests fail, for example, the parent test should also fail.

Additionally, all lines in a subtest should be indented. One level of indentation is two spaces: " ". The indentation should begin at the version line and should end before the parent test's result line.

"Unknown lines" are not considered to be lines in a subtest and thus are allowed to be either indented or not indented.

An example of a test with two nested subtests:

```
KTAP version 1
1..1
    KTAP version 1
    1..2
        ok 1 test_1
        not ok 2 test_2
# example failed
not ok 1 example
```

An example format with multiple levels of nested testing:

```
KTAP version 1
1..2
    KTAP version 1
    1..2
        KTAP version 1
        1..2
            not ok 1 test_1
            ok 2 test_2
            not ok 1 test_3
            ok 2 test_4 # SKIP
```

```
not ok 1 example_test_1
ok 2 example_test_2
```

17.7 Major differences between TAP and KTAP

Feature	TAP	KTAP
yaml and json in diagnostic message	ok	not recommended
TODO directive	ok	not recognized
allows an arbitrary number of tests to be nested	no	yes
"Unknown lines" are in category of "Anything else"	yes	no
"Unknown lines" are	incorrect	allowed

The TAP14 specification does permit nested tests, but instead of using another nested version line, uses a line of the form "Subtest: <name>" where <name> is the name of the parent test.

17.8 Example KTAP output

```
KTAP version 1
1..1
  KTAP version 1
    1..3
      KTAP version 1
        1..1
        # test_1: initializing test_1
        ok 1 test_1
      ok 1 example_test_1
      KTAP version 1
        1..2
        ok 1 test_1 # SKIP test_1 skipped
        ok 2 test_2
      ok 2 example_test_2
      KTAP version 1
        1..3
        ok 1 test_1
        # test_2: FAIL
        not ok 2 test_2
        ok 3 test_3 # SKIP test_3 skipped
      not ok 3 example_test_3
    not ok 1 main_test
```

This output defines the following hierarchy:

A single test called "main_test", which fails, and has three subtests: - "example_test_1", which passes, and has one subtest:

- "test_1", which passes, and outputs the diagnostic message "test_1: initializing test_1"
- "example_test_2", which passes, and has two subtests:

- "test_1", which is skipped, with the explanation "test_1 skipped"
- "test_2", which passes
- "example_test_3", which fails, and has three subtests
 - "test_1", which passes
 - "test_2", which outputs the diagnostic line "test_2: FAIL", and fails.
 - "test_3", which is skipped with the explanation "test_3 skipped"

Note that the individual subtests with the same names do not conflict, as they are found in different parent tests. This output also exhibits some sensible rules for "bubbling up" test results: a test fails if any of its subtests fail. Skipped tests do not affect the result of the parent test (though it often makes sense for a test to be marked skipped if `_all_` of its subtests have been skipped).

17.9 See also:

- The TAP specification: <https://testanything.org/tap-version-13-specification.html>
- The (stagnant) TAP version 14 specification: <https://github.com/TestAnything/Specification/blob/tap-14-specification/specification.md>
- The kselftest documentation: *Linux Kernel Selftests*
- The KUnit documentation: *KUnit - Linux Kernel Unit Testing*

UAPI CHECKER

The UAPI checker (`scripts/check-uapi.sh`) is a shell script which checks UAPI header files for userspace backwards-compatibility across the git tree.

18.1 Options

This section will describe the options with which `check-uapi.sh` can be run.

Usage:

```
check-uapi.sh [-b BASE_REF] [-p PAST_REF] [-j N] [-l ERROR_LOG] [-i] [-q] [-v]
```

Available options:

<code>-b BASE_REF</code>	Base git reference to use for comparison. If unspecified or → empty,
	will use any dirty changes in tree to UAPI files. If there are → no
	dirty changes, HEAD will be used.
<code>-p PAST_REF</code>	Compare BASE_REF to PAST_REF (e.g. -p v6.1). If unspecified or → empty,
	will use BASE_REF ¹ . Must be an ancestor of BASE_REF. Only → headers
	that exist on PAST_REF will be checked for compatibility.
<code>-j JOBS</code>	Number of checks to run in parallel (default: number of CPU → cores).
<code>-l ERROR_LOG</code>	Write error log to file (default: no error log is generated).
<code>-i</code>	Ignore ambiguous changes that may or may not break UAPI → compatibility.
<code>-q</code>	Quiet operation.
<code>-v</code>	Verbose operation (print more information about each header → being checked).

Environmental args:

<code>ABIDIFF</code>	Custom path to abidiff binary
<code>CC</code>	C compiler (default is "gcc")
<code>ARCH</code>	Target architecture of C compiler (default is host arch)

Exit codes:

- 0) Success
- 1) ABI difference detected
- 2) Prerequisite not met

18.2 Examples

18.2.1 Basic Usage

First, let's try making a change to a UAPI header file that obviously won't break userspace:

```
cat << 'EOF' | patch -l -p1
--- a/include/uapi/linux/acct.h
+++ b/include/uapi/linux/acct.h
@@ -21,7 +21,9 @@
#include <asm/param.h>
#include <asm/byteorder.h>

-/*
+#define FOO
+
+/*
 * comp_t is a 16-bit "floating" point number with a 3-bit base 8
 * exponent and a 13-bit fraction.
 * comp2_t is 24-bit with 5-bit base 2 exponent and 20 bit fraction
diff --git a/include/uapi/linux/bpf.h b/include/uapi/linux/bpf.h
EOF
```

Now, let's use the script to validate:

```
% ./scripts/check-uapi.sh
Installing user-facing UAPI headers from dirty tree... OK
Installing user-facing UAPI headers from HEAD... OK
Checking changes to UAPI headers between HEAD and dirty tree...
All 912 UAPI headers compatible with x86 appear to be backwards compatible
```

Let's add another change that *might* break userspace:

```
cat << 'EOF' | patch -l -p1
--- a/include/uapi/linux/bpf.h
+++ b/include/uapi/linux/bpf.h
@@ -74,7 +74,7 @@ struct bpf_insn {
    __u8    dst_reg:4;      /* dest register */
    __u8    src_reg:4;      /* source register */
    __s16   off;           /* signed offset */
-    __s32   imm;           /* signed immediate constant */
+    __u32   imm;           /* unsigned immediate constant */
};
```

```
/* Key of an a BPF_MAP_TYPE_LPM_TRIE entry */
EOF
```

The script will catch this:

```
% ./scripts/check-uapi.sh
Installing user-facing UAPI headers from dirty tree... OK
Installing user-facing UAPI headers from HEAD... OK
Checking changes to UAPI headers between HEAD and dirty tree...
===== ABI differences detected in include/linux/bpf.h from HEAD -> dirty tree
=====
[C] 'struct bpf_insn' changed:
  type size hasn't changed
  1 data member change:
    type of '__s32 imm' changed:
      typedef name changed from __s32 to __u32 at int-ll64.h:27:1
      underlying type 'int' changed:
        type name changed from 'int' to 'unsigned int'
        type size hasn't changed
=====
error - 1/912 UAPI headers compatible with x86 appear _not_ to be backwards
compatible
```

In this case, the script is reporting the type change because it could break a userspace program that passes in a negative number. Now, let's say you know that no userspace program could possibly be using a negative value in `imm`, so changing to an unsigned type there shouldn't hurt anything. You can pass the `-i` flag to the script to ignore changes in which the userspace backwards compatibility is ambiguous:

```
% ./scripts/check-uapi.sh -i
Installing user-facing UAPI headers from dirty tree... OK
Installing user-facing UAPI headers from HEAD... OK
Checking changes to UAPI headers between HEAD and dirty tree...
All 912 UAPI headers compatible with x86 appear to be backwards compatible
```

Now, let's make a similar change that *will* break userspace:

```
cat << 'EOF' | patch -l -pl
--- a/include/uapi/linux/bpf.h
+++ b/include/uapi/linux/bpf.h
@@ -71,8 +71,8 @@ enum {

  struct bpf_insn {
    __u8    code;          /* opcode */
-    __u8    dst_reg:4;     /* dest register */
-    __u8    src_reg:4;     /* source register */
+    __u8    dst_reg:4;     /* dest register */
    __s16   off;           /* signed offset */
    __s32   imm;           /* signed immediate constant */
  };
EOF
```

EOF

Since we're re-ordering an existing struct member, there's no ambiguity, and the script will report the breakage even if you pass `-i`:

```
% ./scripts/check-uapi.sh -i
Installing user-facing UAPI headers from dirty tree... OK
Installing user-facing UAPI headers from HEAD... OK
Checking changes to UAPI headers between HEAD and dirty tree...
==== ABI differences detected in include/linux/bpf.h from HEAD -> dirty tree
=====[C] 'struct bpf_insn' changed:
    type size hasn't changed
    2 data member changes:
        '__u8 dst_reg' offset changed from 8 to 12 (in bits) (by +4 bits)
        '__u8 src_reg' offset changed from 12 to 8 (in bits) (by -4 bits)
=====
error - 1/912 UAPI headers compatible with x86 appear _not_ to be backwards compatible
```

Let's commit the breaking change, then commit the innocuous change:

```
% git commit -m 'Breaking UAPI change' include/uapi/linux/bpf.h
[detached HEAD f758e574663a] Breaking UAPI change
 1 file changed, 1 insertion(+), 1 deletion(-)
% git commit -m 'Innocuous UAPI change' include/uapi/linux/acct.h
[detached HEAD 2e87df769081] Innocuous UAPI change
 1 file changed, 3 insertions(+), 1 deletion(-)
```

Now, let's run the script again with no arguments:

```
% ./scripts/check-uapi.sh
Installing user-facing UAPI headers from HEAD... OK
Installing user-facing UAPI headers from HEAD^1... OK
Checking changes to UAPI headers between HEAD^1 and HEAD...
All 912 UAPI headers compatible with x86 appear to be backwards compatible
```

It doesn't catch any breaking change because, by default, it only compares HEAD to HEAD¹. The breaking change was committed on HEAD~2. If we wanted the search scope to go back further, we'd have to use the `-p` option to pass a different past reference. In this case, let's pass `-p HEAD~2` to the script so it checks UAPI changes between HEAD~2 and HEAD:

```
% ./scripts/check-uapi.sh -p HEAD~2
Installing user-facing UAPI headers from HEAD... OK
Installing user-facing UAPI headers from HEAD~2... OK
Checking changes to UAPI headers between HEAD~2 and HEAD...
==== ABI differences detected in include/linux/bpf.h from HEAD~2 -> HEAD ====
[C] 'struct bpf_insn' changed:
    type size hasn't changed
    2 data member changes:
        '__u8 dst_reg' offset changed from 8 to 12 (in bits) (by +4 bits)
```

```
'__u8 src_reg' offset changed from 12 to 8 (in bits) (by -4 bits)
=====
error - 1/912 UAPI headers compatible with x86 appear _not_ to be backwardsu
compatible
```

Alternatively, we could have also run with `-b HEAD~`. This would set the base reference to `HEAD~` so then the script would compare it to `HEAD~^1`.

18.2.2 Architecture-specific Headers

Consider this change:

```
cat << 'EOF' | patch -l -p1
--- a/arch/arm64/include/uapi/asm/sigcontext.h
+++ b/arch/arm64/include/uapi/asm/sigcontext.h
@@ -70,6 +70,7 @@ struct sigcontext {
    struct _aarch64_ctx {
        __u32 magic;
        __u32 size;
+       __u32 new_var;
    };

#define FPSIMD_MAGIC 0x46508001
EOF
```

This is a change to an arm64-specific UAPI header file. In this example, I'm running the script from an x86 machine with an x86 compiler, so, by default, the script only checks x86-compatible UAPI header files:

```
% ./scripts/check-uapi.sh
Installing user-facing UAPI headers from dirty tree... OK
Installing user-facing UAPI headers from HEAD... OK
No changes to UAPI headers were applied between HEAD and dirty tree
```

With an x86 compiler, we can't check header files in `arch/arm64`, so the script doesn't even try. If we want to check the header file, we'll have to use an arm64 compiler and set `ARCH` accordingly:

```
% CC=aarch64-linux-gnu-gcc ARCH=arm64 ./scripts/check-uapi.sh
Installing user-facing UAPI headers from dirty tree... OK
Installing user-facing UAPI headers from HEAD... OK
Checking changes to UAPI headers between HEAD and dirty tree...
===== ABI differences detected in include/asm/sigcontext.h from HEAD -> dirtyu
tree ====
[C] 'struct _aarch64_ctx' changed:
    type size changed from 64 to 96 (in bits)
    1 data member insertion:
        '__u32 new_var', at offset 64 (in bits) at sigcontext.h:73:1
-- snip --
```

```
[C] 'struct zt_context' changed:
  type size changed from 128 to 160 (in bits)
  2 data member changes (1 filtered):
    '__u16 nregs' offset changed from 64 to 96 (in bits) (by +32 bits)
    '__u16 __reserved[3]' offset changed from 80 to 112 (in bits) (by +32 bits)
=====

error - 1/884 UAPI headers compatible with arm64 appear _not_ to be backwards-compatible
```

We can see with ARCH and CC set properly for the file, the ABI change is reported properly. Also notice that the total number of UAPI header files checked by the script changes. This is because the number of headers installed for arm64 platforms is different than x86.

18.2.3 Cross-Dependency Breakages

Consider this change:

```
cat << 'EOF' | patch -l -p1
--- a/include/uapi/linux/types.h
+++ b/include/uapi/linux/types.h
@@ -52,7 +52,7 @@ typedef __u32 __bitwise __wsum;
#define __aligned_be64 __be64 __attribute__((aligned(8)))
#define __aligned_le64 __le64 __attribute__((aligned(8)))

-typedef unsigned __bitwise __poll_t;
+typedef unsigned short __bitwise __poll_t;

#endif /* __ASSEMBLY__ */
#endif /* _UAPI_LINUX_TYPES_H */
EOF
```

Here, we're changing a `typedef` in `types.h`. This doesn't break a UAPI in `types.h`, but other UAPIs in the tree may break due to this change:

```
% ./scripts/check-uapi.sh
Installing user-facing UAPI headers from dirty tree... OK
Installing user-facing UAPI headers from HEAD... OK
Checking changes to UAPI headers between HEAD and dirty tree...
==== ABI differences detected in include/linux/eventpoll.h from HEAD -> dirty
tree ====
[C] 'struct epoll_event' changed:
  type size changed from 96 to 80 (in bits)
  2 data member changes:
    type of '__poll_t events' changed:
      underlying type 'unsigned int' changed:
        type name changed from 'unsigned int' to 'unsigned short int'
        type size changed from 32 to 16 (in bits)
    '__u64 data' offset changed from 32 to 16 (in bits) (by -16 bits)
```

```
=====
include/linux/eventpoll.h did not change between HEAD and dirty tree...
It's possible a change to one of the headers it includes caused this error:
#include <linux/fcntl.h>
#include <linux/types.h>
```

Note that the script noticed the failing header file did not change, so it assumes one of its includes must have caused the breakage. Indeed, we can see `linux/types.h` is used from `eventpoll.h`.

18.2.4 UAPI Header Removals

Consider this change:

```
cat << 'EOF' | patch -l -p1
diff --git a/include/uapi/asm-generic/Kbuild b/include/uapi/asm-generic/Kbuild
index ebb180aac74e..a9c88b0a8b3b 100644
--- a/include/uapi/asm-generic/Kbuild
+++ b/include/uapi/asm-generic/Kbuild
@@ -31,6 +31,6 @@ mandatory-y += stat.h
 mandatory-y += statfs.h
 mandatory-y += swab.h
 mandatory-y += termbits.h
-mandatory-y += termios.h
+#mandatory-y += termios.h
 mandatory-y += types.h
 mandatory-y += unistd.h
EOF
```

This script removes a UAPI header file from the install list. Let's run the script:

```
% ./scripts/check-uapi.sh
Installing user-facing UAPI headers from dirty tree... OK
Installing user-facing UAPI headers from HEAD... OK
Checking changes to UAPI headers between HEAD and dirty tree...
===== UAPI header include/asm/termios.h was removed between HEAD and dirty tree_
→=====

error - 1/912 UAPI headers compatible with x86 appear _not_ to be backwards_
→compatible
```

Removing a UAPI header is considered a breaking change, and the script will flag it as such.

18.2.5 Checking Historic UAPI Compatibility

You can use the `-b` and `-p` options to examine different chunks of your git tree. For example, to check all changed UAPI header files between tags v6.0 and v6.1, you'd run:

```
% ./scripts/check-uapi.sh -b v6.1 -p v6.0
Installing user-facing UAPI headers from v6.1... OK
Installing user-facing UAPI headers from v6.0... OK
Checking changes to UAPI headers between v6.0 and v6.1...

--- snip ---
error - 37/907 UAPI headers compatible with x86 appear _not_ to be backwardsin
compatible
```

Note: Before v5.3, a header file needed by the script is not present, so the script is unable to check changes before then.

You'll notice that the script detected many UAPI changes that are not backwards compatible. Knowing that kernel UAPIs are supposed to be stable forever, this is an alarming result. This brings us to the next section: caveats.

18.3 Caveats

The UAPI checker makes no assumptions about the author's intention, so some types of changes may be flagged even though they intentionally break UAPI.

18.3.1 Removals For Refactoring or Deprecation

Sometimes drivers for very old hardware are removed, such as in this example:

```
% ./scripts/check-uapi.sh -b ba47652ba655
Installing user-facing UAPI headers from ba47652ba655... OK
Installing user-facing UAPI headers from ba47652ba655^1... OK
Checking changes to UAPI headers between ba47652ba655^1 and ba47652ba655...
===== UAPI header include/linux/meye.h was removed between ba47652ba655^1 and ba47652ba655
=====

error - 1/910 UAPI headers compatible with x86 appear _not_ to be backwardsin
compatible
```

The script will always flag removals (even if they're intentional).

18.3.2 Struct Expansions

Depending on how a structure is handled in kernelspace, a change which expands a struct could be non-breaking.

If a struct is used as the argument to an ioctl, then the kernel driver must be able to handle ioctl commands of any size. Beyond that, you need to be careful when copying data from the user. Say, for example, that `struct foo` is changed like this:

```
struct foo {
    __u64 a; /* added in version 1 */
+    __u32 b; /* added in version 2 */
+    __u32 c; /* added in version 2 */
}
```

By default, the script will flag this kind of change for further review:

```
[C] 'struct foo' changed:
  type size changed from 64 to 128 (in bits)
  2 data member insertions:
    '__u32 b', at offset 64 (in bits)
    '__u32 c', at offset 96 (in bits)
```

However, it is possible that this change was made safely.

If a userspace program was built with version 1, it will think `sizeof(struct foo)` is 8. That size will be encoded in the ioctl value that gets sent to the kernel. If the kernel is built with version 2, it will think the `sizeof(struct foo)` is 16.

The kernel can use the `_IOC_SIZE` macro to get the size encoded in the ioctl code that the user passed in and then use `copy_struct_from_user()` to safely copy the value:

```
int handle_ioctl(unsigned long cmd, unsigned long arg)
{
    switch _IOC_NR(cmd) {
    0x01: {
        struct foo my_cmd; /* size 16 in the kernel */

        ret = copy_struct_from_user(&my_cmd, arg, sizeof(struct foo), _IOC_
        ↪SIZE(cmd));
        ...
    }
}
```

`copy_struct_from_user` will zero the struct in the kernel and then copy only the bytes passed in from the user (leaving new members zeroized). If the user passed in a larger struct, the extra members are ignored.

If you know this situation is accounted for in the kernel code, you can pass `-i` to the script, and struct expansions like this will be ignored.

18.3.3 Flex Array Migration

While the script handles expansion into an existing flex array, it does still flag initial migration to flex arrays from 1-element fake flex arrays. For example:

```
struct foo {
    __u32 x;
-    __u32 flex[1]; /* fake flex */
+    __u32 flex[]; /* real flex */
};
```

This change would be flagged by the script:

```
[C] 'struct foo' changed:
  type size changed from 64 to 32 (in bits)
  1 data member change:
    type of '__u32 flex[1]' changed:
      type name changed from '__u32[1]' to '__u32[]'
      array type size changed from 32 to 'unknown'
      array type subrange 1 changed length from 1 to 'unknown'
```

At this time, there's no way to filter these types of changes, so be aware of this possible false positive.

18.3.4 Summary

While many types of false positives are filtered out by the script, it's possible there are some cases where the script flags a change which does not break UAPI. It's also possible a change which *does* break userspace would not be flagged by this script. While the script has been run on much of the kernel history, there could still be corner cases that are not accounted for.

The intention is for this script to be used as a quick check for maintainers or automated tooling, not as the end-all authority on patch compatibility. It's best to remember: use your best judgment (and ideally a unit test in userspace) to make sure your UAPI changes are backwards-compatible!

INDEX

Symbols

`__kfence_free (C function)`, 100
`__kunit_add_resource (C function)`, 200

A

`ASSERT_EQ (C macro)`, 140
`ASSERT_EXCLUSIVE_ACCESS (C macro)`, 88
`ASSERT_EXCLUSIVE_ACCESS_SCOPED (macro)`, 89
`ASSERT_EXCLUSIVE_BITS (C macro)`, 89
`ASSERT_EXCLUSIVE_WRITER (C macro)`, 87
`ASSERT_EXCLUSIVE_WRITER_SCOPED (macro)`, 87
`ASSERT_FALSE (C macro)`, 142
`ASSERT_GE (C macro)`, 142
`ASSERT_GT (C macro)`, 141
`ASSERT_LE (C macro)`, 141
`ASSERT_LT (C macro)`, 141
`ASSERT_NE (C macro)`, 141
`ASSERT_NULL (C macro)`, 142
`ASSERT_STREQ (C macro)`, 142
`ASSERT_STRNE (C macro)`, 143
`ASSERT_TRUE (C macro)`, 142

E

`EXPECT_EQ (C macro)`, 143
`EXPECT_FALSE (C macro)`, 145
`EXPECT_GE (C macro)`, 144
`EXPECT_GT (C macro)`, 144
`EXPECT_LE (C macro)`, 144
`EXPECT_LT (C macro)`, 143
`EXPECT_NE (C macro)`, 143
`EXPECT_NULL (C macro)`, 144
`EXPECT_STREQ (C macro)`, 145
`EXPECT_STRNE (C macro)`, 145
`EXPECT_TRUE (C macro)`, 145

F

`Fixture (C macro)`, 138
`Fixture_Data (C macro)`, 137
`Fixture_Setup (C macro)`, 138

`Fixture_TearDown (C macro)`, 138
`Fixture_Variant (C macro)`, 139
`Fixture_Variant_Add (C macro)`, 139

G

`gdb_regs_to_pt_regs (C function)`, 122

`(C) I`
`is_kfence_address (C function)`, 98

K

`(C) kfence_alloc (C function)`, 99
`kfence_free (C function)`, 100
`kfence_handle_page_fault (C function)`, 100
`kfence_ksize (C function)`, 99
`kfence_object_start (C function)`, 100
`kfence_shutdown_cache (C function)`, 99
`kgdb_arch (C struct)`, 123
`kgdb_arch_exit (C function)`, 121
`kgdb_arch_handle_exception (C function)`, 122
`kgdb_arch_handle_qxfer_pkt (C function)`, 122
`kgdb_arch_init (C function)`, 121
`kgdb_arch_late (C function)`, 123
`kgdb_arch_set_pc (C function)`, 123
`kgdb_breakpoint (C function)`, 120
`kgdb_call_nmi_hook (C function)`, 122
`kgdb_io (C struct)`, 124
`kgdb_roundup_cpus (C function)`, 123
`kgdb_skipexception (C function)`, 120
`kunit (C struct)`, 181
`kunit_activate_static_stub (C macro)`, 211
`kunit_add_action (C function)`, 204
`kunit_add_action_or_reset (C function)`, 205
`kunit_add_named_resource (C function)`, 201
`kunit_add_resource (C function)`, 201
`kunit_alloc_and_get_resource (C function)`, 202
`kunit_alloc_resource (C function)`, 202

KUNIT_ARRAY_PARAM (*C macro*), 198
 KUNIT_ARRAY_PARAM_DESC (*C macro*), 198
 KUNIT_ASSERT_EQ (*C macro*), 193
 KUNIT_ASSERT_FALSE (*C macro*), 193
 KUNIT_ASSERT_GE (*C macro*), 196
 KUNIT_ASSERT_GT (*C macro*), 195
 KUNIT_ASSERT_LE (*C macro*), 195
 KUNIT_ASSERT_LT (*C macro*), 195
 KUNIT_ASSERT_NE (*C macro*), 194
 KUNIT_ASSERT_NOT_ERR_OR_NULL (*C macro*),
 197
 KUNIT_ASSERT_NOT_NULL (*C macro*), 197
 KUNIT_ASSERT_NULL (*C macro*), 197
 KUNIT_ASSERT_PTR_EQ (*C macro*), 193
 KUNIT_ASSERT_PTR_NE (*C macro*), 194
 KUNIT_ASSERT_STREQ (*C macro*), 196
 KUNIT_ASSERT_STRNEQ (*C macro*), 196
 KUNIT_ASSERT_TRUE (*C macro*), 192
 KUNIT_CASE (*C macro*), 179
 kunit_case (*C struct*), 178
 KUNIT_CASE_ATTR (*C macro*), 179
 KUNIT_CASE_PARAM (*C macro*), 180
 KUNIT_CASE_PARAM_ATTR (*C macro*), 180
 KUNIT_CASE_SLOW (*C macro*), 180
 kunit_deactivate_static_stub (*C function*),
 211
 KUNIT_DEFINE_ACTION_WRAPPER (*C macro*),
 204
 kunit_destroy_resource (*C function*), 203
 kunit_device_register (*C function*), 206
 kunit_device_register_with_driver (*C
 function*), 207
 kunit_device_unregister (*C function*), 207
 kunit_driver_create (*C function*), 206
 kunit_err (*C macro*), 185
 KUNIT_EXPECT_EQ (*C macro*), 187
 KUNIT_EXPECT_FALSE (*C macro*), 187
 KUNIT_EXPECT_GE (*C macro*), 189
 KUNIT_EXPECT_GT (*C macro*), 189
 KUNIT_EXPECT_LE (*C macro*), 189
 KUNIT_EXPECT_LT (*C macro*), 188
 KUNIT_EXPECT_MEMEQ (*C macro*), 191
 KUNIT_EXPECT_MEMNEQ (*C macro*), 191
 KUNIT_EXPECT_NE (*C macro*), 188
 KUNIT_EXPECT_NOT_ERR_OR_NULL (*C macro*),
 192
 KUNIT_EXPECT_NOT_NULL (*C macro*), 192
 KUNIT_EXPECT_NULL (*C macro*), 192
 KUNIT_EXPECT_PTR_EQ (*C macro*), 187
 KUNIT_EXPECT_PTR_NE (*C macro*), 188
 KUNIT_EXPECT_STREQ (*C macro*), 190
 KUNIT_EXPECT_STRNEQ (*C macro*), 190
 KUNIT_EXPECT_TRUE (*C macro*), 186
 KUNIT_FAIL (*C macro*), 186
 kunit_find_named_resource (*C function*),
 203
 kunit_find_resource (*C function*), 203
 kunit_get_resource (*C function*), 200
 kunit_info (*C macro*), 185
 kunit_kcalloc (*C function*), 184
 kunit_kfree (*C function*), 183
 kunit_kmalloc (*C function*), 183
 kunit_kmalloc_array (*C function*), 182
 kunit_kzalloc (*C function*), 183
 kunit_mark_skipped (*C macro*), 184
 kunit_put_resource (*C function*), 200
 kunit_release_action (*C function*), 206
 kunit_remove_action (*C function*), 205
 kunit_remove_resource (*C function*), 204
 kunit_resource (*C struct*), 199
 kunit_resource_name_match (*C function*),
 203
 kunit_skip (*C macro*), 184
 KUNIT_STATIC_STUB_REDIRECT (*C macro*), 210
 kunit_status (*C enum*), 178
 KUNIT_SUCCEED (*C macro*), 186
 kunit_suite (*C struct*), 180
 kunit_test_init_section_suites (*C
 macro*), 182
 kunit_test_suites (*C macro*), 182
 kunit_warn (*C macro*), 185

P

pt_regs_to_gdb_regs (*C function*), 121

S

sleeping_thread_to_gdb_regs (*C function*),
 121

T

TEST (*C macro*), 137
 TEST_F (*C macro*), 140
 TEST_HARNESS_MAIN (*C macro*), 140
 TEST_SIGNAL (*C macro*), 137
 TH_LOG (*C macro*), 136