
Linux Peci Documentation

Release 6.8.0

The kernel development community

Jan 16, 2026

CONTENTS

1	Overview	1
1.1	PECI Wire	1
1.2	PECI subsystem internals	1
1.3	PECI CPU Driver API	6
Index		11

OVERVIEW

The Platform Environment Control Interface (PECI) is a communication interface between Intel processor and management controllers (e.g. Baseboard Management Controller, BMC). PECI provides services that allow the management controller to configure, monitor and debug platform by accessing various registers. It defines a dedicated command protocol, where the management controller is acting as a PECI originator and the processor - as a PECI responder. PECI can be used in both single processor and multiple-processor based systems.

NOTE: Intel PECI specification is not released as a dedicated document, instead it is a part of External Design Specification (EDS) for given Intel CPU. External Design Specifications are usually not publicly available.

1.1 PECI Wire

PECI Wire interface uses a single wire for self-clocking and data transfer. It does not require any additional control lines - the physical layer is a self-clocked one-wire bus signal that begins each bit with a driven, rising edge from an idle near zero volts. The duration of the signal driven high allows to determine whether the bit value is logic '0' or logic '1'. PECI Wire also includes variable data rate established with every message.

For PECI Wire, each processor package will utilize unique, fixed addresses within a defined range and that address should have a fixed relationship with the processor socket ID - if one of the processors is removed, it does not affect addresses of remaining processors.

1.2 PECI subsystem internals

struct peci_controller_ops

 PECI controller specific methods

Definition:

```
struct peci_controller_ops {  
    int (*xfer)(struct peci_controller *controller, u8 addr, struct peci_  
    -request *req);  
};
```

Members

xfer

 PECI transfer function

Description

PECI controllers may have different hardware interfaces - the drivers implementing PECI controllers can use this structure to abstract away those differences by exposing a common interface for PECI core.

struct peci_controller

 PECI controller

Definition:

```
struct peci_controller {  
    struct device dev;  
    const struct peci_controller_ops *ops;  
    struct mutex bus_lock;  
    u8 id;  
};
```

Members

dev

device object to register PECI controller to the device model

ops

pointer to device specific controller operations

bus_lock

lock used to protect multiple callers

id

PECI controller ID

Description

PECI controllers usually connect to their drivers using non-PECI bus, such as the platform bus. Each PECI controller can communicate with one or more PECI devices.

struct peci_device

 PECI device

Definition:

```
struct peci_device {  
    struct device dev;  
    struct {  
        u16 family;  
        u8 model;  
        u8 peci_revision;  
        u8 socket_id;  
    } info;  
    u8 addr;  
    bool deleted;  
};
```

Members

dev

device object to register PECl device to the device model

info

PECI device characteristics

info.family

device family

info.model

device model

info.peci_revision

PECI revision supported by the PECl device

info.socket_id

the socket ID represented by the PECl device

addr

address used on the PECl bus connected to the parent controller

deleted

indicates that PECl device was already deleted

Description

A peci_device identifies a single device (i.e. CPU) connected to a PECl bus. The behaviour exposed to the rest of the system is defined by the PECl driver managing the device.

struct peci_request

PECI request

Definition:

```
struct peci_request {
    struct peci_device *device;
    struct {
        u8 buf[PECI_REQUEST_MAX_BUF_SIZE];
        u8 len;
    } rx, tx;
};
```

Members**device**

PECI device to which the request is sent

rx

RX buffer specific data

rx.buf

RX buffer

rx.len

received data length in bytes

tx

TX buffer specific data

tx.buf

TX buffer

tx.len

transfer data length in bytes

Description

A peci_request represents a request issued by PECI originator (TX) and a response received from PECI responder (RX).

struct peci_device_id

PECI device data to match

Definition:

```
struct peci_device_id {  
    const void *data;  
    u16 family;  
    u8 model;  
};
```

Members

data

pointer to driver private data specific to device

family

device family

model

device model

struct peci_driver

PECI driver

Definition:

```
struct peci_driver {  
    struct device_driver driver;  
    int (*probe)(struct peci_device *device, const struct peci_device_id *id);  
    void (*remove)(struct peci_device *device);  
    const struct peci_device_id *id_table;  
};
```

Members

driver

inherit device driver

probe

probe callback

remove

remove callback

id_table

PECI device match table to decide which device to bind

peci_driver_register

```
peci_driver_register (driver)
```

register PECI driver

Parameters**driver**

the driver to be registered

Description

PECI drivers that don't need to do anything special in module init should use the convenience "module_peci_driver" macro instead

Return

zero on success, else a negative error code.

module_peci_driver

```
module_peci_driver (__peci_driver)
```

helper macro for registering a modular PECI driver

Parameters**__peci_driver**

peci_driver struct

Description

Helper macro for PECI drivers which do not do anything special in module init/exit. This eliminates a lot of boilerplate. Each module may only use this macro once, and calling it replaces module_init() and module_exit()

```
struct peci_controller *devm_peci_controller_add(struct device *dev, const struct  
                                                 peci_controller_ops *ops)
```

add PECI controller

Parameters**struct device *dev**

device for devm operations

const struct peci_controller_ops *ops

pointer to controller specific methods

Description

In final stage of its probe(), peci_controller driver calls *devm_peci_controller_add()* to register itself with the PECI bus.

Return

Pointer to the newly allocated controller or ERR_PTR() in case of failure.

```
int peci_request_status(struct peci_request *req)
```

return -errno based on PECI completion code

Parameters

struct peci_request *req

the PECI request that contains response data with completion code

Description

It can't be used for Ping(), GetDIB() and GetTemp() - for those commands we don't expect completion code in the response.

Return

-errno

struct *peci_request* *peci_request_alloc(*struct peci_device* *device, u8 tx_len, u8 rx_len)
allocate struct peci_requests

Parameters

struct peci_device *device

PECI device to which request is going to be sent

u8 tx_len

TX length

u8 rx_len

RX length

Return

A pointer to a newly allocated *struct peci_request* on success or NULL otherwise.

void peci_request_free(*struct peci_request* *req)
free peci_request

Parameters

struct peci_request *req

the PECI request to be freed

1.3 PECI CPU Driver API

int peci_temp_read(*struct peci_device* *device, s16 *temp_raw)

read the maximum die temperature from PECI target device

Parameters

struct peci_device *device

PECI device to which request is going to be sent

s16 *temp_raw

where to store the read temperature

Description

It uses GetTemp PECI command.

Return

0 if succeeded, other values in case errors.

```
int peci_pcs_read(struct peci_device *device, u8 index, u16 param, u32 *data)
    read PCS register
```

Parameters**struct peci_device *device**

PECI device to which request is going to be sent

u8 index

PCS index

u16 param

PCS parameter

u32 *data

where to store the read data

Description

It uses RdPkgConfig PECI command.

Return

0 if succeeded, other values in case errors.

```
int peci_pci_local_read(struct peci_device *device, u8 bus, u8 dev, u8 func, u16 reg, u32
    *data)
```

read 32-bit memory location using raw address

Parameters**struct peci_device *device**

PECI device to which request is going to be sent

u8 bus

bus

u8 dev

device

u8 func

function

u16 reg

register

u32 *data

where to store the read data

Description

It uses RdPCIConfigLocal PECI command.

Return

0 if succeeded, other values in case errors.

```
int peci_ep_pci_local_read(struct peci_device *device, u8 seg, u8 bus, u8 dev, u8 func, u16
    reg, u32 *data)
```

read 32-bit memory location using raw address

Parameters

```
struct peci_device *device
    PECI device to which request is going to be sent

u8 seg
    PCI segment

u8 bus
    bus

u8 dev
    device

u8 func
    function

u16 reg
    register

u32 *data
    where to store the read data
```

Description

Like [*peci_pci_local_read*](#), but it uses RdEndpointConfig PECI command.

Return

0 if succeeded, other values in case errors.

```
int peci_mmio_read(struct peci_device *device, u8 bar, u8 seg, u8 bus, u8 dev, u8 func, u64
                    address, u32 *data)

    read 32-bit memory location using 64-bit bar offset address
```

Parameters

```
struct peci_device *device
    PECI device to which request is going to be sent

u8 bar
    PCI bar

u8 seg
    PCI segment

u8 bus
    bus

u8 dev
    device

u8 func
    function

u64 address
    64-bit MMIO address

u32 *data
    where to store the read data
```

Description

It uses RdEndpointConfig PECI command.

Return

0 if succeeded, other values in case errors.

INDEX

D

`devm_peci_controller_add (C function)`, 5

M

`module_peci_driver (C macro)`, 5

P

`peci_controller (C struct)`, 2

`peci_controller_ops (C struct)`, 1

`peci_device (C struct)`, 2

`peci_device_id (C struct)`, 4

`peci_driver (C struct)`, 4

`peci_driver_register (C macro)`, 4

`peci_ep_pci_local_read (C function)`, 7

`peci_mmio_read (C function)`, 8

`peci_pci_local_read (C function)`, 7

`peci_pcs_read (C function)`, 6

`peci_request (C struct)`, 3

`peci_request_alloc (C function)`, 6

`peci_request_free (C function)`, 6

`peci_request_status (C function)`, 5

`peci_temp_read (C function)`, 6